

Programming Assignment 2: Convolutional Neural Networks

Version 1.5

- Fixed the bug in the `compute_loss` function in part A

Version Release Date: 2022-02-05

Due Date: Friday, Feb. 18, at 11:59pm

Based on an assignment by Lisa Zhang

For CSC413/2516 in Winter 2022 with Professors Jimmy Ba and Bo Wang

Submission: You must submit two files through [MarkUs](#): a PDF file containing your writeup, titled *a2-writeup.pdf*, and your code file *a2-code.ipynb*. Your writeup must be typeset.

The programming assignments are individual work. See the Course Syllabus for detailed policies.

Introduction:

This assignment will focus on the applications of convolutional neural networks in various image processing tasks. First, we will train a convolutional neural network for a task known as image colourization. Given a greyscale image, we will predict the colour at each pixel. This is a difficult problem for many reasons, one of which being that it is ill-posed: for a single greyscale image, there can be multiple, equally valid colourings.

In the second half of the assignment, we switch gears and perform object detection by fine-tuning a pre-trained model. Specifically, we use the YOLOv3 ([Redmon and Farhadi, 2018](#)) pre-trained model and fine-tune it on the COCO ([Lin et al., 2014](#)) dataset.

Colab FAQ and Using GPU

For some basic overview and features offered in Colab notebooks, check out: [Overview of Colaboratory Features](#).

You need to use the Colab GPU for this assignment by selecting:

Runtime → Change runtime type → Hardware Accelerator: GPU

Download CIFAR and Colour dictionary

We will use the [CIFAR-10 data set](#), which consists of images of size 32x32 pixels. For most of the questions we will use a subset of the dataset. To make the problem easier, we will only use the “Horse” category from this data set. Now let’s learn to colour some horses!

The data loading script is included below. It can take up to a couple of minutes to download everything the first time.

All files are stored at `/content/csc413/a2/data/` folder.

Helper code

You can ignore the restart warning.

```
#####
# Setup working directory
#####
%mkdir -p /content/csc413/a2/
%cd /content/csc413/a2

#####
# Helper functions for loading data
#####
# adapted from
#
https://github.com/fchollet/keras/blob/master/keras/datasets/cifar10.p
y

import os
import pickle
import sys
import tarfile

import numpy as np
from PIL import Image
from six.moves.urllib.request import urlretrieve

def get_file(fname, origin, untar=False, extract=False,
archive_format="auto", cache_dir="data"):
    datadir = os.path.join(cache_dir)
    if not os.path.exists(datadir):
        os.makedirs(datadir)

    if untar:
        untar_fpath = os.path.join(datadir, fname)
        fpath = untar_fpath + ".tar.gz"
    else:
        fpath = os.path.join(datadir, fname)

    print("File path: %s" % fpath)
    if not os.path.exists(fpath):
        print("Downloading data from", origin)

        error_msg = "URL fetch failure on {}: {} -- {}"
```

```

        try:
            try:
                urlretrieve(origin, fpath)
            except URLError as e:
                raise Exception(error_msg.format(origin, e.errno,
e.reason))
            except HTTPError as e:
                raise Exception(error_msg.format(origin, e.code,
e.msg))
        except (Exception, KeyboardInterrupt) as e:
            if os.path.exists(fpath):
                os.remove(fpath)
            raise

    if untar:
        if not os.path.exists(untar_fpath):
            print("Extracting file.")
            with tarfile.open(fpath) as archive:
                archive.extractall(datadir)
            return untar_fpath

    if extract:
        _extract_archive(fpath, datadir, archive_format)

    return fpath


def load_batch(fpath, label_key="labels"):
    """Internal utility for parsing CIFAR data.
    # Arguments
        fpath: path the file to parse.
        label_key: key for label data in the retrieve
        dictionary.
    # Returns
        A tuple `(data, labels)`.
    """
    f = open(fpath, "rb")
    if sys.version_info < (3,):
        d = pickle.load(f)
    else:
        d = pickle.load(f, encoding="bytes")
        # decode utf8
        d_decoded = {}
        for k, v in d.items():
            d_decoded[k.decode("utf8")] = v
        d = d_decoded
    f.close()
    data = d["data"]
    labels = d[label_key]

```

```
data = data.reshape(data.shape[0], 3, 32, 32)
return data, labels
```

```
def load_cifar10(transpose=False):
    """Loads CIFAR10 dataset.
    # Returns
    Tuple of Numpy arrays: `(x_train, y_train), (x_test, y_test)`.
    """
    dirname = "cifar-10-batches-py"
    origin = "http://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz"
    path = get_file(dirname, origin=origin, untar=True)

    num_train_samples = 50000

    x_train = np.zeros((num_train_samples, 3, 32, 32), dtype="uint8")
    y_train = np.zeros((num_train_samples,), dtype="uint8")

    for i in range(1, 6):
        fpath = os.path.join(path, "data_batch_" + str(i))
        data, labels = load_batch(fpath)
        x_train[(i - 1) * 10000 : i * 10000, :, :, :] = data
        y_train[(i - 1) * 10000 : i * 10000] = labels

    fpath = os.path.join(path, "test_batch")
    x_test, y_test = load_batch(fpath)

    y_train = np.reshape(y_train, (len(y_train), 1))
    y_test = np.reshape(y_test, (len(y_test), 1))

    if transpose:
        x_train = x_train.transpose(0, 2, 3, 1)
        x_test = x_test.transpose(0, 2, 3, 1)
    return (x_train, y_train), (x_test, y_test)
```

/content/csc413/a2

Download files

This may take 1 or 2 mins for the first time.

```
# Download cluster centers for k-means over colours
colours_fpath = get_file(
    fname="colours",
    origin="http://www.cs.toronto.edu/~jba/kmeans_colour_a2.tar.gz",
    untar=True
)
# Download CIFAR dataset
m = load_cifar10()
```

```
File path: data/colours.tar.gz
Downloading data from
http://www.cs.toronto.edu/~jba/kmeans_colour_a2.tar.gz
Extracting file.
File path: data/cifar-10-batches-py.tar.gz
Downloading data from http://www.cs.toronto.edu/~kriz/cifar-10-
python.tar.gz
Extracting file.
```

Image Colourization as Classification

We will select a subset of 24 colours and frame colourization as a pixel-wise classification problem, where we label each pixel with one of 24 colours. The 24 colours are selected using [k-means clustering](#) over colours, and selecting cluster centers.

This was already done for you, and cluster centers are provided in http://www.cs.toronto.edu/~jba/kmeans_colour_a2.tar.gz, which was downloaded by the helper functions above. For simplicity, we will measure distance in RGB space. This is not ideal but reduces the software dependencies for this assignment.

Helper code

```
"""
Colourization of CIFAR-10 Horses via classification.
"""
```

```
import argparse
import math
import time

import matplotlib
import matplotlib.pyplot as plt
import numpy as np
import numpy.random as npr
import scipy.misc
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.autograd import Variable

# from load_data import load_cifar10

HORSE_CATEGORY = 7
```

Data related code

```
def get_rgb_cat(xs, colours):
    """
    Get colour categories given RGB values. This function doesn't
    actually do the work, instead it splits the work into smaller
    chunks that can fit into memory, and calls helper function
    """
```

`_get_rgb_cat`

Args:

xs: float numpy array of RGB images in [B, C, H, W] format

colours: numpy array of colour categories and their RGB values

Returns:

result: int numpy array of shape [B, 1, H, W]

"""

if np.shape(xs)[0] < 100:

return _get_rgb_cat(xs)

batch_size = 100

nexts = []

for i **in** range(0, np.shape(xs)[0], batch_size):

 next = _get_rgb_cat(xs[i : i + batch_size, :, :, :], colours)

 nexts.append(next)

result = np.concatenate(nexts, axis=0)

return result

def _get_rgb_cat(xs, colours):

"""

Get colour categories given RGB values. This is done by choosing the colour in `colours` that is the closest (in RGB space) to each point in the image `xs`. This function is a little memory intensive, and so the size of `xs` should not be too large.

Args:

xs: float numpy array of RGB images in [B, C, H, W] format

colours: numpy array of colour categories and their RGB values

Returns:

result: int numpy array of shape [B, 1, H, W]

"""

num_colours = np.shape(colours)[0]

xs = np.expand_dims(xs, 0)

cs = np.reshape(colours, [num_colours, 1, 3, 1, 1])

dists = np.linalg.norm(xs - cs, axis=2) # 2 = colour axis

cat = np.argmin(dists, axis=0)

cat = np.expand_dims(cat, axis=1)

return cat

def get_cat_rgb(cats, colours):

"""

Get RGB colours given the colour categories

Args:

cats: integer numpy array of colour categories

colours: numpy array of colour categories and their RGB values

Returns:

```

    numpy tensor of RGB colours
    """
    return colours[cats]

def process(xs, ys, max_pixel=256.0, downsize_input=False):
    """
    Pre-process CIFAR10 images by taking only the horse category,
    shuffling, and have colour values be bound between 0 and 1

    Args:
        xs: the colour RGB pixel values
        ys: the category labels
        max_pixel: maximum pixel value in the original data
    Returns:
        xs: value normalized and shuffled colour images
        grey: greyscale images, also normalized so values are between 0
and 1
    """
    xs = xs / max_pixel
    xs = xs[np.where(ys == HORSE_CATEGORY)[0], :, :, :]
    npr.shuffle(xs)

    grey = np.mean(xs, axis=1, keepdims=True)

    if downsize_input:
        downsize_module = nn.Sequential(
            nn.AvgPool2d(2),
            nn.AvgPool2d(2),
            nn.Upsample(scale_factor=2),
            nn.Upsample(scale_factor=2),
        )
        xs_downsized =
downsize_module.forward(torch.from_numpy(xs).float())
        xs_downsized = xs_downsized.data.numpy()
        return (xs, xs_downsized)
    else:
        return (xs, grey)

def get_batch(x, y, batch_size):
    """
    Generated that yields batches of data

    Args:
        x: input values
        y: output values
        batch_size: size of each batch
    Yields:

```

```

    batch_x: a batch of inputs of size at most batch_size
    batch_y: a batch of outputs of size at most batch_size
    """
    N = np.shape(x)[0]
    assert N == np.shape(y)[0]
    for i in range(0, N, batch_size):
        batch_x = x[i : i + batch_size, :, :, :]
        batch_y = y[i : i + batch_size, :, :, :]
        yield (batch_x, batch_y)

```

Torch helper

```

def get_torch_vars(xs, ys, gpu=False):
    """

```

*Helper function to convert numpy arrays to pytorch tensors.
If GPU is used, move the tensors to GPU.*

Args:

*xs (float numpy tenosor): greyscale input
ys (int numpy tenosor): categorical labels
gpu (bool): whether to move pytorch tensor to GPU*

Returns:

Variable(xs), Variable(ys)

```

    """

```

```

xs = torch.from_numpy(xs).float()
ys = torch.from_numpy(ys).long()
if gpu:
    xs = xs.cuda()
    ys = ys.cuda()
return Variable(xs), Variable(ys)

```

```

def compute_loss(criterion, outputs, labels, batch_size, num_colours):
    """

```

Helper function to compute the loss. Since this is a pixelwise prediction task we need to reshape the output and ground truth tensors into a 2D tensor before passing it in to the loss

criterion.

Args:

*criterion: pytorch loss criterion
outputs (pytorch tensor): predicted labels from the model
labels (pytorch tensor): ground truth labels
batch_size (int): batch size used for training
num_colours (int): number of colour categories*

Returns:

pytorch tensor for loss

```

    """

```

```

batch = outputs.size(0)
loss_out = outputs.transpose(1, 3).contiguous().view([batch * 32 *
32, num_colours])

```



```

    loss_lab = labels.transpose(1, 3).contiguous().view([batch * 32 *
32])
    return criterion(loss_out, loss_lab)

def run_validation_step(
    cnn,
    criterion,
    test_grey,
    test_rgb_cat,
    batch_size,
    colours,
    plotpath=None,
    visualize=True,
    downsize_input=False
):
    correct = 0.0
    total = 0.0
    losses = []
    num_colours = np.shape(colours)[0]
    for i, (xs, ys) in enumerate(get_batch(test_grey, test_rgb_cat,
batch_size)):
        images, labels = get_torch_vars(xs, ys, args.gpu)
        outputs = cnn(images)
        val_loss = compute_loss(
            criterion, outputs, labels, batch_size=args.batch_size,
num_colours=num_colours
        )
        losses.append(val_loss.data.item())

        _, predicted = torch.max(outputs.data, 1, keepdim=True)
        total += labels.size(0) * 32 * 32
        correct += (predicted == labels.data).sum()

    if plotpath: # only plot if a path is provided
        plot(
            xs,
            ys,
            predicted.cpu().numpy(),
            colours,
            plotpath,
            visualize=visualize,
            compare_bilinear=downsize_input,
        )

    val_loss = np.mean(losses)
    val_acc = 100 * correct / total
    return val_loss, val_acc

```

Visualization

```
def plot(input, gtlabel, output, colours, path, visualize,
        compare_bilinear=False):
    """
    Generate png plots of input, ground truth, and outputs

    Args:
        input: the greyscale input to the colourization CNN
        gtlabel: the growth truth categories for each pixel
        output: the predicted categories for each pixel
        colours: numpy array of colour categories and their RGB values
        path: output path
        visualize: display the figures inline or save the figures in
    path
    """
    grey = np.transpose(input[:10, :, :, :], [0, 2, 3, 1])
    gtcolor = get_cat_rgb(gtlabel[:10, 0, :, :], colours)
    predcolor = get_cat_rgb(output[:10, 0, :, :], colours)

    img_stack = [np.hstack(np.tile(grey, [1, 1, 1, 3])),
                  np.hstack(gtcolor), np.hstack(predcolor)]

    if compare_bilinear:
        downsize_module = nn.Sequential(
            nn.AvgPool2d(2),
            nn.AvgPool2d(2),
            nn.Upsample(scale_factor=2, mode="bilinear"),
            nn.Upsample(scale_factor=2, mode="bilinear"),
        )
        gt_input = np.transpose(
            gtcolor,
            [
                0,
                3,
                1,
                2
            ],
        )
        color_bilinear =
        downsize_module.forward(torch.from_numpy(gt_input).float())
        color_bilinear = np.transpose(color_bilinear.data.numpy(), [0,
        2, 3, 1])
        img_stack = [
            np.hstack(np.transpose(input[:10, :, :, :], [0, 2, 3,
        1])),
            np.hstack(gtcolor),
            np.hstack(predcolor),
            np.hstack(color_bilinear),
        ]
    img = np.vstack(img_stack)
```

```

plt.grid(None)
plt.imshow(img, vmin=0.0, vmax=1.0)
if visualize:
    plt.show()
else:
    plt.savefig(path)

def toimage(img, cmin, cmax):
    return Image.fromarray((img.clip(cmin, cmax) *
255).astype(np.uint8))

def plot_activation(args, cnn):
    # LOAD THE COLOURS CATEGORIES
    colours = np.load(args.colours, allow_pickle=True)[0]
    num_colours = np.shape(colours)[0]

    (x_train, y_train), (x_test, y_test) = load_cifar10()
    test_rgb, test_grey = process(x_test, y_test,
downsize_input=args.downsize_input)
    test_rgb_cat = get_rgb_cat(test_rgb, colours)

    # Take the idnex of the test image
    id = args.index
    outdir = "outputs/" + args.experiment_name + "/act" + str(id)
    if not os.path.exists(outdir):
        os.makedirs(outdir)
    images, labels = get_torch_vars(
        np.expand_dims(test_grey[id], 0),
np.expand_dims(test_rgb_cat[id], 0)
    )
    cnn.cpu()
    outputs = cnn(images)
    _, predicted = torch.max(outputs.data, 1, keepdim=True)
    predcolor = get_cat_rgb(predicted.cpu().numpy()[0, 0, :, :],
colours)
    img = predcolor
    toimage(predcolor, cmin=0, cmax=1).save(os.path.join(outdir,
"output_%d.png" % id))

    if not args.downsize_input:
        img = np.tile(np.transpose(test_grey[id], [1, 2, 0]), [1, 1,
3])
    else:
        img = np.transpose(test_grey[id], [1, 2, 0])
    toimage(img, cmin=0, cmax=1).save(os.path.join(outdir, "input_
%d.png" % id))

```

```

img = np.transpose(test_rgb[id], [1, 2, 0])
toimage(img, cmin=0, cmax=1).save(os.path.join(outdir, "input_
%d_gt.png" % id))

def add_border(img):
    return np.pad(img, 1, "constant", constant_values=1.0)

def draw_activations(path, activation, imgwidth=4):
    img = np.vstack(
        [
            np.hstack(
                [
                    add_border(filter)
                    for filter in activation[i * imgwidth : (i +
1) * imgwidth, :, :]
                ]
            )
            for i in range(activation.shape[0] // imgwidth)
        ]
    )
    scipy.misc.imsave(path, img)

for i, tensor in enumerate([cnn.out1, cnn.out2, cnn.out3,
cnn.out4, cnn.out5]):
    draw_activations(
        os.path.join(outdir, "conv%d_out_%d.png" % (i + 1, id)),
        tensor.data.cpu().numpy()[0]
    )
    print("visualization results are saved to %s" % outdir)

```

Training

```

class AttrDict(dict):
    def __init__(self, *args, **kwargs):
        super(AttrDict, self).__init__(*args, **kwargs)
        self.__dict__ = self

def train(args, cnn=None):
    # Set the maximum number of threads to prevent crash in Teaching
    # Labs
    # TODO: necessary?
    torch.set_num_threads(5)
    # Numpy random seed
    npr.seed(args.seed)

    # Save directory
    save_dir = "outputs/" + args.experiment_name

```

```

# LOAD THE COLOURS CATEGORIES
colours = np.load(args.colours, allow_pickle=True,
encoding="bytes")[0]
num_colours = np.shape(colours)[0]
# INPUT CHANNEL
num_in_channels = 1 if not args.downsize_input else 3
# LOAD THE MODEL
if cnn is None:
    Net = globals()[args.model]
    cnn = Net(args.kernel, args.num_filters, num_colours,
num_in_channels)

# LOSS FUNCTION
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(cnn.parameters(), lr=args.learn_rate)

# DATA
print("Loading data...")
(x_train, y_train), (x_test, y_test) = load_cifar10()

print("Transforming data...")
train_rgb, train_grey = process(x_train, y_train,
downsize_input=args.downsize_input)
train_rgb_cat = get_rgb_cat(train_rgb, colours)
test_rgb, test_grey = process(x_test, y_test,
downsize_input=args.downsize_input)
test_rgb_cat = get_rgb_cat(test_rgb, colours)

# Create the outputs folder if not created already
if not os.path.exists(save_dir):
    os.makedirs(save_dir)

print("Beginning training ...")
if args.gpu:
    cnn.cuda()
start = time.time()

train_losses = []
valid_losses = []
valid_accs = []
for epoch in range(args.epochs):
    # Train the Model
    cnn.train() # Change model to 'train' mode
    losses = []
    for i, (xs, ys) in enumerate(get_batch(train_grey,
train_rgb_cat, args.batch_size)):
        images, labels = get_torch_vars(xs, ys, args.gpu)
        # Forward + Backward + Optimize
        optimizer.zero_grad()
        outputs = cnn(images)

```

```

        loss = compute_loss(
            criterion, outputs, labels,
            batch_size=args.batch_size, num_colours=num_colours
        )
        loss.backward()
        optimizer.step()
        losses.append(loss.data.item())

    # plot training images
    if args.plot:
        _, predicted = torch.max(outputs.data, 1, keepdim=True)
        plot(
            xs,
            ys,
            predicted.cpu().numpy(),
            colours,
            save_dir + "/train_%d.png" % epoch,
            args.visualize,
            args.downsize_input,
        )

    # plot training images
    avg_loss = np.mean(losses)
    train_losses.append(avg_loss)
    time_elapsed = time.time() - start
    print(
        "Epoch [%d/%d], Loss: %.4f, Time (s): %d"
        % (epoch + 1, args.epochs, avg_loss, time_elapsed)
    )

    # Evaluate the model
    cnn.eval() # Change model to 'eval' mode (BN uses moving
    mean/var).
    val_loss, val_acc = run_validation_step(
        cnn,
        criterion,
        test_grey,
        test_rgb_cat,
        args.batch_size,
        colours,
        save_dir + "/test_%d.png" % epoch,
        args.visualize,
        args.downsize_input,
    )

    time_elapsed = time.time() - start
    valid_losses.append(val_loss)
    valid_accs.append(val_acc)
    print(
        "Epoch [%d/%d], Val Loss: %.4f, Val Acc: %.1f%%, Time(s):

```

```

%.2f"
        % (epoch + 1, args.epochs, val_loss, val_acc,
time_elapsed)
    )

    # Plot training curve
    plt.figure()
    plt.plot(train_losses, "ro-", label="Train")
    plt.plot(valid_losses, "go-", label="Validation")
    plt.legend()
    plt.title("Loss")
    plt.xlabel("Epochs")
    plt.savefig(save_dir + "/training_curve.png")

    if args.checkpoint:
        print("Saving model...")
        torch.save(cnn.state_dict(), args.checkpoint)

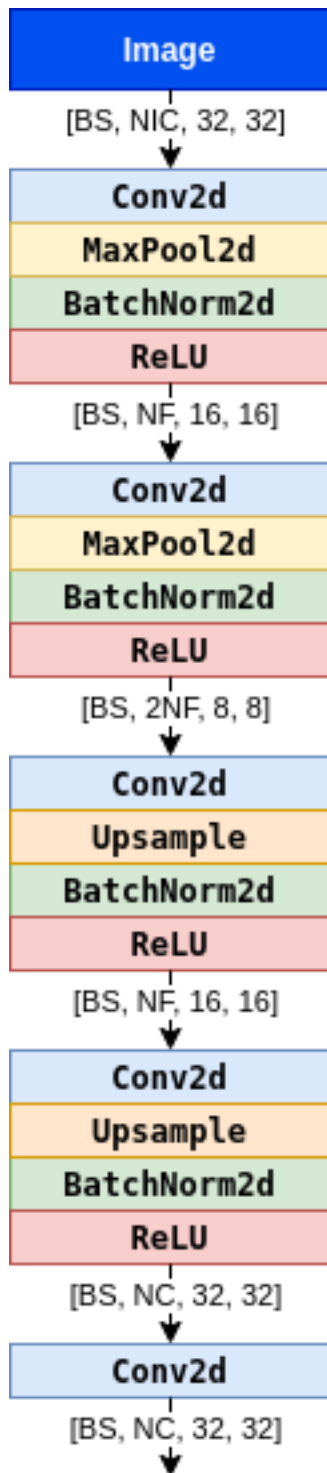
    return cnn

```

Part A: Pooling and Upsampling (2 pts)

Question 1

Complete the PoolUpsampleNet CNN model following the architecture described in the assignment handout.



In the diagram above, we denote the number of filters as **NF**. Further layers double the number of filters, denoted as **2NF**. In the final layers, the number of filters will be equivalent to the number of colour classes, denoted as **NC**. Consequently, your constructed neural network should define the number of input/output layers with respect to the variables `num_filters` and `num_colours`, as opposed to a constant value.

The specific modules to use are listed below. If parameters are not otherwise specified, use the default PyTorch parameters.

- `nn.Conv2d` — The number of input filters should match the second dimension of the *input* tensor (e.g. the first `nn.Conv2d` layer has **NIC** input filters). The number of output filters should match the second dimension of the *output* tensor (e.g. the first `nn.Conv2d` layer has **NF** output filters). Set kernel size to parameter `kernel`. Set padding to the `padding` variable included in the starter code.
- `nn.MaxPool2d` — Use `kernel_size=2` for all layers.
- `nn.BatchNorm2d` — The number of features is specified after the hyphen in the diagram as a multiple of **NF** or **NC**.
- `nn.Upsample` — Use `scaling_factor=2` for all layers.
- `nn.ReLU`

We recommend grouping each block of operations (those adjacent without whitespace in the diagram) into `nn.Sequential` containers. Grouping up relevant operations will allow for easier implementation of the forward method.

```
class PoolUpsampleNet(nn.Module):
    def __init__(self, kernel, num_filters, num_colours,
num_in_channels):
        super().__init__()

        # Useful parameters
        padding = kernel // 2

        ##### YOUR CODE GOES HERE #####
        #####

        self.layer1 = torch.nn.Sequential(
            torch.nn.Conv2d(num_in_channels, num_filters,
kernel_size=kernel, padding = padding),
            torch.nn.MaxPool2d(kernel_size = 2),
            torch.nn.BatchNorm2d(num_filters),
            torch.nn.ReLU()
        )

        self.layer2 = torch.nn.Sequential(
            torch.nn.Conv2d(num_filters, 2 * num_filters,
kernel_size=kernel, padding = padding),
            torch.nn.MaxPool2d(kernel_size = 2),
            torch.nn.BatchNorm2d(2 * num_filters),
            torch.nn.ReLU()
        )

        self.layer3 = torch.nn.Sequential(
            torch.nn.Conv2d(2 * num_filters, num_filters,
kernel_size=kernel, padding = padding),
            torch.nn.Upsample(scale_factor = 2),
```

```

        torch.nn.BatchNorm2d(num_filters),
        torch.nn.ReLU()
    )

    self.layer4 = torch.nn.Sequential(
        torch.nn.Conv2d(num_filters, num_colours,
kernel_size=kernel, padding = padding),
        torch.nn.Upsample(scale_factor = 2),
        torch.nn.BatchNorm2d(num_colours),
        torch.nn.ReLU()
    )

    self.layer5 = torch.nn.Conv2d(num_colours, num_colours,
kernel_size=kernel, padding = padding)

def forward(self, x):
    ##### YOUR CODE GOES HERE #####
    #####
    x = self.layer1(x)
    x = self.layer2(x)
    x = self.layer3(x)
    x = self.layer4(x)
    x = self.layer5(x)
    return x

```

Question 2

Run main training loop of PoolUpsampleNet. This will train the CNN for a few epochs using the cross-entropy objective. It will generate some images showing the trained result at the end. Do these results look good to you? Why or why not?

```

args = AttrDict()
args_dict = {
    "gpu": True,
    "valid": False,
    "checkpoint": "",
    "colours": "./data/colours/colour_kmeans24_cat7.npy",
    "model": "PoolUpsampleNet",
    "kernel": 3,
    "num_filters": 32,
    'learn_rate': 0.001,
    "batch_size": 100,
    "epochs": 25,
    "seed": 0,
    "plot": True,
    "experiment_name": "colourization_cnn",
    "visualize": False,
    "downsize_input": False,
}

```

```
args.update(args_dict)
cnn = train(args)
```

Loading data...


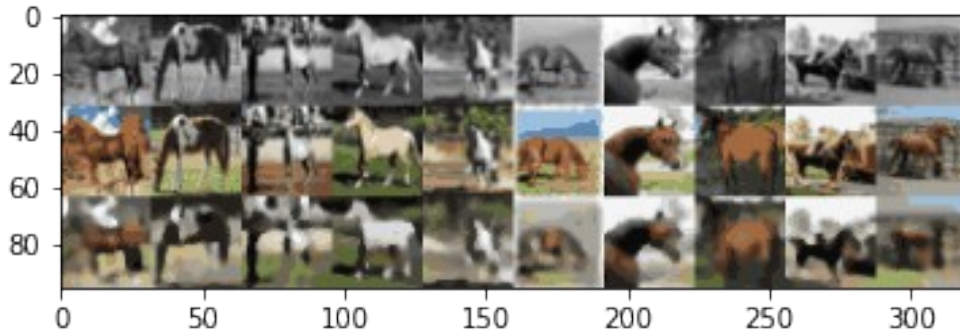
File path: data/cifar-10-batches-py.tar.gz

Transforming data...

Beginning training ...

```
Epoch [1/25], Loss: 2.4397, Time (s): 2
Epoch [1/25], Val Loss: 2.1542, Val Acc: 26.3%, Time(s): 2.50
Epoch [2/25], Loss: 1.9967, Time (s): 4
Epoch [2/25], Val Loss: 1.9079, Val Acc: 33.0%, Time(s): 4.93
Epoch [3/25], Loss: 1.8744, Time (s): 7
Epoch [3/25], Val Loss: 1.8081, Val Acc: 35.8%, Time(s): 7.40
Epoch [4/25], Loss: 1.8130, Time (s): 9
Epoch [4/25], Val Loss: 1.7627, Val Acc: 36.9%, Time(s): 9.92
Epoch [5/25], Loss: 1.7743, Time (s): 12
Epoch [5/25], Val Loss: 1.7331, Val Acc: 37.6%, Time(s): 12.48
Epoch [6/25], Loss: 1.7465, Time (s): 14
Epoch [6/25], Val Loss: 1.7094, Val Acc: 38.2%, Time(s): 15.08
Epoch [7/25], Loss: 1.7253, Time (s): 17
Epoch [7/25], Val Loss: 1.6919, Val Acc: 38.7%, Time(s): 17.72
Epoch [8/25], Loss: 1.7084, Time (s): 19
Epoch [8/25], Val Loss: 1.6786, Val Acc: 39.0%, Time(s): 20.42
Epoch [9/25], Loss: 1.6947, Time (s): 22
Epoch [9/25], Val Loss: 1.6678, Val Acc: 39.2%, Time(s): 23.17
Epoch [10/25], Loss: 1.6831, Time (s): 25
Epoch [10/25], Val Loss: 1.6588, Val Acc: 39.4%, Time(s): 25.94
Epoch [11/25], Loss: 1.6733, Time (s): 28
Epoch [11/25], Val Loss: 1.6504, Val Acc: 39.6%, Time(s): 28.88
Epoch [12/25], Loss: 1.6647, Time (s): 31
Epoch [12/25], Val Loss: 1.6439, Val Acc: 39.8%, Time(s): 32.13
Epoch [13/25], Loss: 1.6573, Time (s): 34
Epoch [13/25], Val Loss: 1.6380, Val Acc: 40.0%, Time(s): 35.06
Epoch [14/25], Loss: 1.6506, Time (s): 37
Epoch [14/25], Val Loss: 1.6325, Val Acc: 40.1%, Time(s): 38.03
Epoch [15/25], Loss: 1.6441, Time (s): 40
Epoch [15/25], Val Loss: 1.6254, Val Acc: 40.2%, Time(s): 41.04
Epoch [16/25], Loss: 1.6379, Time (s): 43
Epoch [16/25], Val Loss: 1.6209, Val Acc: 40.4%, Time(s): 44.11
Epoch [17/25], Loss: 1.6318, Time (s): 46
Epoch [17/25], Val Loss: 1.6153, Val Acc: 40.5%, Time(s): 47.23
Epoch [18/25], Loss: 1.6260, Time (s): 49
Epoch [18/25], Val Loss: 1.6099, Val Acc: 40.6%, Time(s): 50.36
Epoch [19/25], Loss: 1.6207, Time (s): 52
Epoch [19/25], Val Loss: 1.6052, Val Acc: 40.8%, Time(s): 53.55
Epoch [20/25], Loss: 1.6156, Time (s): 56
Epoch [20/25], Val Loss: 1.6016, Val Acc: 40.9%, Time(s): 56.78
Epoch [21/25], Loss: 1.6108, Time (s): 59
Epoch [21/25], Val Loss: 1.5963, Val Acc: 41.0%, Time(s): 60.36
Epoch [22/25], Loss: 1.6065, Time (s): 62
```

Epoch [22/25], Val Loss: 1.5967, Val Acc: 40.9%, Time(s): 63.67
 Epoch [23/25], Loss: 1.6025, Time (s): 66
 Epoch [23/25], Val Loss: 1.5912, Val Acc: 41.1%, Time(s): 67.01
 Epoch [24/25], Loss: 1.5988, Time (s): 69
 Epoch [24/25], Val Loss: 1.5898, Val Acc: 41.1%, Time(s): 70.59
 Epoch [25/25], Loss: 1.5952, Time (s): 73
 Epoch [25/25], Val Loss: 1.5852, Val Acc: 41.2%, Time(s): 74.29



Question 3

Original weight input dimension (width/height)

Number of Weights in 5 convolution layers in their respective order is

$$k^2 * [NIC * NF + 2 * NF^2 + 2 * NF^2 + NF * NC + NC^2]$$

$$= k^2 * [NIC * NF + 4NF^2 + NF * NC + NC^2]$$

Number of weights in max pooling layers is 0 as there are no trainable parameters for retrieving maximum. Number of weights in upsampling layers is also 0 as it is used to double the dimension.

Therefore, the total number of weights in the model is

$$k^2 * [NIC * NF + 4NF^2 + NF * NC + NC^2]$$

Number of Outputs

The outputs from 4 maxpooling/upsampling layers are $16^2 NF$, $2(8^2) NF$, $16^2 NF$, and $32^2 NC$ respectively.

The outputs from the convolution layers are $32^2 NF$, $(2)16^2 NF$, $2(8^2) NF$, $16^2 NC$ and $32^2 NC$

Number of Connections

Number of connections between the first convolution layer and the first maxpooling layer is $(32k)^2 NIC * NF$. Number of connections between the first maxpooling layer and ReLu is $(16)^2 * NF$.

Number of connections between the second convolution layer and the second maxpooling layer is $(16k)^2 (2)NF^2$. Number of connections between the second maxpooling layer and ReLu is $(8)^2 * (2)NF$.

Number of connections between the third convolution layer and the first upsampling layer is $(8k)^2 (2)NF^2$. Number of connections between the first upsampling layer and ReLu is $(16)^2 * NF$.

Number of connections between the fourth convolution layer and the second upsampling layer is $(16k)^2 NF * NC$. Number of connections between the first upsampling layer and ReLu is $(32)^2 * NC$.

Number of connections between the last ReLu and the fifth convolution layer is $(32k * NC)^2$

Each weight input dimension is doubled

Number of Weights remains the same as the original weight input as the number of weights for convolutional layers does not depend on width or height. Therefore, the total number of weights in the model is still $k^2 * [NIC * NF + 4NF^2 + NF * NC + NC^2]$

In terms of **Number of Outputs**, as both width and height are doubled, the number of outputs will be 4 times the number of output of the original dimension. Therefore, the outputs from 4 maxpooling/upsampling layers are $(4)16^2 NF$, $(8^3)NF$, $(4)16^2 NF$, and $(4)32^2 NC$ respectively.

The outputs from the convolution layers are $(4)32^2 NF$, $(8)16^2 NF$, $(8^3)NF$, $(4)16^2 NC$ and $(4)32^2 NC$

Similarly, the **Number of Connections** will be 4 times the original dimension.

Thus, number of connections between the first convolution layer and the first maxpooling layer is $(4)(32k)^2 N I C * N F$. Number of connections between the first maxpooling layer and ReLu is $(4)(16)^2 * N F$.

Number of connections between the second convolution layer and the second maxpooling layer is $(4)(16k)^2 (2) N F^2$. Number of connections between the second maxpooling layer and ReLu is $(8)^3 N F$.

Number of connections between the third convolution layer and the first upsampling layer is $(4)(8k)^2 (2) N F^2$. Number of connections between the first upsampling layer and ReLu is $(4)(16)^2 * N F$.

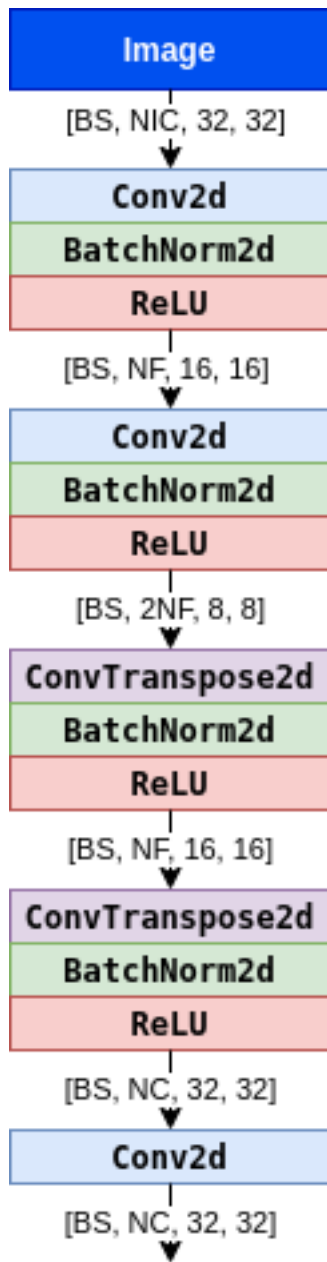
Number of connections between the fourth convolution layer and the second upsampling layer is $(4)(16k)^2 N F * N C$. Number of connections between the first upsampling layer and ReLu is $(4)(32)^2 * N C$.

Number of connections between the last ReLu and the fifth convolution layer is $(4)(32k * N C)^2$

Part B: Strided and Transposed Dilated Convolutions (3 pts)

Question 1

Complete the ConvTransposeNet CNN model following the architecture described in the assignment handout.



An excellent visualization of convolutions and transposed convolutions with strides can be found here: https://github.com/vdumoulin/conv_arithmetic.

The specific modules to use are listed below. If parameters are not otherwise specified, use the default PyTorch parameters.

- `nn.Conv2d` — The number of input and output filters, and the kernel size, should be set in the same way as Part A. For the first two `nn.Conv2d` layers, set `stride` to 2 and set `padding` to 1.
- `nn.BatchNorm2d` — The number of features should be specified in the same way as for Part A.

- `nn.ConvTranspose2d` — The number of input filters should match the second dimension of the *input* tensor. The number of output filters should match the second dimension of the *output* tensor. Set `kernel_size` to parameter `kernel`. Set `stride` to 2, set `dilation` to 1, and set both `padding` and `output_padding` to 1.
- `nn.ReLU`

```
class ConvTransposeNet(nn.Module):
    def __init__(self, kernel, num_filters, num_colours,
num_in_channels):
        super().__init__()

        # Useful parameters
        stride = 2
        padding = kernel // 2
        output_padding = 1

        ##### YOUR CODE GOES HERE #####
        #####

        self.layer1 = torch.nn.Sequential(
            torch.nn.Conv2d(num_in_channels, num_filters,
kernel_size=kernel, padding = 1, stride = 2),
            torch.nn.BatchNorm2d(num_filters),
            torch.nn.ReLU()
        )

        self.layer2 = torch.nn.Sequential(
            torch.nn.Conv2d(num_filters, 2 * num_filters,
kernel_size=kernel, padding = 1, stride = 2),
            torch.nn.BatchNorm2d(2 * num_filters),
            torch.nn.ReLU()
        )

        self.layer3 = torch.nn.Sequential(
            torch.nn.ConvTranspose2d(2 * num_filters, num_filters,
kernel_size = kernel, stride = 2, dilation = 1, padding = 1,
output_padding = 1),
            torch.nn.BatchNorm2d(num_filters),
            torch.nn.ReLU()
        )

        self.layer4 = torch.nn.Sequential(
            torch.nn.ConvTranspose2d(num_filters, num_colours,
kernel_size = kernel, stride = 2, dilation = 1, padding = 1,
output_padding = 1),
            torch.nn.BatchNorm2d(num_colours),
            torch.nn.ReLU()
        )

        self.layer5 = torch.nn.Conv2d(num_colours, num_colours,
```



```
kernel_size = kernel, padding = padding)
```

```
def forward(self, x):  
    ##### YOUR CODE GOES HERE #####  
    #####  
    x = self.layer1(x)  
    x = self.layer2(x)  
    x = self.layer3(x)  
    x = self.layer4(x)  
    x = self.layer5(x)  
    return x
```

Question 2

Train the model for at least 25 epochs using a batch size of 100 and a kernel size of 3. Plot the training curve, and include this plot in your write-up. How do the results compare to the previous model?

```
args = AttrDict()  
args_dict = {  
    "gpu": True,  
    "valid": False,  
    "checkpoint": "",  
    "colours": "./data/colours/colour_kmeans24_cat7.npy",  
    "model": "ConvTransposeNet",  
    "kernel": 3,  
    "num_filters": 32,  
    'learn_rate': 0.001,  
    "batch_size": 100,  
    "epochs": 25,  
    "seed": 0,  
    "plot": True,  
    "experiment_name": "colourization_cnn",  
    "visualize": False,  
    "downsize_input": False,  
}  
args.update(args_dict)  
cnn = train(args)
```

Loading data...

File path: data/cifar-10-batches-py.tar.gz

Transforming data...

Beginning training ...

Epoch [1/25], Loss: 2.4820, Time (s): 2

Epoch [1/25], Val Loss: 2.0605, Val Acc: 30.9%, Time(s): 2.32

Epoch [2/25], Loss: 1.8601, Time (s): 4


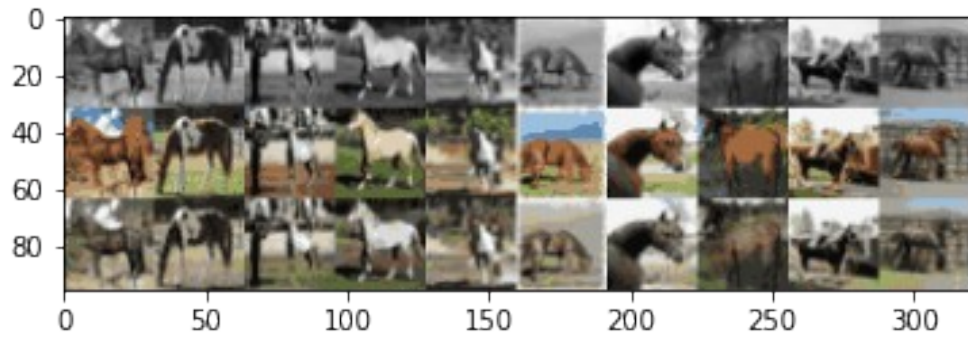
Epoch [2/25], Val Loss: 1.7481, Val Acc: 37.6%, Time(s): 4.51

Epoch [3/25], Loss: 1.7111, Time (s): 6

Epoch [3/25], Val Loss: 1.6406, Val Acc: 40.3%, Time(s): 6.77

Epoch [4/25], Loss: 1.6305, Time (s): 8

Epoch [4/25], Val Loss: 1.5723, Val Acc: 42.0%, Time(s): 9.06
Epoch [5/25], Loss: 1.5703, Time (s): 11
Epoch [5/25], Val Loss: 1.5130, Val Acc: 43.9%, Time(s): 11.41
Epoch [6/25], Loss: 1.5221, Time (s): 13
Epoch [6/25], Val Loss: 1.4685, Val Acc: 45.2%, Time(s): 13.78
Epoch [7/25], Loss: 1.4828, Time (s): 15
Epoch [7/25], Val Loss: 1.4350, Val Acc: 46.1%, Time(s): 16.20
Epoch [8/25], Loss: 1.4506, Time (s): 18
Epoch [8/25], Val Loss: 1.4063, Val Acc: 46.9%, Time(s): 18.68
Epoch [9/25], Loss: 1.4230, Time (s): 20
Epoch [9/25], Val Loss: 1.3814, Val Acc: 47.6%, Time(s): 21.20
Epoch [10/25], Loss: 1.3986, Time (s): 23
Epoch [10/25], Val Loss: 1.3593, Val Acc: 48.3%, Time(s): 23.78
Epoch [11/25], Loss: 1.3771, Time (s): 25
Epoch [11/25], Val Loss: 1.3340, Val Acc: 49.2%, Time(s): 26.40
Epoch [12/25], Loss: 1.3575, Time (s): 28
Epoch [12/25], Val Loss: 1.3128, Val Acc: 49.9%, Time(s): 29.06
Epoch [13/25], Loss: 1.3398, Time (s): 31
Epoch [13/25], Val Loss: 1.2947, Val Acc: 50.5%, Time(s): 31.77
Epoch [14/25], Loss: 1.3236, Time (s): 33
Epoch [14/25], Val Loss: 1.2778, Val Acc: 51.0%, Time(s): 34.50
Epoch [15/25], Loss: 1.3089, Time (s): 36
Epoch [15/25], Val Loss: 1.2631, Val Acc: 51.4%, Time(s): 37.28
Epoch [16/25], Loss: 1.2954, Time (s): 39
Epoch [16/25], Val Loss: 1.2527, Val Acc: 51.8%, Time(s): 40.11
Epoch [17/25], Loss: 1.2829, Time (s): 42
Epoch [17/25], Val Loss: 1.2416, Val Acc: 52.1%, Time(s): 42.94
Epoch [18/25], Loss: 1.2714, Time (s): 45
Epoch [18/25], Val Loss: 1.2262, Val Acc: 52.6%, Time(s): 45.82
Epoch [19/25], Loss: 1.2607, Time (s): 48
Epoch [19/25], Val Loss: 1.2159, Val Acc: 52.9%, Time(s): 48.76
Epoch [20/25], Loss: 1.2507, Time (s): 51
Epoch [20/25], Val Loss: 1.2053, Val Acc: 53.2%, Time(s): 51.73
Epoch [21/25], Loss: 1.2415, Time (s): 54
Epoch [21/25], Val Loss: 1.1999, Val Acc: 53.4%, Time(s): 54.76
Epoch [22/25], Loss: 1.2328, Time (s): 57
Epoch [22/25], Val Loss: 1.1906, Val Acc: 53.7%, Time(s): 57.83
Epoch [23/25], Loss: 1.2247, Time (s): 60
Epoch [23/25], Val Loss: 1.1812, Val Acc: 54.0%, Time(s): 60.96
Epoch [24/25], Loss: 1.2172, Time (s): 63
Epoch [24/25], Val Loss: 1.1754, Val Acc: 54.2%, Time(s): 64.15
Epoch [25/25], Loss: 1.2101, Time (s): 66
Epoch [25/25], Val Loss: 1.1684, Val Acc: 54.4%, Time(s): 67.42



Question 3

ConvTransposeNet has lower validation loss, higher validation accuracy and shorter computational time than PoolUpsampleNet. It has a higher accuracy as it is built up from a low resolution image to an image with higher resolution. Therefore, after training, the network has the ability to identify images with rough resolution and modifies them to have higher resolution for prediction and/or classification. These steps allow the model to make better predictions, which results in lower loss.

Question 4

The padding parameter for the convolution layers will increase. As we increase the kernel size, the output layer's receptive field is larger. Therefore, the output layer would have a smaller size. Therefore, if we want to maintain the output size, padding needs to be increased.

Question 5

When batch size is 100, the validation loss is 1.1566, the validation accuracy is 54.9% and it takes about 70 seconds to train.


When batch size is 32, validation loss decreases and validation accuracy rises to 56.9% but it takes 5 more seconds for the program to execute. With a smaller batch size, the predicted pixel can closely resemble the color of the original image. This means the predicted image is very similar to the original RGB image.

On the other hand, if batch size is 128, the validation loss is higher and the validation accuracy drops to 53.4%. However, the training finished in 66 seconds. With larger batch size, the predicted pixel is more similar to the black and white image and the image is less colourful.

Part C. Skip Connections (1 pts)

A skip connection in a neural network is a connection which skips one or more layer and connects to a later layer. We will introduce skip connections to our previous model.

Question 1



In this question, we will be adding a skip connection from the first layer to the last, second layer to the second last, etc. That is, the final convolution should have both the output of the previous layer and the initial greyscale input as input. This type of skip-connection is introduced by [Ronneberger et al.\[2015\]](#), and is called a "UNet".

Just like the `ConvTransposeNet` class that you have completed in the previous part, complete the `__init__` and forward methods of the `UNet` class below.

Hint: You will need to use the function `torch.cat`.

```

class UNet(nn.Module):
    def __init__(self, kernel, num_filters, num_colours,
num_in_channels):
        super().__init__()

        # Useful parameters
        stride = 2
        padding = kernel // 2
        output_padding = 1

        ##### YOUR CODE GOES HERE #####
        #####

        self.layer1 = torch.nn.Sequential(
            torch.nn.Conv2d(num_in_channels, num_filters,
kernel_size=kernel, padding = 1, stride = 2),
            torch.nn.BatchNorm2d(num_filters),
            torch.nn.ReLU()
        )

        self.layer2 = torch.nn.Sequential(
            torch.nn.Conv2d(num_filters, 2 * num_filters,
kernel_size=kernel, padding = 1, stride = 2),
            torch.nn.BatchNorm2d(2 * num_filters),
            torch.nn.ReLU()
        )

        self.layer3 = torch.nn.Sequential(
            torch.nn.ConvTranspose2d(2 * num_filters, num_filters,
kernel_size = kernel, stride = 2, dilation = 1, padding = 1,
output_padding = 1),
            torch.nn.BatchNorm2d(num_filters),
            torch.nn.ReLU()
        )

        # the input dim. is 2 * num_filters because layer 3_prime has
2 * num_filters input units
        self.layer4 = torch.nn.Sequential(
            torch.nn.ConvTranspose2d(2 * num_filters, num_colours,
kernel_size = kernel, stride = stride, dilation = 1, padding = 1,
output_padding = 1),
            torch.nn.BatchNorm2d(num_colours),
            torch.nn.ReLU()
        )

        self.layer5 = torch.nn.Conv2d(num_colours, num_colours,
kernel_size = kernel, padding = padding)

    def forward(self, x):
        ##### YOUR CODE GOES HERE #####

```

```
#####
x1 = self.layer1(x)
x2 = self.layer2(x1)
x3 = self.layer3(x2)
x3_prime = torch.cat((x1, x3), 1)
x4 = self.layer4(x3_prime)
x4_prime = torch.cat((x, x4), 1)
x5 = self.layer5(x4)
return x5
```

Question 2

Train the model for at least 25 epochs using a batch size of 100 and a kernel size of 3. Plot the training curve, and include this plot in your write-up.

```
args = AttrDict()
args_dict = {
    "gpu": True,
    "valid": False,
    "checkpoint": "",
    "colours": "./data/colours/colour_kmeans24_cat7.npy",
    "model": "UNet",
    "kernel": 3,
    "num_filters": 32,
    'learn_rate': 0.001,
    "batch_size": 100,
    "epochs": 25,
    "seed": 0,
    "plot": True,
    "experiment_name": "colourization_cnn",
    "visualize": False,
    "downsize_input": False,
}
args.update(args_dict)
cnn = train(args)
```

Loading data...


File path: data/cifar-10-batches-py.tar.gz


Transforming data...

Beginning training ...

```
Epoch [1/25], Loss: 2.4398, Time (s): 2
Epoch [1/25], Val Loss: 2.0491, Val Acc: 32.6%, Time(s): 2.43
Epoch [2/25], Loss: 1.8276, Time (s): 4
Epoch [2/25], Val Loss: 1.7302, Val Acc: 37.4%, Time(s): 4.75
Epoch [3/25], Loss: 1.6447, Time (s): 6
Epoch [3/25], Val Loss: 1.5793, Val Acc: 42.2%, Time(s): 7.13
Epoch [4/25], Loss: 1.5504, Time (s): 9
Epoch [4/25], Val Loss: 1.4961, Val Acc: 44.6%, Time(s): 9.57
Epoch [5/25], Loss: 1.4825, Time (s): 11
Epoch [5/25], Val Loss: 1.4349, Val Acc: 46.5%, Time(s): 12.03
Epoch [6/25], Loss: 1.4283, Time (s): 14
```

Epoch [6/25], Val Loss: 1.3872, Val Acc: 47.8%, Time(s): 14.52
 Epoch [7/25], Loss: 1.3839, Time (s): 16
 Epoch [7/25], Val Loss: 1.3441, Val Acc: 49.2%, Time(s): 17.07
 Epoch [8/25], Loss: 1.3466, Time (s): 19
 Epoch [8/25], Val Loss: 1.3065, Val Acc: 50.5%, Time(s): 19.66
 Epoch [9/25], Loss: 1.3146, Time (s): 21
 Epoch [9/25], Val Loss: 1.2743, Val Acc: 51.5%, Time(s): 22.30
 Epoch [10/25], Loss: 1.2871, Time (s): 24
 Epoch [10/25], Val Loss: 1.2469, Val Acc: 52.5%, Time(s): 25.29
 Epoch [11/25], Loss: 1.2635, Time (s): 27
 Epoch [11/25], Val Loss: 1.2219, Val Acc: 53.3%, Time(s): 27.99
 Epoch [12/25], Loss: 1.2431, Time (s): 30
 Epoch [12/25], Val Loss: 1.2013, Val Acc: 54.0%, Time(s): 30.76
 Epoch [13/25], Loss: 1.2255, Time (s): 33
 Epoch [13/25], Val Loss: 1.1821, Val Acc: 54.6%, Time(s): 33.58
 Epoch [14/25], Loss: 1.2098, Time (s): 35
 Epoch [14/25], Val Loss: 1.1649, Val Acc: 55.2%, Time(s): 36.45
 Epoch [15/25], Loss: 1.1962, Time (s): 38
 Epoch [15/25], Val Loss: 1.1514, Val Acc: 55.5%, Time(s): 39.35
 Epoch [16/25], Loss: 1.1838, Time (s): 41
 Epoch [16/25], Val Loss: 1.1365, Val Acc: 56.0%, Time(s): 42.28
 Epoch [17/25], Loss: 1.1730, Time (s): 44
 Epoch [17/25], Val Loss: 1.1282, Val Acc: 56.2%, Time(s): 45.27
 Epoch [18/25], Loss: 1.1637, Time (s): 47
 Epoch [18/25], Val Loss: 1.1232, Val Acc: 56.3%, Time(s): 48.30
 Epoch [19/25], Loss: 1.1544, Time (s): 50
 Epoch [19/25], Val Loss: 1.1140, Val Acc: 56.6%, Time(s): 51.38
 Epoch [20/25], Loss: 1.1458, Time (s): 53
 Epoch [20/25], Val Loss: 1.1081, Val Acc: 56.8%, Time(s): 54.49
 Epoch [21/25], Loss: 1.1379, Time (s): 56
 Epoch [21/25], Val Loss: 1.0965, Val Acc: 57.2%, Time(s): 57.62
 Epoch [22/25], Loss: 1.1305, Time (s): 60
 Epoch [22/25], Val Loss: 1.0877, Val Acc: 57.5%, Time(s): 61.32
 Epoch [23/25], Loss: 1.1234, Time (s): 63
 Epoch [23/25], Val Loss: 1.0795, Val Acc: 57.8%, Time(s): 64.55
 Epoch [24/25], Loss: 1.1170, Time (s): 67
 Epoch [24/25], Val Loss: 1.0770, Val Acc: 57.8%, Time(s): 67.83
 Epoch [25/25], Loss: 1.1110, Time (s): 70
 Epoch [25/25], Val Loss: 1.0689, Val Acc: 58.0%, Time(s): 71.19





Question 3

The skip connection model performs better than the previous 2 models. The validation loss is lower, the validation accuracy is higher but the training time is about the same as the other 2 models. The predictive image is also more colourful than the previous 2 networks.

The skip connection model merges information from previous layers with the current layer. This recovers some information lost during downsampling. This ensures the flow between each input and output layer is maintained at a maximum level. On the other hand, reusing previous features also stabilize training and convergence.

Object Detection as Regression and Classification - the YOLO approach

In the previous two parts, we worked on training models for image colourization. Now we will switch gears and perform object detection by fine-tuning a pre-trained model.

For the following, you are not expected to read the referenced papers, though the writing is very entertaining (by academic paper standards) and it may help provide additional context.

We use the YOLO (You Only Look Once) approach, as laid out in the [the original paper by Redmon et al.](#) YOLO uses a single neural network to predict bounding boxes (4 coordinates describing the corners of the box bounding a particular object) and class probabilities (what object is in the bounding box) based on a single pass over an image. It first divides the image into a grid, and for each grid cell predicts bounding boxes, confidence for those boxes, and conditional class probabilities.

For the YOLOv3 model, which we use here, we draw from [their YOLOv3 paper](#) which also builds on the previous [YOLO9000 paper](#).

We use the pretrained YOLOv3 model weights and fine-tune it on the COCO ([Lin et al., 2014](#)) dataset.

##Setup

```
!git clone https://github.com/Silent-Zebra/2022
```

fatal: destination path '2022' already exists and is not an empty directory.

Rerun the cd command below if you restart the runtime (but everything should work fine without restarting the runtime anyway)

```
%cd 2022/assets/assignments/pa2-q4-files
```

```
[Errno 2] No such file or directory: '2022/assets/assignments/pa2-q4-files'
/content/csc413/a2/2022/assets/assignments/pa2-q4-files
```

```
def notebook_init():
    # For notebooks
    print('Checking setup...')
    from IPython import display # to display images and clear console output

    from utils.general import emojis
    from utils.torch_utils import select_device # imports

    display.clear_output()
    select_device(newline=False)
    print(emojis('Setup complete ✔'))
    return display
```

```
display = notebook_init()
```

```
YOLOv3 46dad08 torch 1.10.0+cu111 CUDA:0 (Tesla K80, 11441MiB)
```

```
Setup complete ✔
```

In my experience you don't have to restart the runtime after the below installation

```
!pip install -r requirements.txt
```

```
Requirement already satisfied: matplotlib>=3.2.2 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
4)) (3.2.2)
```

```
Requirement already satisfied: numpy>=1.18.5 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
5)) (1.21.5)
```

Requirement already satisfied: opencv-python>=4.1.2 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
6)) (4.1.2.30)

Requirement already satisfied: Pillow>=7.1.2 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
7)) (7.1.2)

Requirement already satisfied: PyYAML>=5.3.1 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
8)) (6.0)

Requirement already satisfied: requests>=2.23.0 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
9)) (2.23.0)

Requirement already satisfied: scipy>=1.4.1 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
10)) (1.4.1)

Requirement already satisfied: torch>=1.7.0 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
11)) (1.10.0+cu111)

Requirement already satisfied: torchvision>=0.8.1 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
12)) (0.11.1+cu111)

Requirement already satisfied: tqdm>=4.41.0 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
13)) (4.62.3)

Requirement already satisfied: tensorboard>=2.4.1 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
16)) (2.8.0)

Requirement already satisfied: wandb in /usr/local/lib/python3.7/dist-
packages (from -r requirements.txt (line 17)) (0.12.10)

Requirement already satisfied: pandas>=1.1.4 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
20)) (1.3.5)

Requirement already satisfied: seaborn>=0.11.0 in
/usr/local/lib/python3.7/dist-packages (from -r requirements.txt (line
21)) (0.11.2)

Requirement already satisfied: thop in /usr/local/lib/python3.7/dist-
packages (from -r requirements.txt (line 36)) (0.0.31.post2005241907)

Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=3.2.2->-r
requirements.txt (line 4)) (0.11.0)

Requirement already satisfied: python-dateutil>=2.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=3.2.2->-r
requirements.txt (line 4)) (2.8.2)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!
=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from
matplotlib>=3.2.2->-r requirements.txt (line 4)) (3.0.7)

Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=3.2.2->-r
requirements.txt (line 4)) (1.3.2)

Requirement already satisfied: urllib3!=1.25.0,!=1.25.1,<1.26,>=1.21.1

in /usr/local/lib/python3.7/dist-packages (from requests>=2.23.0->-r requirements.txt (line 9)) (1.24.3)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests>=2.23.0->-r requirements.txt (line 9)) (2.10)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests>=2.23.0->-r requirements.txt (line 9)) (2021.10.8)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests>=2.23.0->-r requirements.txt (line 9)) (3.0.4)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch>=1.7.0->-r requirements.txt (line 11)) (3.10.0.2)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (0.4.6)
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (57.4.0)
Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (1.8.1)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (3.3.6)
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (0.6.1)
Requirement already satisfied: protobuf>=3.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (3.17.3)
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (1.0.1)
Requirement already satisfied: wheel>=0.26 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (0.37.1)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (1.35.0)
Requirement already satisfied: grpcio>=1.24.3 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (1.43.0)
Requirement already satisfied: absl-py>=0.4 in /usr/local/lib/python3.7/dist-packages (from tensorboard>=2.4.1->-r requirements.txt (line 16)) (1.0.0)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-packages (from pandas>=1.1.4->-r requirements.txt (line 20)) (2018.9)

Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from absl-py>=0.4->tensorboard>=2.4.1->-r requirements.txt (line 16)) (1.15.0)

Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard>=2.4.1->-r requirements.txt (line 16)) (0.2.8)

Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard>=2.4.1->-r requirements.txt (line 16)) (4.8)

Requirement already satisfied: cachetools<5.0,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard>=2.4.1->-r requirements.txt (line 16)) (4.2.4)

Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.7/dist-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorboard>=2.4.1->-r requirements.txt (line 16)) (1.3.1)

Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.7/dist-packages (from markdown>=2.6.8->tensorboard>=2.4.1->-r requirements.txt (line 16)) (4.11.0)

Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata>=4.4->markdown>=2.6.8->tensorboard>=2.4.1->-r requirements.txt (line 16)) (3.7.0)

Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python3.7/dist-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorboard>=2.4.1->-r requirements.txt (line 16)) (0.4.8)

Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.7/dist-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorboard>=2.4.1->-r requirements.txt (line 16)) (3.2.0)

Requirement already satisfied: Click!=8.0.0,>=7.0 in /usr/local/lib/python3.7/dist-packages (from wandb->-r requirements.txt (line 17)) (7.1.2)

Requirement already satisfied: pathtools in /usr/local/lib/python3.7/dist-packages (from wandb->-r requirements.txt (line 17)) (0.1.2)

Requirement already satisfied: yaspin>=1.0.0 in /usr/local/lib/python3.7/dist-packages (from wandb->-r requirements.txt (line 17)) (2.1.0)

Requirement already satisfied: GitPython>=1.0.0 in /usr/local/lib/python3.7/dist-packages (from wandb->-r requirements.txt (line 17)) (3.1.26)

Requirement already satisfied: psutil>=5.0.0 in /usr/local/lib/python3.7/dist-packages (from wandb->-r requirements.txt (line 17)) (5.4.8)

Requirement already satisfied: sentry-sdk>=1.0.0 in /usr/local/lib/python3.7/dist-packages (from wandb->-r requirements.txt (line 17)) (1.5.5)

Requirement already satisfied: promise<3,>=2.0 in

```

/usr/local/lib/python3.7/dist-packages (from wandb->-r
requirements.txt (line 17)) (2.3)
Requirement already satisfied: docker-pycreds>=0.4.0 in
/usr/local/lib/python3.7/dist-packages (from wandb->-r
requirements.txt (line 17)) (0.4.0)
Requirement already satisfied: shortuuid>=0.5.0 in
/usr/local/lib/python3.7/dist-packages (from wandb->-r
requirements.txt (line 17)) (1.0.8)
Requirement already satisfied: gitdb<5,>=4.0.1 in
/usr/local/lib/python3.7/dist-packages (from GitPython>=1.0.0->wandb-
->-r requirements.txt (line 17)) (4.0.9)
Requirement already satisfied: smmap<6,>=3.0.1 in
/usr/local/lib/python3.7/dist-packages (from gitdb<5,>=4.0.1-
>GitPython>=1.0.0->wandb->-r requirements.txt (line 17)) (5.0.0)
Requirement already satisfied: termcolor<2.0.0,>=1.1.0 in
/usr/local/lib/python3.7/dist-packages (from yaspin>=1.0.0->wandb->-r
requirements.txt (line 17)) (1.1.0)

```

Part D.1: Freezing Parameters

A common practice in computer vision tasks is to take a pre-trained model trained on a large dataset and finetune only parts of the model for a specific usecase. This can be helpful, for example, for preventing overfitting if the dataset we fine-tune on is small.

In this notebook, we are finetuning on the COCO dataset, and freezing model parameters is not strictly necessary here. However, it still allows for faster training and is meant to be instructive.

Fill in the section in `train.py` (2022/assets/assignments/pa2-q4-files/train.py) for freezing model parameters (line 129). The key idea here is to set `requires_grad` to be `False` for all frozen layers `v` and `True` for all other layers `v`. Hint: it might be helpful to check a condition such as: `any(x in k for x in freeze)`.

Part D.2: Classification Loss

The YOLO model's loss function consists of several components, including a regression loss for the bounding box as well as a classification loss for the object in the bounding box. We will consider only the classification loss here.

For the classification loss, we will work in `utils/loss.py`.

First define in line 97 the `BCEcls` by calling `nn.BCEWithLogitsLoss` (see PyTorch documentation [here](#) for reference) using `h['cls_pw']` as the positive weight for the classification, and passing in `device` as well. Then, in line 150, add to `lcls` the loss using `self.BCEcls` called on the respective parts of the prediction related to the classification component (`ps[:, 5:]`) and the target `t`.

Training

Train the YOLOv3 model on COCO128 for 5 epochs, freezing 10 layers, by running the below cell.

You can set up an account for wandb (click on the output area at the flashing cursor where it asks you to enter your choice, and type whatever input number you like, followed by hitting enter), which provides lots of cool visualizations (during training you will see live updates at <https://wandb.ai/home>). For the purpose of this assignment, you can enter 3 (skipping wandb).

NOTE: This cell below should take around 6 minutes. If it is taking much longer, please double check your work on Part D.1 (freezing the model parameters)

```
!python train.py --img 640 --batch 16 --epochs 5 --data coco128.yaml
--weights yolov3.pt --cache --freeze 10
```

```
wandb: (1) Create a W&B account
wandb: (2) Use an existing W&B account
wandb: (3) Don't visualize my results
wandb: Enter your choice: (30 second timeout) 3
wandb: You chose 'Don't visualize my results'
train: weights=yolov3.pt, cfg=, data=coco128.yaml,
hyp=data/hyps/hyp.scratch.yaml, epochs=5, batch_size=16, imgsz=640,
rect=False, resume=False, nosave=False, noval=False,
noautoanchor=False, evolve=None, bucket=, cache=ram,
image_weights=False, device=, multi_scale=False, single_cls=False,
adam=False, sync_bn=False, workers=8, project=runs/train, name=exp,
exist_ok=False, quad=False, linear_lr=False, label_smoothing=0.0,
patience=100, freeze=10, save_period=-1, local_rank=-1, entity=None,
upload_dataset=False, bbox_interval=-1, artifact_alias=latest
github: skipping check (not a git repository), for updates see
https://github.com/ultralytics/yolov3
YOLOv3 46dad08 torch 1.10.0+cud11 CUDA:0 (Tesla K80, 11441MiB)
```

```
hyperparameters: lr0=0.01, lrf=0.1, momentum=0.937,
weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8,
warmup_bias_lr=0.1, box=0.05, cls=0.5, cls_pw=1.0, obj=1.0,
obj_pw=1.0, iou_t=0.2, anchor_t=4.0, fl_gamma=0.0, hsv_h=0.015,
hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.1, scale=0.5,
shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, mosaic=1.0,
mixup=0.0, copy_paste=0.0
Weights & Biases: run 'pip install wandb' to automatically track and
visualize YOLOv3 runs (RECOMMENDED)
TensorBoard: Start with 'tensorboard --logdir runs/train', view at
http://localhost:6006/
```

	from	n	params	module
arguments				
0	-1	1	928	models.common.Conv

[3, 32, 3, 1]					
1	-1	1	18560	models.common.Conv	
[32, 64, 3, 2]					
2	-1	1	20672	models.common.Bottleneck	
[64, 64]					
3	-1	1	73984	models.common.Conv	
[64, 128, 3, 2]					
4	-1	2	164608	models.common.Bottleneck	
[128, 128]					
5	-1	1	295424	models.common.Conv	
[128, 256, 3, 2]					
6	-1	8	2627584	models.common.Bottleneck	
[256, 256]					
7	-1	1	1180672	models.common.Conv	
[256, 512, 3, 2]					
8	-1	8	10498048	models.common.Bottleneck	
[512, 512]					
9	-1	1	4720640	models.common.Conv	
[512, 1024, 3, 2]					
10	-1	4	20983808	models.common.Bottleneck	
[1024, 1024]					
11	-1	1	5245952	models.common.Bottleneck	
[1024, 1024, False]					
12	-1	1	525312	models.common.Conv	
[1024, 512, [1, 1]]					
13	-1	1	4720640	models.common.Conv	
[512, 1024, 3, 1]					
14	-1	1	525312	models.common.Conv	
[1024, 512, 1, 1]					
15	-1	1	4720640	models.common.Conv	
[512, 1024, 3, 1]					
16	-2	1	131584	models.common.Conv	
[512, 256, 1, 1]					
17	-1	1	0		
torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']					
18	[-1, 8]	1	0	models.common.Concat	
[1]					
19	-1	1	1377792	models.common.Bottleneck	
[768, 512, False]					
20	-1	1	1312256	models.common.Bottleneck	
[512, 512, False]					
21	-1	1	131584	models.common.Conv	
[512, 256, 1, 1]					
22	-1	1	1180672	models.common.Conv	
[256, 512, 3, 1]					
23	-2	1	33024	models.common.Conv	
[256, 128, 1, 1]					
24	-1	1	0		
torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']					
25	[-1, 6]	1	0	models.common.Concat	


```

[1]
 26          -1  1    344832  models.common.Bottleneck
[384, 256, False]
 27          -1  2    656896  models.common.Bottleneck
[256, 256, False]
 28      [27, 22, 15]  1    457725  models.yolo.Detect
[80, [[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 90,
156, 198, 373, 326]], [256, 512, 1024]]
Model Summary: 333 layers, 61949149 parameters, 61949149 gradients,
156.3 GFLOPs

```

Transferred 439/439 items from yolov3.pt

```

freezing model.0.conv.weight
freezing model.0.bn.weight
freezing model.0.bn.bias
freezing model.1.conv.weight
freezing model.1.bn.weight
freezing model.1.bn.bias
freezing model.2.cv1.conv.weight
freezing model.2.cv1.bn.weight
freezing model.2.cv1.bn.bias
freezing model.2.cv2.conv.weight
freezing model.2.cv2.bn.weight
freezing model.2.cv2.bn.bias
freezing model.3.conv.weight
freezing model.3.bn.weight
freezing model.3.bn.bias
freezing model.4.0.cv1.conv.weight
freezing model.4.0.cv1.bn.weight
freezing model.4.0.cv1.bn.bias
freezing model.4.0.cv2.conv.weight
freezing model.4.0.cv2.bn.weight
freezing model.4.0.cv2.bn.bias
freezing model.4.1.cv1.conv.weight
freezing model.4.1.cv1.bn.weight
freezing model.4.1.cv1.bn.bias
freezing model.4.1.cv2.conv.weight
freezing model.4.1.cv2.bn.weight
freezing model.4.1.cv2.bn.bias
freezing model.5.conv.weight
freezing model.5.bn.weight
freezing model.5.bn.bias
freezing model.6.0.cv1.conv.weight
freezing model.6.0.cv1.bn.weight
freezing model.6.0.cv1.bn.bias
freezing model.6.0.cv2.conv.weight
freezing model.6.0.cv2.bn.weight
freezing model.6.0.cv2.bn.bias
freezing model.6.1.cv1.conv.weight
freezing model.6.1.cv1.bn.weight

```

freezing model.6.1.cv1.bn.bias
freezing model.6.1.cv2.conv.weight
freezing model.6.1.cv2.bn.weight
freezing model.6.1.cv2.bn.bias
freezing model.6.2.cv1.conv.weight
freezing model.6.2.cv1.bn.weight
freezing model.6.2.cv1.bn.bias
freezing model.6.2.cv2.conv.weight
freezing model.6.2.cv2.bn.weight
freezing model.6.2.cv2.bn.bias
freezing model.6.3.cv1.conv.weight
freezing model.6.3.cv1.bn.weight
freezing model.6.3.cv1.bn.bias
freezing model.6.3.cv2.conv.weight
freezing model.6.3.cv2.bn.weight
freezing model.6.3.cv2.bn.bias
freezing model.6.4.cv1.conv.weight
freezing model.6.4.cv1.bn.weight
freezing model.6.4.cv1.bn.bias
freezing model.6.4.cv2.conv.weight
freezing model.6.4.cv2.bn.weight
freezing model.6.4.cv2.bn.bias
freezing model.6.5.cv1.conv.weight
freezing model.6.5.cv1.bn.weight
freezing model.6.5.cv1.bn.bias
freezing model.6.5.cv2.conv.weight
freezing model.6.5.cv2.bn.weight
freezing model.6.5.cv2.bn.bias
freezing model.6.6.cv1.conv.weight
freezing model.6.6.cv1.bn.weight
freezing model.6.6.cv1.bn.bias
freezing model.6.6.cv2.conv.weight
freezing model.6.6.cv2.bn.weight
freezing model.6.6.cv2.bn.bias
freezing model.6.7.cv1.conv.weight
freezing model.6.7.cv1.bn.weight
freezing model.6.7.cv1.bn.bias
freezing model.6.7.cv2.conv.weight
freezing model.6.7.cv2.bn.weight
freezing model.6.7.cv2.bn.bias
freezing model.7.conv.weight
freezing model.7.bn.weight
freezing model.7.bn.bias
freezing model.8.0.cv1.conv.weight
freezing model.8.0.cv1.bn.weight
freezing model.8.0.cv1.bn.bias
freezing model.8.0.cv2.conv.weight
freezing model.8.0.cv2.bn.weight
freezing model.8.0.cv2.bn.bias
freezing model.8.1.cv1.conv.weight

```
freezing model.8.1.cv1.bn.weight
freezing model.8.1.cv1.bn.bias
freezing model.8.1.cv2.conv.weight
freezing model.8.1.cv2.bn.weight
freezing model.8.1.cv2.bn.bias
freezing model.8.2.cv1.conv.weight
freezing model.8.2.cv1.bn.weight
freezing model.8.2.cv1.bn.bias
freezing model.8.2.cv2.conv.weight
freezing model.8.2.cv2.bn.weight
freezing model.8.2.cv2.bn.bias
freezing model.8.3.cv1.conv.weight
freezing model.8.3.cv1.bn.weight
freezing model.8.3.cv1.bn.bias
freezing model.8.3.cv2.conv.weight
freezing model.8.3.cv2.bn.weight
freezing model.8.3.cv2.bn.bias
freezing model.8.4.cv1.conv.weight
freezing model.8.4.cv1.bn.weight
freezing model.8.4.cv1.bn.bias
freezing model.8.4.cv2.conv.weight
freezing model.8.4.cv2.bn.weight
freezing model.8.4.cv2.bn.bias
freezing model.8.5.cv1.conv.weight
freezing model.8.5.cv1.bn.weight
freezing model.8.5.cv1.bn.bias
freezing model.8.5.cv2.conv.weight
freezing model.8.5.cv2.bn.weight
freezing model.8.5.cv2.bn.bias
freezing model.8.6.cv1.conv.weight
freezing model.8.6.cv1.bn.weight
freezing model.8.6.cv1.bn.bias
freezing model.8.6.cv2.conv.weight
freezing model.8.6.cv2.bn.weight
freezing model.8.6.cv2.bn.bias
freezing model.8.7.cv1.conv.weight
freezing model.8.7.cv1.bn.weight
freezing model.8.7.cv1.bn.bias
freezing model.8.7.cv2.conv.weight
freezing model.8.7.cv2.bn.weight
freezing model.8.7.cv2.bn.bias
freezing model.9.conv.weight
freezing model.9.bn.weight
freezing model.9.bn.bias
Scaled weight_decay = 0.0005
optimizer: SGD with parameter groups 72 weight, 75 weight (no decay),
75 bias
albumentations: version 1.0.3 required by YOLOv3, but version 0.1.12
is currently installed
train: Scanning '../datasets/coco128/labels/train2017.cache' images
```

```

and labels... 128 found, 0 missing, 2 empty, 0 corrupted: 100% 128/128
[00:00<?, ?it/s]
train: Caching images (0.1GB ram): 100% 128/128 [00:00<00:00,
267.11it/s]
val: Scanning '../datasets/coco128/labels/train2017.cache' images and
labels... 128 found, 0 missing, 2 empty, 0 corrupted: 100% 128/128
[00:00<?, ?it/s]
val: Caching images (0.1GB ram): 100% 128/128 [00:01<00:00,
105.91it/s]
Plotting labels to runs/train/exp3/labels.jpg...

```

```

AutoAnchor: 4.27 anchors/target, 0.994 Best Possible Recall (BPR).
Current anchors are a good fit to dataset ✓
Image sizes 640 train, 640 val
Using 2 dataloader workers
Logging results to runs/train/exp3
Starting training for 5 epochs...

```

Epoch	gpu_mem	box	obj	cls	labels	img_size
0/4	3.99G	0.03622	0.05623	0.01037	250	
640: 100% 8/8	[00:29<00:00,	3.65s/it]				
	Class	Images	Labels	P		R
mAP@.5 mAP@.5:.95:	100% 4/4	[00:13<00:00,	3.48s/it]			
	all	128	929	0.688		0.79
0.805	0.586					

Epoch	gpu_mem	box	obj	cls	labels	img_size
1/4	8.5G	0.0353	0.05446	0.009651	219	
640: 100% 8/8	[00:27<00:00,	3.40s/it]				
	Class	Images	Labels	P		R
mAP@.5 mAP@.5:.95:	100% 4/4	[00:13<00:00,	3.47s/it]			
	all	128	929	0.702		0.784
0.811	0.588					

Epoch	gpu_mem	box	obj	cls	labels	img_size
2/4	8.5G	0.03818	0.05692	0.0118	280	
640: 100% 8/8	[00:27<00:00,	3.38s/it]				
	Class	Images	Labels	P		R
mAP@.5 mAP@.5:.95:	100% 4/4	[00:13<00:00,	3.47s/it]			
	all	128	929	0.697		0.786
0.811	0.591					

Epoch	gpu_mem	box	obj	cls	labels	img_size
3/4	8.5G	0.03935	0.05555	0.01291	243	
640: 100% 8/8	[00:26<00:00,	3.37s/it]				
	Class	Images	Labels	P		R
mAP@.5 mAP@.5:.95:	100% 4/4	[00:13<00:00,	3.46s/it]			
	all	128	929	0.71		0.785
0.815	0.593					

Epoch	gpu_mem	box	obj	cls	labels	img_size
4/4	8.5G	0.03757	0.05091	0.01077	188	
640: 100% 8/8 [00:27<00:00, 3.38s/it]						
	Class	Images	Labels		P	R
mAP@.5	mAP@.5:.95: 100%	4/4	[00:13<00:00, 3.45s/it]			
	all	128	929	0.747		0.762
0.816	0.594					

5 epochs completed in 0.062 hours.

Optimizer stripped from runs/train/exp3/weights/last.pt, 124.4MB

Optimizer stripped from runs/train/exp3/weights/best.pt, 124.4MB

Validating runs/train/exp3/weights/best.pt...

Fusing layers...

Model Summary: 261 layers, 61922845 parameters, 0 gradients, 156.1 GFL0Ps

	Class	Images	Labels	P	R
mAP@.5	mAP@.5:.95: 100%	4/4	[00:16<00:00, 4.02s/it]		
	all	128	929	0.748	0.762
0.816	0.594				
	person	128	254	0.854	0.783
0.853	0.626				
	bicycle	128	6	0.655	0.5
0.659	0.38				
	car	128	46	0.831	0.543
0.624	0.312				
	motorcycle	128	5	0.913	1
0.995	0.779				
	airplane	128	6	0.908	1
0.995	0.727				
	bus	128	7	0.882	0.714
0.843	0.757				
	train	128	3	0.86	1
0.995	0.895				
	truck	128	12	0.676	0.583
0.645	0.43				
	boat	128	6	0.671	0.5
0.711	0.451				
	traffic light	128	14	1	0.428
0.566	0.291				
	stop sign	128	2	0.649	1
0.995	0.821				
	bench	128	9	1	0.755
0.841	0.422				
	bird	128	16	0.966	1
0.995	0.684				
	cat	128	4	0.893	1
0.995	0.937				
	dog	128	9	0.889	0.889
0.984	0.778				

	horse	128	2	0.548	1
0.995	0.796				
	elephant	128	17	0.927	0.941
0.936	0.797				
	bear	128	1	0.671	1
0.995	0.895				
	zebra	128	4	0.865	1
0.995	0.971				
	giraffe	128	9	0.868	1
0.984	0.781				
	backpack	128	6	0.779	0.5
0.652	0.394				
	umbrella	128	18	0.871	0.889
0.917	0.625				
	handbag	128	19	0.707	0.383
0.574	0.339				
	tie	128	7	0.888	0.857
0.857	0.665				
	suitcase	128	4	0.751	1
0.995	0.722				
	frisbee	128	5	0.722	0.8
0.761	0.661				
	skis	128	1	0.698	1
0.995	0.796				
	snowboard	128	7	0.881	0.714
0.848	0.613				
	sports ball	128	6	0.817	0.751
0.809	0.483				
	kite	128	10	0.541	0.8
0.664	0.213				
	baseball bat	128	4	0.525	0.563
0.503	0.246				
	baseball glove	128	7	0.58	0.714
0.718	0.41				
	skateboard	128	5	0.816	0.899
0.962	0.434				
	tennis racket	128	7	0.635	0.571
0.603	0.378				
	bottle	128	18	0.61	0.778
0.75	0.488				
	wine glass	128	16	0.775	0.861
0.873	0.529				
	cup	128	36	0.868	0.861
0.915	0.628				
	fork	128	6	1	0.33
0.706	0.445				
	knife	128	16	0.764	0.812
0.811	0.545				
	spoon	128	22	0.786	0.591
0.631	0.441				

	bowl	128	28	0.833	0.713
0.758	0.627				
	banana	128	1	0.704	1
0.995	0.597				
	sandwich	128	2	1	0
0.638	0.61				
	orange	128	4	0.563	1
0.995	0.678				
	broccoli	128	11	0.483	0.364
0.411	0.315				
	carrot	128	24	0.751	0.75
0.802	0.516				
	hot dog	128	2	0.576	1
0.995	0.995				
	pizza	128	5	0.787	1
0.995	0.729				
	donut	128	14	0.771	1
0.941	0.85				
	cake	128	4	0.735	1
0.995	0.902				
	chair	128	35	0.776	0.695
0.778	0.465				
	couch	128	6	0.692	0.833
0.942	0.588				
	potted plant	128	14	0.796	0.857
0.897	0.618				
	bed	128	3	1	0.464
0.863	0.659				
	dining table	128	13	0.627	0.519
0.571	0.344				
	toilet	128	2	0.681	1
0.995	0.945				
	tv	128	2	0.501	1
0.995	0.821				
	laptop	128	3	0.429	0.333
0.597	0.295				
	mouse	128	2	1	0
0.662	0.217				
	remote	128	8	0.752	0.625
0.665	0.587				
	cell phone	128	8	0.694	0.625
0.635	0.389				
	microwave	128	3	0.682	1
0.995	0.864				
	oven	128	5	0.476	0.6
0.477	0.374				
	sink	128	6	0.409	0.5
0.565	0.379				
	refrigerator	128	5	0.623	0.8
0.855	0.705				

0.372	book	128	29	0.626	0.276
0.205	clock	128	9	0.915	1
0.995	vase	128	2	0.381	1
0.945	scissors	128	1	0.499	1
0.0995	teddy bear	128	21	0.946	0.841
0.609	toothbrush	128	5	0.838	1
0.777					

Results saved to runs/train/exp3

Visualization

Set up the visualization by running the below cell. Note that if you ran the training loop multiple times, you would have additional folder exp2, exp3, etc.

```
!python detect.py --weights runs/train/exp3/weights/best.pt --img 640
--conf 0.25 --source data/images
```

```
detect: weights=['runs/train/exp3/weights/best.pt'],
source=data/images, imgsz=[640, 640], conf_thres=0.25, iou_thres=0.45,
max_det=1000, device=, view_img=False, save_txt=False,
save_conf=False, save_crop=False, nosave=False, classes=None,
agnostic_nms=False, augment=False, visualize=False, update=False,
project=runs/detect, name=exp, exist_ok=False, line_thickness=3,
hide_labels=False, hide_conf=False, half=False, dnn=False
YOLOv3 46dad08 torch 1.10.0+cu111 CUDA:0 (Tesla K80, 11441MiB)
```

Fusing layers...

Model Summary: 261 layers, 61922845 parameters, 0 gradients, 156.1 GFLOPs

image 1/2

/content/csc413/a2/2022/assets/assignments/pa2-q4-files/data/images/Cats_and_dog.jpg: 480x640 3 cats, 1 dog, 1 potted plant, Done. (0.149s)

image 2/2


/content/csc413/a2/2022/assets/assignments/pa2-q4-files/data/images/bus.jpg: 640x480 4 persons, 1 bus, 2 ties, Done. (0.151s)

Speed: 0.6ms pre-process, 150.0ms inference, 1.9ms NMS per image at shape (1, 3, 640, 640)


Results saved to runs/detect/exp4

Display Results

```
display.Image(filename='runs/detect/exp4/Cats_and_dog.jpg', width=600)
```

```
display.Image(filename='runs/detect/exp4/bus.jpg', width=600)
```

bus 0.91

person 0.9

person 0.93

person 0

tie 0.47

tie 0.39

person 0.64