

Colab FAQ

For some basic overview and features offered in Colab notebooks, check out: [Overview of Colaboratory Features](#)

You need to use the colab GPU for this assignment by selecting:

Runtime → Change runtime type → Hardware Accelerator: GPU

Setup PyTorch

All files will be stored at /content/csc421/a3/ folder

```
#####  
# Setup python environment and change the current working directory  
#####  
!pip install Pillow  
%mkdir -p ./content/csc421/a3/  
%cd ./content/csc421/a3
```

Requirement already satisfied: Pillow in
/usr/local/lib/python3.7/dist-packages (7.1.2)
/content/content/csc421/a3

Helper code

Utility functions

```
%matplotlib inline
```

```
import os  
import pdb  
import argparse  
import pickle as pkl  
from pathlib import Path
```

```
from collections import defaultdict
```

```
import numpy as np  
import matplotlib as mpl  
import matplotlib.pyplot as plt  
import matplotlib.ticker as ticker
```

```
import torch  
import torch.nn as nn  
import torch.optim as optim  
import torch.nn.functional as F
```

```

from torch.autograd import Variable

from six.moves.urllib.request import urlretrieve
import tarfile
import pickle
import sys

def get_file(
    fname, origin, untar=False, extract=False, archive_format="auto",
    cache_dir="data"
):
    datadir = os.path.join(cache_dir)
    if not os.path.exists(datadir):
        os.makedirs(datadir)

    if untar:
        untar_fpath = os.path.join(datadir, fname)
        fpath = untar_fpath + ".tar.gz"
    else:
        fpath = os.path.join(datadir, fname)

    print(fpath)
    if not os.path.exists(fpath):
        print("Downloading data from", origin)

        error_msg = "URL fetch failure on {}: {} -- {}"
        try:
            try:
                urlretrieve(origin, fpath)
            except URLError as e:
                raise Exception(error_msg.format(origin, e.errno,
e.reason))
            except HTTPError as e:
                raise Exception(error_msg.format(origin, e.code,
e.msg))
            except (Exception, KeyboardInterrupt) as e:
                if os.path.exists(fpath):
                    os.remove(fpath)
                raise

        if untar:
            if not os.path.exists(untar_fpath):
                print("Extracting file.")
                with tarfile.open(fpath) as archive:
                    archive.extractall(datadir)
            return untar_fpath

    if extract:

```

```

        _extract_archive(fpath, datadir, archive_format)

    return fpath

class AttrDict(dict):
    def __init__(self, *args, **kwargs):
        super(AttrDict, self).__init__(*args, **kwargs)
        self.__dict__ = self

def to_var(tensor, cuda):
    """Wraps a Tensor in a Variable, optionally placing it on the GPU.

    Arguments:
        tensor: A Tensor object.
        cuda: A boolean flag indicating whether to use the GPU.

    Returns:
        A Variable object, on the GPU if cuda==True.
    """
    if cuda:
        return Variable(tensor.cuda())
    else:
        return Variable(tensor)

def create_dir_if_not_exists(directory):
    """Creates a directory if it doesn't already exist."""
    if not os.path.exists(directory):
        os.makedirs(directory)

def save_loss_plot(train_losses, val_losses, opts):
    """Saves a plot of the training and validation loss curves."""
    plt.figure()
    plt.plot(range(len(train_losses)), train_losses)
    plt.plot(range(len(val_losses)), val_losses)
    plt.title("BS={}, nhid={}".format(opts.batch_size,
opts.hidden_size), fontsize=20)
    plt.xlabel("Epochs", fontsize=16)
    plt.ylabel("Loss", fontsize=16)
    plt.xticks(fontsize=14)
    plt.yticks(fontsize=14)
    plt.tight_layout()
    plt.savefig(os.path.join(opts.checkpoint_path, "loss_plot.pdf"))
    plt.close()

```

```

def save_loss_comparison_gru(l1, l2, o1, o2, fn, s=500):
    """Plot comparison of training and val loss curves from GRU runs.

    Arguments:
        l1: Tuple of lists containing training / val losses for model
1.
        l2: Tuple of lists containing training / val losses for model
2.
        o1: Options for model 1.
        o2: Options for model 2.
        fn: Output file name.
        s: Number of training iterations to average over.
    """
    mean_l1 = [np.mean(l1[0][i * s : (i + 1) * s]) for i in
range(len(l1[0]) // s)]
    mean_l2 = [np.mean(l2[0][i * s : (i + 1) * s]) for i in
range(len(l2[0]) // s)]

    plt.figure()

    fig, ax = plt.subplots(1, 2, figsize=(10, 4))

    ax[0].plot(range(len(mean_l1)), mean_l1, label="ds=" +
o1.data_file_name)
    ax[0].plot(range(len(mean_l2)), mean_l2, label="ds=" +
o2.data_file_name)
    ax[0].title.set_text("Train Loss | GRU Hidden Size =
{}".format(o2.hidden_size))

    # Validation losses are assumed to be by epoch
    ax[1].plot(range(len(l1[1])), l1[1], label="ds=" +
o1.data_file_name)
    ax[1].plot(range(len(l2[1])), l2[1], label="ds=" +
o2.data_file_name)
    ax[1].title.set_text("Val Loss | GRU Hidden Size =
{}".format(o2.hidden_size))

    ax[0].set_xlabel("Iterations (x{})".format(s), fontsize=10)
    ax[0].set_ylabel("Loss", fontsize=10)
    ax[1].set_xlabel("Epochs", fontsize=10)
    ax[1].set_ylabel("Loss", fontsize=10)
    ax[0].legend(loc="upper right")
    ax[1].legend(loc="upper right")

    fig.suptitle("GRU Performance by Dataset", fontsize=14)
    plt.tight_layout()
    fig.subplots_adjust(top=0.85)
    plt.legend()

```

```

plt_path = "./loss_plot_{}.pdf".format(fn)
plt.savefig(plt_path)
print(f"Plot saved to: {Path(plt_path).resolve()}")

def save_loss_comparison_by_dataset(l1, l2, l3, l4, o1, o2, o3, o4,
fn, s=500):
    """Plot comparison of training and validation loss curves from all
four
runs in Part 3, comparing by dataset while holding hidden size
constant.

Models within each pair (l1, l2) and (l3, l4) have the same hidden
sizes.

Arguments:
1. l1: Tuple of lists containing training / val losses for model
2. l2: Tuple of lists containing training / val losses for model
3. l3: Tuple of lists containing training / val losses for model
4. l4: Tuple of lists containing training / val losses for model
o1: Options for model 1.
o2: Options for model 2.
o3: Options for model 3.
o4: Options for model 4.
fn: Output file name.
s: Number of training iterations to average over.
"""
    mean_l1 = [np.mean(l1[0][i * s : (i + 1) * s]) for i in
range(len(l1[0]) // s)]
    mean_l2 = [np.mean(l2[0][i * s : (i + 1) * s]) for i in
range(len(l2[0]) // s)]
    mean_l3 = [np.mean(l3[0][i * s : (i + 1) * s]) for i in
range(len(l3[0]) // s)]
    mean_l4 = [np.mean(l4[0][i * s : (i + 1) * s]) for i in
range(len(l4[0]) // s)]

    plt.figure()
    fig, ax = plt.subplots(2, 2, figsize=(10, 8))

    ax[0][0].plot(range(len(mean_l1)), mean_l1, label="ds=" +
o1.data_file_name)
    ax[0][0].plot(range(len(mean_l2)), mean_l2, label="ds=" +
o2.data_file_name)
    ax[0][0].title.set_text(
        "Train Loss | Model Hidden Size = {}".format(o1.hidden_size)
    )

```

```

        # Validation losses are assumed to be by epoch
        ax[0][1].plot(range(len(l1[1])), l1[1], label="ds=" +
o1.data_file_name)
        ax[0][1].plot(range(len(l2[1])), l2[1], label="ds=" +
o2.data_file_name)
        ax[0][1].title.set_text("Val Loss | Model Hidden Size =
{}".format(o1.hidden_size))

        ax[1][0].plot(range(len(mean_l3)), mean_l3, label="ds=" +
o3.data_file_name)
        ax[1][0].plot(range(len(mean_l4)), mean_l4, label="ds=" +
o4.data_file_name)
        ax[1][0].title.set_text(
            "Train Loss | Model Hidden Size = {}".format(o3.hidden_size)
        )

        ax[1][1].plot(range(len(l3[1])), l3[1], label="ds=" +
o3.data_file_name)
        ax[1][1].plot(range(len(l4[1])), l4[1], label="ds=" +
o4.data_file_name)
        ax[1][1].title.set_text("Val Loss | Model Hidden Size =
{}".format(o4.hidden_size))

    for i in range(2):
        ax[i][0].set_xlabel("Iterations (x{})".format(s), fontsize=10)
        ax[i][0].set_ylabel("Loss", fontsize=10)
        ax[i][1].set_xlabel("Epochs", fontsize=10)
        ax[i][1].set_ylabel("Loss", fontsize=10)
        ax[i][0].legend(loc="upper right")
        ax[i][1].legend(loc="upper right")

    fig.suptitle("Performance by Dataset Size", fontsize=16)
    plt.tight_layout()
    fig.subplots_adjust(top=0.9)
    plt.legend()
    plt.savefig("./loss_plot_{}.pdf".format(fn))
    plt.close()

```

```

def save_loss_comparison_by_hidden(l1, l2, l3, l4, o1, o2, o3, o4, fn,
s=500):
    """Plot comparison of training and validation loss curves from all
four
    runs in Part 3, comparing by hidden size while holding dataset
constant.

```

Models within each pair (l1, l3) and (l2, l4) have the same dataset.

```

Arguments:
    l1: Tuple of lists containing training / val losses for model
1.
    l2: Tuple of lists containing training / val losses for model
2.
    l3: Tuple of lists containing training / val losses for model
3.
    l4: Tuple of lists containing training / val losses for model
4.
    o1: Options for model 1.
    o2: Options for model 2.
    o3: Options for model 3.
    o4: Options for model 4.
    fn: Output file name.
    s: Number of training iterations to average over.
"""
    mean_l1 = [np.mean(l1[0][i * s : (i + 1) * s]) for i in
range(len(l1[0]) // s)]
    mean_l2 = [np.mean(l2[0][i * s : (i + 1) * s]) for i in
range(len(l2[0]) // s)]
    mean_l3 = [np.mean(l3[0][i * s : (i + 1) * s]) for i in
range(len(l3[0]) // s)]
    mean_l4 = [np.mean(l4[0][i * s : (i + 1) * s]) for i in
range(len(l4[0]) // s)]

    plt.figure()
    fig, ax = plt.subplots(2, 2, figsize=(10, 8))

    ax[0][0].plot(range(len(mean_l1)), mean_l1, label="hid_size=" +
str(o1.hidden_size))
    ax[0][0].plot(range(len(mean_l3)), mean_l3, label="hid_size=" +
str(o3.hidden_size))
    ax[0][0].title.set_text("Train Loss | Dataset = " +
o1.data_file_name)

    # Validation losses are assumed to be by epoch
    ax[0][1].plot(range(len(l1[1])), l1[1], label="hid_size=" +
str(o1.hidden_size))
    ax[0][1].plot(range(len(l3[1])), l3[1], label="hid_size=" +
str(o3.hidden_size))
    ax[0][1].title.set_text("Val Loss | Dataset = " +
o1.data_file_name)

    ax[1][0].plot(range(len(mean_l2)), mean_l2, label="hid_size=" +
str(o2.hidden_size))
    ax[1][0].plot(range(len(mean_l4)), mean_l4, label="hid_size=" +
str(o4.hidden_size))
    ax[1][0].title.set_text("Train Loss | Dataset = " +
o3.data_file_name)

```

```

    ax[1][1].plot(range(len(l2[1])), l2[1], label="hid_size=" +
str(o2.hidden_size))
    ax[1][1].plot(range(len(l4[1])), l4[1], label="hid_size=" +
str(o4.hidden_size))
    ax[1][1].title.set_text("Val Loss | Dataset = " +
o4.data_file_name)

    for i in range(2):
        ax[i][0].set_xlabel("Iterations (x{})".format(s), fontsize=10)
        ax[i][0].set_ylabel("Loss", fontsize=10)
        ax[i][1].set_xlabel("Epochs", fontsize=10)
        ax[i][1].set_ylabel("Loss", fontsize=10)
        ax[i][0].legend(loc="upper right")
        ax[i][1].legend(loc="upper right")

fig.suptitle("Performance by Hidden State Size", fontsize=16)
plt.tight_layout()
fig.subplots_adjust(top=0.9)
plt.legend()
plt.savefig("./loss_plot_{}.pdf".format(fn))
plt.close()

```

```

def checkpoint(encoder, decoder, idx_dict, opts):
    """Saves the current encoder and decoder models, along with
    idx_dict, which
    contains the char_to_index and index_to_char mappings, and the
    start_token
    and end_token values.
    """
    with open(os.path.join(opts.checkpoint_path, "encoder.pt"), "wb")
as f:
        torch.save(encoder, f)

    with open(os.path.join(opts.checkpoint_path, "decoder.pt"), "wb")
as f:
        torch.save(decoder, f)

    with open(os.path.join(opts.checkpoint_path, "idx_dict.pkl"),
"wb") as f:
        pkl.dump(idx_dict, f)

```

Data loader

```

def read_lines(filename):
    """Read a file and split it into lines."""
    lines = open(filename).read().strip().lower().split("\n")
    return lines

```



```

def read_pairs(filename):
    """Reads lines that consist of two words, separated by a space.

    Returns:
        source_words: A list of the first word in each line of the
        file.
        target_words: A list of the second word in each line of the
        file.
    """
    lines = read_lines(filename)
    source_words, target_words = [], []
    for line in lines:
        line = line.strip()
        if line:
            source, target = line.split()
            source_words.append(source)
            target_words.append(target)
    return source_words, target_words


def all_alpha_or_dash(s):
    """Helper function to check whether a string is alphabetic,
    allowing dashes '-'. """
    return all(c.isalpha() or c == "-" for c in s)


def filter_lines(lines):
    """Filters lines to consist of only alphabetic characters or
    dashes "-". """
    return [line for line in lines if all_alpha_or_dash(line)]


def load_data(file_name):
    """Loads (English, Pig-Latin) word pairs, and creates mappings
    from characters to indexes. """
    path = "./data/{}.txt".format(file_name)
    source_lines, target_lines = read_pairs(path)

    # Filter lines
    source_lines = filter_lines(source_lines)
    target_lines = filter_lines(target_lines)

    all_characters = set("".join(source_lines)) |
    set("".join(target_lines))

    # Create a dictionary mapping each character to a unique index
    char_to_index = {
        char: index for (index, char) in

```

```

enumerate(sorted(list(all_characters)))
}

# Add start and end tokens to the dictionary
start_token = len(char_to_index)
end_token = len(char_to_index) + 1
char_to_index["SOS"] = start_token
char_to_index["EOS"] = end_token

# Create the inverse mapping, from indexes to characters (used to
decode the model's predictions)
index_to_char = {index: char for (char, index) in
char_to_index.items()}

# Store the final size of the vocabulary
vocab_size = len(char_to_index)

line_pairs = list(set(zip(source_lines, target_lines))) # Python
3

idx_dict = {
    "char_to_index": char_to_index,
    "index_to_char": index_to_char,
    "start_token": start_token,
    "end_token": end_token,
}

return line_pairs, vocab_size, idx_dict

def create_dict(pairs):
    """Creates a mapping { (source_length, target_length): [list of
(source, target) pairs]
This is used to make batches: each batch consists of two parallel
tensors, one containing
all source indexes and the other containing all corresponding
target indexes.
Within a batch, all the source words are the same length, and all
the target words are
the same length.
"""
    unique_pairs = list(set(pairs)) # Find all unique (source,
target) pairs

    d = defaultdict(list)
    for (s, t) in unique_pairs:
        d[(len(s), len(t))].append((s, t))

    return d

```

Training and evaluation code

```
def string_to_index_list(s, char_to_index, end_token):
    """Converts a sentence into a list of indexes (for each
    character)."""
    return [char_to_index[char] for char in s] + [
        end_token
    ] # Adds the end token to each index list

def translate_sentence(sentence, encoder, decoder, idx_dict, opts):
    """Translates a sentence from English to Pig-Latin, by splitting
    the sentence into
    words (whitespace-separated), running the encoder-decoder model to
    translate each
    word independently, and then stitching the words back together
    with spaces between them.
    """
    if idx_dict is None:
        line_pairs, vocab_size, idx_dict =
load_data(opts["data_file_name"])
    return " ".join(
        [translate(word, encoder, decoder, idx_dict, opts) for word in
sentence.split()]
    )

def translate(input_string, encoder, decoder, idx_dict, opts):
    """Translates a given string from English to Pig-Latin."""

    char_to_index = idx_dict["char_to_index"]
    index_to_char = idx_dict["index_to_char"]
    start_token = idx_dict["start_token"]
    end_token = idx_dict["end_token"]

    max_generated_chars = 20
    gen_string = ""

    indexes = string_to_index_list(input_string, char_to_index,
end_token)
    indexes = to_var(
        torch.LongTensor(indexes).unsqueeze(0), opts.cuda
    ) # Unsqueeze to make it like BS = 1

    encoder_annotations, encoder_last_hidden = encoder(indexes)

    decoder_hidden = encoder_last_hidden
    decoder_input = to_var(torch.LongTensor([[start_token]]),
opts.cuda) # For BS = 1
    decoder_inputs = decoder_input
```

```

for i in range(max_generated_chars):
    ## slow decoding, recompute everything at each time
    decoder_outputs, attention_weights = decoder(
        decoder_inputs, encoder_annotations, decoder_hidden
    )

    generated_words = F.softmax(decoder_outputs, dim=2).max(2)[1]
    ni = generated_words.cpu().numpy().reshape(-1) # LongTensor
of size 1
    ni = ni[-1] # latest output token

    decoder_inputs = torch.cat([decoder_input, generated_words],
dim=1)

    if ni == end_token:
        break
    else:
        gen_string = "".join(
            [
                index_to_char[int(item)]
                for item in
generated_words.cpu().numpy().reshape(-1)
            ]
        )

return gen_string

```

```

def visualize_attention(input_string, encoder, decoder, idx_dict,
opts):
    """Generates a heatmap to show where attention is focused in each
decoder step."""
    if idx_dict is None:
        line_pairs, vocab_size, idx_dict =
load_data(opts["data_file_name"])
        char_to_index = idx_dict["char_to_index"]
        index_to_char = idx_dict["index_to_char"]
        start_token = idx_dict["start_token"]
        end_token = idx_dict["end_token"]

        max_generated_chars = 20
        gen_string = ""

        indexes = string_to_index_list(input_string, char_to_index,
end_token)
        indexes = to_var(
            torch.LongTensor(indexes).unsqueeze(0), opts.cuda
        ) # Unsqueeze to make it like BS = 1

```

```

encoder_annotations, encoder_hidden = encoder(indexes)

decoder_hidden = encoder_hidden
decoder_input = to_var(torch.LongTensor([[start_token]]),
opts.cuda) # For BS = 1
decoder_inputs = decoder_input

produced_end_token = False

for i in range(max_generated_chars):
    ## slow decoding, recompute everything at each time
    decoder_outputs, attention_weights = decoder(
        decoder_inputs, encoder_annotations, decoder_hidden
    )
    generated_words = F.softmax(decoder_outputs, dim=2).max(2)[1]
    ni = generated_words.cpu().numpy().reshape(-1) # LongTensor
of size 1
    ni = ni[-1] # latest output token

    decoder_inputs = torch.cat([decoder_input, generated_words],
dim=1)

    if ni == end_token:
        break
    else:
        gen_string = "".join(
            [
                index_to_char[int(item)]
                for item in
generated_words.cpu().numpy().reshape(-1)
            ]
        )

    if isinstance(attention_weights, tuple):
        ## transformer's attention mweights
        attention_weights, self_attention_weights = attention_weights

    all_attention_weights = attention_weights.data.cpu().numpy()

    for i in range(len(all_attention_weights)):
        attention_weights_matrix = all_attention_weights[i].squeeze()
        fig = plt.figure()
        ax = fig.add_subplot(111)
        cax = ax.matshow(attention_weights_matrix, cmap="bone")
        fig.colorbar(cax)

    # Set up axes
    ax.set_yticklabels([""] + list(input_string) + ["EOS"],

```

```

rotation=90)
    ax.set_xticklabels(
        [""] + list(gen_string) + (["EOS"] if produced_end_token
else [])
    )

    # Show label at every tick
    ax.xaxis.set_major_locator(ticker.MultipleLocator(1))
    ax.yaxis.set_major_locator(ticker.MultipleLocator(1))
    # Add title
    plt.xlabel("Attention weights to the source sentence in layer
{}".format(i + 1))
    plt.tight_layout()
    plt.grid("off")
    plt.show()

    return gen_string

```

```

def compute_loss(data_dict, encoder, decoder, idx_dict, criterion,
optimizer, opts):

```

"""Train/Evaluate the model on a dataset.

Arguments:

data_dict: The validation/test word pairs, organized by source and target lengths.

encoder: An encoder model to produce annotations for each step of the input sequence.

decoder: A decoder model (with or without attention) to generate output tokens.

idx_dict: Contains char-to-index and index-to-char mappings, and start & end token indexes.

criterion: Used to compute the CrossEntropyLoss for each decoder output.

optimizer: Train the weights if an optimizer is given. None if only evaluate the model.

opts: The command-line arguments.

Returns:

mean_loss: The average loss over all batches from data_dict.
"""

```

start_token = idx_dict["start_token"]
end_token = idx_dict["end_token"]
char_to_index = idx_dict["char_to_index"]

```

```

losses = []

```

```

for key in data_dict:

```

```

    input_strings, target_strings = zip(*data_dict[key])

```

```

    input_tensors = [

```

```

        torch.LongTensor(string_to_index_list(s, char_to_index,
end_token))
        for s in input_strings
    ]
    target_tensors = [
        torch.LongTensor(string_to_index_list(s, char_to_index,
end_token))
        for s in target_strings
    ]

    num_tensors = len(input_tensors)
    num_batches = int(np.ceil(num_tensors /
float(opts.batch_size)))

    for i in range(num_batches):

        start = i * opts.batch_size
        end = start + opts.batch_size

        inputs = to_var(torch.stack(input_tensors[start:end]),
opts.cuda)
        targets = to_var(torch.stack(target_tensors[start:end]),
opts.cuda)

        # The batch size may be different in each epoch
        BS = inputs.size(0)

        encoder_annotations, encoder_hidden = encoder(inputs)

        # The last hidden state of the encoder becomes the first
hidden state of the decoder
        decoder_hidden = encoder_hidden

        start_vector = (
            torch.ones(BS).long().unsqueeze(1) * start_token
        ) # BS x 1 --> 16x1 CHECKED
        decoder_input = to_var(start_vector, opts.cuda) # BS x 1
--> 16x1 CHECKED

        loss = 0.0

        seq_len = targets.size(1) # Gets seq_len from BS x
seq_len

        decoder_inputs = torch.cat(
            [decoder_input, targets[:, 0:-1]], dim=1
        ) # Gets decoder inputs by shifting the targets to the
right

```

```

        decoder_outputs, attention_weights = decoder(
            decoder_inputs, encoder_annotations, decoder_hidden
        )
        decoder_outputs_flatten = decoder_outputs.view(-1,
decoder_outputs.size(2))
        targets_flatten = targets.view(-1)

        loss = criterion(decoder_outputs_flatten, targets_flatten)

        losses.append(loss.item())

## training if an optimizer is provided
    if optimizer:
        # Zero gradients
        optimizer.zero_grad()
        # Compute gradients
        loss.backward()
        # Update the parameters of the encoder and decoder
        optimizer.step()

    return losses

```

```

def training_loop(
    train_dict, val_dict, idx_dict, encoder, decoder, criterion,
    optimizer, opts
):
    """Runs the main training loop; evaluates the model on the val set
every epoch.
    * Prints training and val loss each epoch.
    * Prints qualitative translation results each epoch using
TEST_SENTENCE
    * Saves an attention map for TEST_WORD_ATTN each epoch
    * Returns loss curves for comparison


```

Arguments:

- train_dict: The training word pairs, organized by source and target lengths.*
- val_dict: The validation word pairs, organized by source and target lengths.*
- idx_dict: Contains char-to-index and index-to-char mappings, and start & end token indexes.*
- encoder: An encoder model to produce annotations for each step of the input sequence.*
- decoder: A decoder model (with or without attention) to generate output tokens.*
- criterion: Used to compute the CrossEntropyLoss for each decoder output.*
- optimizer: Implements a step rule to update the parameters of*

the encoder and decoder.

opts: The command-line arguments.

Returns:

losses: Lists containing training and validation loss curves.
"""

```
start_token = idx_dict["start_token"]
end_token = idx_dict["end_token"]
char_to_index = idx_dict["char_to_index"]

loss_log = open(os.path.join(opts.checkpoint_path,
"loss_log.txt"), "w")

best_val_loss = 1e6
train_losses = []
val_losses = []

mean_train_losses = []
mean_val_losses = []

early_stopping_counter = 0

for epoch in range(opts.nepochs):

    optimizer.param_groups[0]["lr"] *= opts.lr_decay

    train_loss = compute_loss(
        train_dict, encoder, decoder, idx_dict, criterion,
optimizer, opts
    )
    val_loss = compute_loss(
        val_dict, encoder, decoder, idx_dict, criterion, None,
opts
    )

    mean_train_loss = np.mean(train_loss)
    mean_val_loss = np.mean(val_loss)

    if mean_val_loss < best_val_loss:
        checkpoint(encoder, decoder, idx_dict, opts)
        best_val_loss = mean_val_loss
        early_stopping_counter = 0
    else:
        early_stopping_counter += 1

    if early_stopping_counter > opts.early_stopping_patience:
        print(
            "Validation loss has not improved in {} epochs,
```

```

        stopping_early".format(
            opts.early_stopping_patience
        )
    )
    print("Obtained lowest validation loss of:
{}".format(best_val_loss))
    return (train_losses, mean_val_losses)

    gen_string = translate_sentence(TEST_SENTENCE, encoder,
decoder, idx_dict, opts)
    print(
        "Epoch: {:3d} | Train loss: {:.3f} | Val loss: {:.3f} |
Gen: {:20s}".format(
            epoch, mean_train_loss, mean_val_loss, gen_string
        )
    )

    loss_log.write("{} {} {}\n".format(epoch, train_loss,
val_loss))
    loss_log.flush()

    train_losses += train_loss
    val_losses += val_loss

    mean_train_losses.append(mean_train_loss)
    mean_val_losses.append(mean_val_loss)

    save_loss_plot(mean_train_losses, mean_val_losses, opts)

    print("Obtained lowest validation loss of:
{}".format(best_val_loss))
    return (train_losses, mean_val_losses)

def print_data_stats(line_pairs, vocab_size, idx_dict):
    """Prints example word pairs, the number of data points, and the
vocabulary."""
    print("=" * 80)
    print("Data Stats".center(80))
    print("-" * 80)
    for pair in line_pairs[:5]:
        print(pair)
    print("Num unique word pairs: {}".format(len(line_pairs)))
    print("Vocabulary: {}".format(idx_dict["char_to_index"].keys()))
    print("Vocab size: {}".format(vocab_size))
    print("=" * 80)

def train(opts):

```

```

    line_pairs, vocab_size, idx_dict =
load_data(opts["data_file_name"])
    print_data_stats(line_pairs, vocab_size, idx_dict)

    # Split the line pairs into an 80% train and 20% val split
    num_lines = len(line_pairs)
    num_train = int(0.8 * num_lines)
    train_pairs, val_pairs = line_pairs[:num_train],
line_pairs[num_train:]

    # Group the data by the lengths of the source and target words, to
form batches
    train_dict = create_dict(train_pairs)
    val_dict = create_dict(val_pairs)

#####
####
    ### Setup: Create Encoder, Decoder, Learning Criterion, and
Optimizers ###

#####
####
    if opts.encoder_type == "rnn":
        encoder = GRUEncoder(
            vocab_size=vocab_size, hidden_size=opts.hidden_size,
opts=opts
        )
    elif opts.encoder_type == "transformer":
        encoder = TransformerEncoder(
            vocab_size=vocab_size,
            hidden_size=opts.hidden_size,
            num_layers=opts.num_transformer_layers,
            opts=opts,
        )
    elif opts.encoder_type == "attention":
        encoder = AttentionEncoder(
            vocab_size=vocab_size,
            hidden_size=opts.hidden_size,
            opts=opts,
        )
    else:
        raise NotImplementedError

    if opts.decoder_type == "rnn":
        decoder = RNNDecoder(vocab_size=vocab_size,
hidden_size=opts.hidden_size)
    elif opts.decoder_type == "rnn_attention":
        decoder = RNNAttentionDecoder(
            vocab_size=vocab_size,

```

```

        hidden_size=opts.hidden_size,
        attention_type=opts.attention_type,
    )
elif opts.decoder_type == "transformer":
    decoder = TransformerDecoder(
        vocab_size=vocab_size,
        hidden_size=opts.hidden_size,
        num_layers=opts.num_transformer_layers,
    )
elif opts.encoder_type == "attention":
    decoder = AttentionDecoder(
        vocab_size=vocab_size,
        hidden_size=opts.hidden_size,
    )
else:
    raise NotImplementedError

#### setup checkpoint path
model_name = "h{}-bs{}-{}-{}".format(
    opts.hidden_size, opts.batch_size, opts.decoder_type,
    opts.data_file_name
)
opts.checkpoint_path = model_name
create_dir_if_not_exists(opts.checkpoint_path)
####

if opts.cuda:
    encoder.cuda()
    decoder.cuda()
    print("Moved models to GPU!")

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(
    list(encoder.parameters()) + list(decoder.parameters()),
    lr=opts.learning_rate
)

try:
    losses = training_loop(
        train_dict, val_dict, idx_dict, encoder, decoder,
        criterion, optimizer, opts
    )
except KeyboardInterrupt:
    print("Exiting early from training.")
    return encoder, decoder, losses

return encoder, decoder, losses

```

```
def print_opts(opts):
    """Prints the values of all command-line arguments."""
    print("=" * 80)
    print("Opts".center(80))
    print("-" * 80)
    for key in opts.__dict__:
        print("{:>30}: {:<30}".format(key,
opts.__dict__[key]).center(80))
    print("=" * 80)
```

Download dataset

```
#####
# Download Translation datasets
#####
data_fpath = get_file(
    fname="pig_latin_small.txt",
    origin="http://www.cs.toronto.edu/~jba/pig_latin_small.txt",
    untar=False,
)

data_fpath = get_file(
    fname="pig_latin_large.txt",
    origin="http://www.cs.toronto.edu/~jba/pig_latin_large.txt",
    untar=False,
)

data/pig_latin_small.txt
Downloading data from
http://www.cs.toronto.edu/~jba/pig_latin_small.txt
data/pig_latin_large.txt
Downloading data from
http://www.cs.toronto.edu/~jba/pig_latin_large.txt
```

Part 1: Neural machine translation (NMT)

In this section, you will implement a Gated Recurrent Unit (GRU) cell, a common type of recurrent neural network (RNN). The GRU cell is a simplification of the Long Short-Term Memory cell. Therefore, we have provided you with an implemented LSTM cell (MyLSTMCell), which you can reference when completing MyGRUCell.

```
class MyLSTMCell(nn.Module):
    def __init__(self, input_size, hidden_size):
        super(MyLSTMCell, self).__init__()

        self.input_size = input_size
        self.hidden_size = hidden_size

        self.Wif = nn.Linear(input_size, hidden_size)
        self.Whf = nn.Linear(hidden_size, hidden_size)
```

```

self.Wii = nn.Linear(input_size, hidden_size)
self.Whi = nn.Linear(hidden_size, hidden_size)

self.Wic = nn.Linear(input_size, hidden_size)
self.Whc = nn.Linear(hidden_size, hidden_size)

self.Wio = nn.Linear(input_size, hidden_size)
self.Who = nn.Linear(hidden_size, hidden_size)

def forward(self, x, h_prev, c_prev):
    """Forward pass of the LSTM computation for one time step.

    Arguments
        x: batch_size x input_size
        h_prev: batch_size x hidden_size
        c_prev: batch_size x hidden_size

    Returns:
        h_new: batch_size x hidden_size
        c_new: batch_size x hidden_size
    """

    f = torch.sigmoid(self.Wif(x) + self.Whf(h_prev))
    i = torch.sigmoid(self.Wii(x) + self.Whi(h_prev))

    c = torch.tanh(self.Wic(x) + self.Whc(h_prev))
    o = torch.sigmoid(self.Wio(x) + self.Who(h_prev))

    c_new = f * c_prev + i * c
    h_new = o * torch.tanh(c_new)

    return h_new, c_new

```

Step 1: GRU Cell

Please implement the MyGRUCell class defined in the next cell.

```

class MyGRUCell(nn.Module):
    def __init__(self, input_size, hidden_size):
        super(MyGRUCell, self).__init__()

        self.input_size = input_size
        self.hidden_size = hidden_size

        # Input linear layers
        self.Wiz = nn.Linear(input_size, hidden_size)
        self.Wir = nn.Linear(input_size, hidden_size)
        self.Wih = nn.Linear(input_size, hidden_size)

```

```

# Hidden linear layers
self.Whz = nn.Linear(hidden_size, hidden_size)
self.Whr = nn.Linear(hidden_size, hidden_size)
self.Whh = nn.Linear(hidden_size, hidden_size)

def forward(self, x, h_prev):
    """Forward pass of the GRU computation for one time step.

    Arguments
        x: batch_size x input_size
        h_prev: batch_size x hidden_size

    Returns:
        h_new: batch_size x hidden_size
    """

    z = torch.sigmoid(self.Wiz(x) + self.Whz(h_prev))
    r = torch.sigmoid(self.Wir(x) + self.Whr(h_prev))
    g = torch.tanh(self.Wih(x) + r * (self.Whh(h_prev * r)))
    h_new = (1 - z) * h_prev + z * g
    return h_new

```

Step 2: GRU Encoder

The following cells use your MyGRUCell implementation to build a recurrent encoder and decoder. Please read the implementations to understand what they do and run the cells before proceeding.

```

class GRUEncoder(nn.Module):
    def __init__(self, vocab_size, hidden_size, opts):
        super(GRUEncoder, self).__init__()

        self.vocab_size = vocab_size
        self.hidden_size = hidden_size
        self.opts = opts

        self.embedding = nn.Embedding(vocab_size, hidden_size)
        self.gru = MyGRUCell(hidden_size, hidden_size)

    def forward(self, inputs):
        """Forward pass of the encoder RNN.

        Arguments:
            inputs: Input token indexes across a batch for all time
            steps in the sequence. (batch_size x seq_len)

        Returns:
            annotations: The hidden states computed at each step of

```

```

the input sequence. (batch_size x seq_len x hidden_size)
    hidden: The final hidden state of the encoder, for each
sequence in a batch. (batch_size x hidden_size)
    """

    batch_size, seq_len = inputs.size()
    hidden = self.init_hidden(batch_size)

    encoded = self.embedding(inputs) # batch_size x seq_len x
hidden_size
    annotations = []

    for i in range(seq_len):
        x = encoded[:, i, :] # Get the current time step, across
the whole batch
        hidden = self.gru(x, hidden)
        annotations.append(hidden)

    annotations = torch.stack(annotations, dim=1)
    return annotations, hidden

def init_hidden(self, bs):
    """Creates a tensor of zeros to represent the initial hidden
states
of a batch of sequences.

Arguments:
    bs: The batch size for the initial hidden state.

Returns:
    hidden: An initial hidden state of all zeros. (batch_size
x hidden_size)
    """
    return to_var(torch.zeros(bs, self.hidden_size),
self.opts.cuda)

class RNNDecoder(nn.Module):
    def __init__(self, vocab_size, hidden_size):
        super(RNNDecoder, self).__init__()
        self.vocab_size = vocab_size
        self.hidden_size = hidden_size

        self.embedding = nn.Embedding(vocab_size, hidden_size)
        self.rnn = MyGRUCell(input_size=hidden_size,
hidden_size=hidden_size)
        self.out = nn.Linear(hidden_size, vocab_size)

    def forward(self, inputs, annotations, hidden_init):
        """Forward pass of the non-attentional decoder RNN.

```



```

    Arguments:
        inputs: Input token indexes across a batch. (batch_size x
seq_len)
        annotations: This is not used here. It just maintains
consistency with the
        interface used by the AttentionDecoder class.
        hidden_init: The hidden states from the last step of
encoder, across a batch. (batch_size x hidden_size)

    Returns:
        output: Un-normalized scores for each token in the
vocabulary, across a batch for all the decoding time steps.
(batch_size x decoder_seq_len x vocab_size)
        None
    """
    batch_size, seq_len = inputs.size()
    embed = self.embedding(inputs) # batch_size x seq_len x
hidden_size

    hiddens = []
    h_prev = hidden_init

    for i in range(seq_len):
        x = embed[
            :, i, :
        ] # Get the current time step input tokens, across the
whole batch
        h_prev = self.rnn(x, h_prev) # batch_size x hidden_size
        hiddens.append(h_prev)

    hiddens = torch.stack(hiddens, dim=1) # batch_size x seq_len
x hidden_size

    output = self.out(hiddens) # batch_size x seq_len x
vocab_size
    return output, None

```

Step 3: Training and Analysis

Train the encoder-decoder model to perform English --> Pig Latin translation. We will start by training on the smaller dataset.

```
TEST_SENTENCE = "the air conditioning is working"
```

```

rnn_args_s = AttrDict()
args_dict = {
    "data_file_name": "pig_latin_small",
    "cuda": True,
    "nepochs": 50,

```

```

    "checkpoint_dir": "checkpoints",
    "learning_rate": 0.005,
    "lr_decay": 0.99,
    "early_stopping_patience": 20,
    "batch_size": 64,
    "hidden_size": 32,
    "encoder_type": "rnn", # options: rnn / transformer
    "decoder_type": "rnn", # options: rnn / rnn_attention /
transformer
    "attention_type": "", # options: additive / scaled_dot
}
rnn_args_s.update(args_dict)

print_opts(rnn_args_s)
rnn_encode_s, rnn_decoder_s, rnn_losses_s = train(rnn_args_s)

translated = translate_sentence(
    TEST_SENTENCE, rnn_encode_s, rnn_decoder_s, None, rnn_args_s
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```

=====
=====

```

Opts

```

-----
-----

```

```

data_file_name: pig_latin_small
cuda: 1
nepochs: 50
checkpoint_dir: checkpoints
learning_rate: 0.005
lr_decay: 0.99
early_stopping_patience: 20
batch_size: 64
hidden_size: 32
encoder_type: rnn
decoder_type: rnn

```

attention_type:

```
=====
=====
=====
=====
```

Data Stats

```
-----
-----
('folded', 'oldedfay')
('supposition', 'uppositionsay')
('advancing', 'advancingway')
('reconcile', 'econcileray')
('leagued', 'eaguedlay')
Num unique word pairs: 3198
Vocabulary: dict_keys(['-', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h',
'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v',
'w', 'x', 'y', 'z', 'SOS', 'EOS'])
Vocab size: 29
=====
=====
```

Moved models to GPU!

```
Epoch:  0 | Train loss: 2.297 | Val loss: 2.022 | Gen: ingay-onsay-
onsay-on ingay-onsay-onsay-on ingsay-onsay-onsay-o ingay-onsay-onsay-
on ingsay-onsay-onsay-o
Epoch:  1 | Train loss: 1.891 | Val loss: 1.827 | Gen: eway alway
onssay ingsay onsay
Epoch:  2 | Train loss: 1.715 | Val loss: 1.728 | Gen: eway away-
onday oonsay isteray oonsay
Epoch:  3 | Train loss: 1.601 | Val loss: 1.682 | Gen: eday away-
onday onday-onday isteray onday-onday
Epoch:  4 | Train loss: 1.517 | Val loss: 1.644 | Gen: eway away-
onday onday-onday-onday iway otay-onday
Epoch:  5 | Train loss: 1.443 | Val loss: 1.597 | Gen: eway away-
ingsay-onday onday-onday-onday iway otay-onday
Epoch:  6 | Train loss: 1.370 | Val loss: 1.564 | Gen: eway away-otay
ondeday-otay-otay iway otay-onday-atersay
Epoch:  7 | Train loss: 1.312 | Val loss: 1.548 | Gen: eway away-
away-away-aters otay-ingsay-onday-at iway otay
Epoch:  8 | Train loss: 1.261 | Val loss: 1.550 | Gen: eway aiway-
ingsay-oday oncingsay-oday iway otay-away-awlay
Epoch:  9 | Train loss: 1.216 | Val loss: 1.493 | Gen: eway away-
ingsray-outedwa oncitersay-oday isway ondingsray-oteray
Epoch: 10 | Train loss: 1.167 | Val loss: 1.471 | Gen: eway away-
ingsray-outedwa onciinessay-ingway-in isway oringsay-oday
Epoch: 11 | Train loss: 1.132 | Val loss: 1.517 | Gen: eway aindedway
oncienceringsray-ote isway onmouredway
Epoch: 12 | Train loss: 1.113 | Val loss: 1.499 | Gen: eway iway-
```

indway-awlay oncienceday isway ortisedway
Epoch: 13 | Train loss: 1.076 | Val loss: 1.509 | Gen: eway aindway
oncionway-away-awlay isway ortinedway
Epoch: 14 | Train loss: 1.040 | Val loss: 1.468 | Gen: eway anway-
ingsray-ousedw oncienceringway isway omporedway
Epoch: 15 | Train loss: 1.011 | Val loss: 1.491 | Gen: eathay aindway
onciencedway isway orksterway
Epoch: 16 | Train loss: 0.985 | Val loss: 1.466 | Gen: eway amay-
ingsray-oughtay onciencedway isway omprerestencedway
Epoch: 17 | Train loss: 0.965 | Val loss: 1.442 | Gen: eway
aingencedway onciencedway isway onmonsingway-oughtay
Epoch: 18 | Train loss: 0.940 | Val loss: 1.433 | Gen: easedway
aingenay onciecepay-outeway isway offineway
Epoch: 19 | Train loss: 0.937 | Val loss: 1.410 | Gen: eway amay
onciencedway isway offinencedway
Epoch: 20 | Train loss: 0.916 | Val loss: 1.449 | Gen: eatsay
aindway oncondingway-ietient isway offrineway
Epoch: 21 | Train loss: 0.902 | Val loss: 1.360 | Gen: ehay aingeray
ongingday-andway-oda isway offinentray
Epoch: 22 | Train loss: 0.877 | Val loss: 1.348 | Gen: eway aingeray
onciencedway isway orforificationsway
Epoch: 23 | Train loss: 0.852 | Val loss: 1.351 | Gen: ehay aingenay
onciencedway isway offinessay
Epoch: 24 | Train loss: 0.854 | Val loss: 1.358 | Gen: ehay aingsray
ongray-iecepay isway ofrsway
Epoch: 25 | Train loss: 0.830 | Val loss: 1.325 | Gen: eway aingray
ongingday isway offourshay
Epoch: 26 | Train loss: 0.814 | Val loss: 1.325 | Gen: ehay aingray
oncientrway-otedway isway orforfientway
Epoch: 27 | Train loss: 0.794 | Val loss: 1.337 | Gen: ehay aingay
oncingway-oday isway orforingway
Epoch: 28 | Train loss: 0.780 | Val loss: 1.348 | Gen: ehay aingray
ongrandingway isway offouredway
Epoch: 29 | Train loss: 0.763 | Val loss: 1.329 | Gen: ehay aingray
ongringalingway isway orfinedway
Epoch: 30 | Train loss: 0.749 | Val loss: 1.387 | Gen: ehay aingay
oncingday-aturednay isway onkersay-ondway-awla
Epoch: 31 | Train loss: 0.746 | Val loss: 1.356 | Gen: ehay aingray
ongrindedway isway orfourshay
Epoch: 32 | Train loss: 0.739 | Val loss: 1.404 | Gen: ehay aingray
oncingday-ousedway isway onkneway-iecepay
Epoch: 33 | Train loss: 0.737 | Val loss: 1.367 | Gen: ehay ay
ongingday-otecay isway onknersay
Epoch: 34 | Train loss: 0.722 | Val loss: 1.354 | Gen: ehay aingray
oncingday-otecay isway orkway-aturednay
Epoch: 35 | Train loss: 0.705 | Val loss: 1.350 | Gen: ehay aingray
ongingdray-otecay isway onknersway
Epoch: 36 | Train loss: 0.694 | Val loss: 1.393 | Gen: ehay aingray
onginglay-iecepay isway onksay-ondemnationsd
Epoch: 37 | Train loss: 0.698 | Val loss: 1.391 | Gen: ehay ay

```

ongcingday-otecay isway orkneway
Epoch: 38 | Train loss: 0.688 | Val loss: 1.383 | Gen: ehay arway
ongingday-ousedway isway onkersedway
Epoch: 39 | Train loss: 0.675 | Val loss: 1.361 | Gen: ehay ay
ongingday-odesway isway orkway-ightnay
Epoch: 40 | Train loss: 0.672 | Val loss: 1.361 | Gen: ehay ay
ongcinglay-iecepay isway onknersay
Epoch: 41 | Train loss: 0.663 | Val loss: 1.417 | Gen: ehay arway
oncivicationray isway onksay-ortificationm
Epoch: 42 | Train loss: 0.676 | Val loss: 1.396 | Gen: ehay ay
ongingrandway isway orkway-ightnay
Epoch: 43 | Train loss: 0.674 | Val loss: 1.383 | Gen: ehay anay
ongingday isway orkneway
Epoch: 44 | Train loss: 0.681 | Val loss: 1.353 | Gen: ehay aingray
ongingnationday isway orkshedway
Epoch: 45 | Train loss: 0.669 | Val loss: 1.364 | Gen: ehay ay
ongingday-indway-yba isway orkway-indway-ybay
Validation loss has not improved in 20 epochs, stopping early
Obtained lowest validation loss of: 1.3249081877561717
source:          the air conditioning is working
translated:      ehay aingray ongingday-odway isway onkuningway

```

Next, we train on the larger dataset. This experiment investigates if increasing dataset size improves model generalization on the validation set.

For a fair comparison, the number of iterations (not number of epochs) for each run should be similar. This is done in a quick and dirty way by adjusting the batch size so approximately the same number of batches is processed per epoch.

```
TEST_SENTENCE = "the air conditioning is working"
```

```

rnn_args_l = AttrDict()
args_dict = {
    "data_file_name": "pig_latin_large",
    "cuda": True,
    "nepochs": 50,
    "checkpoint_dir": "checkpoints",
    "learning_rate": 0.005,
    "lr_decay": 0.99,
    "early_stopping_patience": 10,
    "batch_size": 512,
    "hidden_size": 32,
    "encoder_type": "rnn", # options: rnn / transformer
    "decoder_type": "rnn", # options: rnn / rnn_attention /
                           transformer
    "attention_type": "", # options: additive / scaled_dot
}
rnn_args_l.update(args_dict)

print_opts(rnn_args_l)

```

```

rnn_encode_l, rnn_decoder_l, rnn_losses_l = train(rnn_args_l)

translated = translate_sentence(
    TEST_SENTENCE, rnn_encode_l, rnn_decoder_l, None, rnn_args_l
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```

=====
=====

```

Opts

```

-----
-----

```

```

data_file_name: pig_latin_large
        cuda: 1
        nepochs: 50
checkpoint_dir: checkpoints
        learning_rate: 0.005
        lr_decay: 0.99
early_stopping_patience: 10
        batch_size: 512
        hidden_size: 32
        encoder_type: rnn
        decoder_type: rnn
attention_type:

```

```

=====
=====
=====
=====

```

Data Stats

```

-----
-----

```

```

('ford', 'ordfay')
('eq', 'eqway')
('needs', 'eedsnay')

```

```
('frontline', 'ontlinefray')
('labor', 'aborlay')
Num unique word pairs: 22402
Vocabulary: dict_keys(['-', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h',
'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v',
'w', 'x', 'y', 'z', 'SOS', 'EOS'])
Vocab size: 29
```

```
=====
=====
```

Moved models to GPU!

```
Epoch:  0 | Train loss: 2.331 | Val loss: 2.083 | Gen: eray-ay-ay ay-
ay estay-ay-eray-ay-ay ay-ay-ay eray-eray-ay-ay
Epoch:  1 | Train loss: 1.893 | Val loss: 1.919 | Gen: esay-esay-ay-
esay-ay away ongay-ingay-ay-ay-ay atersay-ay-esay-ay-e oteray-edway
Epoch:  2 | Train loss: 1.720 | Val loss: 1.814 | Gen: edway away
ontay-ingay-ingay-in atersay-onsay-onsay otingray-ingay-inway
Epoch:  3 | Train loss: 1.599 | Val loss: 1.728 | Gen: edway away
ontingay-inway istay-onsay omontay-inghay
Epoch:  4 | Train loss: 1.491 | Val loss: 1.666 | Gen: eday away
ontingay-otingay-oth issay oomurationtay-intera
Epoch:  5 | Train loss: 1.408 | Val loss: 1.640 | Gen: edway away
ontinglay-inway issay omulintationcay
Epoch:  6 | Train loss: 1.355 | Val loss: 1.579 | Gen: edway away
ontay-inway-inway istay oomurationgray
Epoch:  7 | Train loss: 1.294 | Val loss: 1.546 | Gen: edway away
ontinglay istay orivintationcay
Epoch:  8 | Train loss: 1.221 | Val loss: 1.482 | Gen: edtay away
ontinglay issway orivingray
Epoch:  9 | Train loss: 1.151 | Val loss: 1.458 | Gen: edtay aitionay
ontinglay-imemay-ime istay orivinicationday
Epoch: 10 | Train loss: 1.122 | Val loss: 1.504 | Gen: eday-imetay
away ontinglay-imemay-ime issway orivingmeray
Epoch: 11 | Train loss: 1.094 | Val loss: 1.427 | Gen: etay-edway
aitionway ontinglay istay orivermentay-inedway
Epoch: 12 | Train loss: 1.037 | Val loss: 1.375 | Gen: edtay away
ontinglingway issway orkingedmay
Epoch: 13 | Train loss: 0.993 | Val loss: 1.362 | Gen: edtay away
ontinglingday istay orkingdereway
Epoch: 14 | Train loss: 0.971 | Val loss: 1.391 | Gen: eteday away
ontionday issway orkingdereway
Epoch: 15 | Train loss: 0.967 | Val loss: 1.344 | Gen: edway away
ontininglingway issway orkingedfay
Epoch: 16 | Train loss: 0.939 | Val loss: 1.404 | Gen: edtay away
ontinglay-emingnay issway ormingderedway
Epoch: 17 | Train loss: 0.922 | Val loss: 1.352 | Gen: edtay ailyway
ontinglyfay issway orikedway
Epoch: 18 | Train loss: 0.876 | Val loss: 1.255 | Gen: edtay away
ontinglay-etovingwa issway orkingdereway
Epoch: 19 | Train loss: 0.830 | Val loss: 1.280 | Gen: edtay away
ontingdray issway oridemnay
```

Epoch: 20 | Train loss: 0.818 | Val loss: 1.336 | Gen: edtay airicay
ontinginderentay issway orkingdereway
Epoch: 21 | Train loss: 0.823 | Val loss: 1.242 | Gen: edtay away
ontingdrouway isway orkingdray
Epoch: 22 | Train loss: 0.810 | Val loss: 1.333 | Gen: edtay away
ontinglay-omemnay issway orkingderay
Epoch: 23 | Train loss: 0.830 | Val loss: 1.325 | Gen: edtay away
ontingday isway orkinglay
Epoch: 24 | Train loss: 0.799 | Val loss: 1.248 | Gen: edtay away
ondingnay issway orkingday
Epoch: 25 | Train loss: 0.764 | Val loss: 1.204 | Gen: edtyway away
onondinglay-elitywa issway orkingdereway
Epoch: 26 | Train loss: 0.744 | Val loss: 1.209 | Gen: edtay away
onningnedway isway orkingdray
Epoch: 27 | Train loss: 0.724 | Val loss: 1.140 | Gen: edtay airway
ontinglidegray isway orkingday
Epoch: 28 | Train loss: 0.698 | Val loss: 1.159 | Gen: ethay away
ontinglidegray isway orkingdray
Epoch: 29 | Train loss: 0.686 | Val loss: 1.160 | Gen: ethay airway
ontinondinglay issway oridemenglay
Epoch: 30 | Train loss: 0.677 | Val loss: 1.225 | Gen: ethay airway
onininglinglytay isway orkingedway
Epoch: 31 | Train loss: 0.688 | Val loss: 1.234 | Gen: edtay airway
ondingtay issway oridementway
Epoch: 32 | Train loss: 0.705 | Val loss: 1.199 | Gen: ethay airway
onininglay-elitytay isway orindimentway
Epoch: 33 | Train loss: 0.695 | Val loss: 1.171 | Gen: ethay airway
ontingionday-iecepay issway orkingdray
Epoch: 34 | Train loss: 0.657 | Val loss: 1.227 | Gen: edtay airway
onningedway isway oridemnnway
Epoch: 35 | Train loss: 0.650 | Val loss: 1.216 | Gen: ethay arway
oninondingderay-inem isway orkingedfay
Epoch: 36 | Train loss: 0.674 | Val loss: 1.255 | Gen: edtay airway
ononondingway issway orkionedgray
Epoch: 37 | Train loss: 0.651 | Val loss: 1.143 | Gen: ehtay airway
oninidentray issway orkingeday
Epoch: 38 | Train loss: 0.604 | Val loss: 1.140 | Gen: ehtay airway
onningderobay issway orkingday
Epoch: 39 | Train loss: 0.579 | Val loss: 1.131 | Gen: ehtay airway
ondiningcteray isway orkingday
Epoch: 40 | Train loss: 0.578 | Val loss: 1.174 | Gen: ehtay airway
ondinondingway issway orkingdray
Epoch: 41 | Train loss: 0.569 | Val loss: 1.155 | Gen: ehtay airway
oningonday-inedway isway orkingday
Epoch: 42 | Train loss: 0.583 | Val loss: 1.237 | Gen: edtay airway
oninondingway-eway isway orkingway
Epoch: 43 | Train loss: 0.601 | Val loss: 1.160 | Gen: ehtay airway
oninglay-imetetay isway orkingday
Epoch: 44 | Train loss: 0.569 | Val loss: 1.109 | Gen: ehtay airway
onninondingday isway orkingdray


```

Epoch: 45 | Train loss: 0.547 | Val loss: 1.095 | Gen: ehtay airway
ondinicgancecay isway orkingday
Epoch: 46 | Train loss: 0.527 | Val loss: 1.112 | Gen: ehtay airway
ondiningctorhay isway orkingday
Epoch: 47 | Train loss: 0.516 | Val loss: 1.108 | Gen: ehtay airway
ondiningctorhay isway orkingday
Epoch: 48 | Train loss: 0.510 | Val loss: 1.118 | Gen: ehtay airway
ondinicgablway isway orkingday
Epoch: 49 | Train loss: 0.506 | Val loss: 1.104 | Gen: ehtay airway
ondioningdray isway orkingday
Obtained lowest validation loss of: 1.095056235182042
source:          the air conditioning is working
translated:      ehtay airway ondioningdray isway orkingday

```

The code below plots the training and validation losses of each model, as a function of the number of gradient descent iterations. Are there significant differences in the validation performance of each model? (see follow-up questions in handout)


```

save_loss_comparison_gru(rnn_losses_s, rnn_losses_l, rnn_args_s,
rnn_args_l, "gru")

```

Plot saved to: /content/content/csc421/a3/loss_plot_gru.pdf

<Figure size 432x288 with 0 Axes>



Question 1

Overall, the validation loss for both models are similar. However, `pig_latin_small` has a higher validation loss. Therefore, `pig_latin_large` performs better than `pig_latin_small`.

For code, refer to Step 1.

Select best performing model, and try translating different sentences by changing the variable `TEST_SENTENCE`. Identify a failure mode and briefly describe it (see follow-up questions in handout).

```

best_encoder = rnn_encode_l # Replace with rnn_encode_s or
rnn_encode_l
best_decoder = rnn_decoder_l # Replace with rnn_decoder_s or
rnn_decoder_l
best_args = rnn_args_l # Replace with rnn_args_s or rnn_args_l

TEST_SENTENCE = "the air conditioning is working"
translated = translate_sentence(
    TEST_SENTENCE, best_encoder, best_decoder, None, best_args
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

source:          the air conditioning is working
translated:      ehtay airway ondioningdray isway orkingday

```

Question 2

failed input & output pairs:

[conditioning, ondioningdray] [working, orkingday]

In the TEST_SENTENCE, 2 words failed: conditioning and working. This shows that the model cannot work with words ending in "ing" very well as some characters are missing after translation.

Question 3

Number of parameters of the LSTM encoder is $4HV DK + 4H^2$ and number of parameters of GRU encoder is $3HV DK + 3H^2$.

Part 2: Attention mechanisms

Step 1: Additive attention

In the next cell, the [additive attention](#) mechanism has been implemented for you. Please take a moment to read through it and understand what it is doing. See the assignment handouts for details.

```

class AdditiveAttention(nn.Module):
    def __init__(self, hidden_size):
        super(AdditiveAttention, self).__init__()

        self.hidden_size = hidden_size

        # A two layer fully-connected network
        # hidden_size * 2 --> hidden_size, ReLU, hidden_size --> 1
        self.attention_network = nn.Sequential(
            nn.Linear(hidden_size * 2, hidden_size),

```

```

        nn.ReLU(),
        nn.Linear(hidden_size, 1),
    )

    self.softmax = nn.Softmax(dim=1)

    def forward(self, queries, keys, values):
        """The forward pass of the additive attention mechanism.

        Arguments:
            queries: The current decoder hidden state. (batch_size x
hidden_size)
            keys: The encoder hidden states for each step of the input
sequence. (batch_size x seq_len x hidden_size)
            values: The encoder hidden states for each step of the
input sequence. (batch_size x seq_len x hidden_size)

        Returns:
            context: weighted average of the values (batch_size x 1 x
hidden_size)
            attention_weights: Normalized attention weights for each
encoder hidden state. (batch_size x seq_len x 1)

            The attention_weights must be a softmax weighting over the
seq_len annotations.
        """
        batch_size = keys.size(0)
        expanded_queries = queries.view(batch_size, -1,
self.hidden_size).expand_as(
            keys
        )
        concat_inputs = torch.cat([expanded_queries, keys], dim=2)
        unnormalized_attention = self.attention_network(concat_inputs)
        attention_weights = self.softmax(unnormalized_attention)
        context = torch.bmm(attention_weights.transpose(2, 1), values)
        return context, attention_weights

```

Step 2: RNN + additive attention

In the next cell, a modification of our RNNDecoder that makes use of an additive attention mechanism as been implemented for your. Please take a moment to read through it and understand what it is doing. See the assignment handouts for details.

```

class RNNAttentionDecoder(nn.Module):
    def __init__(self, vocab_size, hidden_size,
attention_type="scaled_dot"):
        super(RNNAttentionDecoder, self).__init__()
        self.vocab_size = vocab_size
        self.hidden_size = hidden_size

```

```

        self.embedding = nn.Embedding(vocab_size, hidden_size)

        self.rnn = MyGRUCell(input_size=hidden_size * 2,
                               hidden_size=hidden_size)
        if attention_type == "additive":
            self.attention =
AdditiveAttention(hidden_size=hidden_size)
        elif attention_type == "scaled_dot":
            self.attention =
ScaledDotAttention(hidden_size=hidden_size)

        self.out = nn.Linear(hidden_size, vocab_size)

    def forward(self, inputs, annotations, hidden_init):
        """Forward pass of the attention-based decoder RNN.

        Arguments:
            inputs: Input token indexes across a batch for all the
time step. (batch_size x decoder_seq_len)
            annotations: The encoder hidden states for each step of
the input.
                                sequence. (batch_size x seq_len x
hidden_size)
            hidden_init: The final hidden states from the encoder,
across a batch. (batch_size x hidden_size)

        Returns:
            output: Un-normalized scores for each token in the
vocabulary, across a batch for all the decoding time steps.
(batch_size x decoder_seq_len x vocab_size)
            attentions: The stacked attention weights applied to the
encoder annotations (batch_size x encoder_seq_len x decoder_seq_len)
        """

        batch_size, seq_len = inputs.size()
        embed = self.embedding(inputs) # batch_size x seq_len x
hidden_size

        hiddens = []
        attentions = []
        h_prev = hidden_init

        for i in range(seq_len):
            embed_current = embed[
                :, i, :
            ] # Get the current time step, across the whole batch
            context, attention_weights = self.attention(
                h_prev, annotations, annotations
            ) # batch_size x 1 x hidden_size

```

```

        embed_and_context = torch.cat(
            [embed_current, context.squeeze(1)], dim=1
        ) # batch_size x (2*hidden_size)
        h_prev = self.rnn(embed_and_context, h_prev) # batch_size
x hidden_size

        hiddens.append(h_prev)
        attentions.append(attention_weights)

        hiddens = torch.stack(hiddens, dim=1) # batch_size x seq_len
x hidden_size
        attentions = torch.cat(attentions, dim=2) # batch_size x
seq_len x seq_len

        output = self.out(hiddens) # batch_size x seq_len x
vocab_size
        return output, attentions

```

Step 3: Training and analysis (with additive attention)

Now, run the following cell to train our recurrent encoder-decoder model with additive attention. How does it perform compared to the recurrent encoder-decoder model without attention?

```
TEST_SENTENCE = "the air conditioning is working"
```

```

rnn_attn_args = AttrDict()
args_dict = {
    "data_file_name": "pig_latin_small",
    "cuda": True,
    "nepochs": 50,
    "checkpoint_dir": "checkpoints",
    "learning_rate": 0.005,
    "lr_decay": 0.99,
    "early_stopping_patience": 10,
    "batch_size": 64,
    "hidden_size": 64,
    "encoder_type": "rnn", # options: rnn / transformer
    "decoder_type": "rnn_attention", # options: rnn / rnn_attention /
transformer
    "attention_type": "additive", # options: additive / scaled_dot
}
rnn_attn_args.update(args_dict)

print_opts(rnn_attn_args)
rnn_attn_encoder, rnn_attn_decoder, rnn_attn_losses =
train(rnn_attn_args)

translated = translate_sentence(
    TEST_SENTENCE, rnn_attn_encoder, rnn_attn_decoder, None,

```

```
rnn_attn_args
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))
```

```
=====
=====
```

Opts

```
-----
-----
```

data_file_name: pig_latin_small

cuda: 1

nepochs: 50

checkpoint_dir: checkpoints

learning_rate: 0.005

lr_decay: 0.99

early_stopping_patience: 10

batch_size: 64

hidden_size: 64

encoder_type: rnn

decoder_type: rnn_attention

attention_type: additive

```
=====
=====
=====
=====
```

Data Stats

```
-----
-----
```

```
('piqued', 'iquedpay')
('learnt', 'earntlay')
('acquainted', 'acquaintedway')
('terms', 'ermstay')
('steady', 'eadystay')
Num unique word pairs: 3198
```

Vocabulary: dict_keys(['-', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z', 'SOS', 'EOS'])

Vocab size: 29

=====

Moved models to GPU!

Epoch: 0 | Train loss: 1.956 | Val loss: 1.843 | Gen: otay-odgay-
oday-oday illlay-illlay-illlay intinday-intinday-in issay-ilssay-
ilssay ingway-ingway-ingway

Epoch: 1 | Train loss: 1.429 | Val loss: 1.589 | Gen: ethay arway
ondingstingstay-ings isssay-isssay-isssay orongway-ingway-ingw

Epoch: 2 | Train loss: 1.135 | Val loss: 1.416 | Gen: athehtay arway
ondincationday issay oongringray

Epoch: 3 | Train loss: 0.958 | Val loss: 1.350 | Gen: elfay-uethay
away ondincingnincingnay isway ovedingway

Epoch: 4 | Train loss: 0.767 | Val loss: 1.092 | Gen: ethay-uenthay-
ehthay arirway onditingnay isway orfingngnay

Epoch: 5 | Train loss: 0.631 | Val loss: 1.028 | Gen: ethay ariway
ondincingingingingnay isway orfingnay

Epoch: 6 | Train loss: 0.544 | Val loss: 1.125 | Gen: ethay airway
omitiondingsay-ition isway orfingsay-ingsay-ing

Epoch: 7 | Train loss: 0.482 | Val loss: 0.934 | Gen: eteway arway
onditingncay isway orfingnay

Epoch: 8 | Train loss: 0.398 | Val loss: 0.837 | Gen: ethay airway
onditingingcay isway orfingningway

Epoch: 9 | Train loss: 0.366 | Val loss: 0.930 | Gen: ethay away
ondcay isway orkingnay

Epoch: 10 | Train loss: 0.317 | Val loss: 0.862 | Gen: eway airway
onditionitionday isway orfingway

Epoch: 11 | Train loss: 0.240 | Val loss: 0.620 | Gen: ethay airway
onditionday isway orkingnay

Epoch: 12 | Train loss: 0.179 | Val loss: 0.540 | Gen: ethtay airway
onditiningcay isway orkingway

Epoch: 13 | Train loss: 0.168 | Val loss: 0.664 | Gen: ehtay away
onditingway iway orfingway

Epoch: 14 | Train loss: 0.171 | Val loss: 0.594 | Gen: ethay airway
onditiongcay isway orfingway

Epoch: 15 | Train loss: 0.150 | Val loss: 0.583 | Gen: ethay airway
onditionminingcay isway orkingway

Epoch: 16 | Train loss: 0.117 | Val loss: 0.443 | Gen: ethay airway
onditiongcay isway orkingway

Epoch: 17 | Train loss: 0.080 | Val loss: 0.402 | Gen: ethay airway
onditiondcay isway orkingway

Epoch: 18 | Train loss: 0.062 | Val loss: 0.390 | Gen: ehthay airway
onditionicaningway isway orkingway

Epoch: 19 | Train loss: 0.046 | Val loss: 0.422 | Gen: ethay airway
onditionday isway orkingway

Epoch: 20 | Train loss: 0.041 | Val loss: 0.358 | Gen: ethay airway
onditiongway way orkingway

```

Epoch: 21 | Train loss: 0.035 | Val loss: 0.404 | Gen: ethay airway
onditiongcay isway orkingway
Epoch: 22 | Train loss: 0.037 | Val loss: 0.361 | Gen: ethay airway
onditioncay isway orkingway
Epoch: 23 | Train loss: 0.024 | Val loss: 0.373 | Gen: ethay airway
onditioningcay isway orkingway
Epoch: 24 | Train loss: 0.016 | Val loss: 0.326 | Gen: ethay airway
onditionicningcay isway orkingway
Epoch: 25 | Train loss: 0.019 | Val loss: 0.433 | Gen: ethay airway
onditioncay isway orkingway
Epoch: 26 | Train loss: 0.038 | Val loss: 0.447 | Gen: ethay airway
onditiningcay isway orkingway
Epoch: 27 | Train loss: 0.064 | Val loss: 0.664 | Gen: ethay airway
onditingcay isway orkingway
Epoch: 28 | Train loss: 0.290 | Val loss: 0.904 | Gen: ekthay-outhay
airway onditicininingcay isway onkingway-ingway
Epoch: 29 | Train loss: 0.264 | Val loss: 0.705 | Gen: ethay arway
onitingcay issay orfingugway
Epoch: 30 | Train loss: 0.186 | Val loss: 0.554 | Gen: ethay airway
onitiningcay isway orkingingway
Epoch: 31 | Train loss: 0.099 | Val loss: 0.444 | Gen: ethay airway
onditiningcay isway orfingway
Epoch: 32 | Train loss: 0.065 | Val loss: 0.492 | Gen: ethay airway
onditioncay isway orfingway
Epoch: 33 | Train loss: 0.067 | Val loss: 0.421 | Gen: ethay airway
onditioncay isway orkingway
Epoch: 34 | Train loss: 0.036 | Val loss: 0.389 | Gen: ethay airway
onditioncay isway orkingway
Validation loss has not improved in 10 epochs, stopping early
Obtained lowest validation loss of: 0.32645468270549405
source:          the air conditioning is working
translated:      ethay airway onditiningcay isway orkingway

```

```

TEST_SENTENCE = "the air conditioning is working"
translated = translate_sentence(
    TEST_SENTENCE, rnn_attn_encoder, rnn_attn_decoder, None,
    rnn_attn_args
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```

source:          the air conditioning is working
translated:      ethay airway onditiningcay isway orkingway

```

Step 4: Implement scaled dot-product attention

In the next cell, you will implement the [scaled dot-product attention](#) mechanism. See the assignment handouts for details.

```

class ScaledDotAttention(nn.Module):
    def __init__(self, hidden_size):

```



```

super(ScaledDotAttention, self).__init__()

self.hidden_size = hidden_size

self.Q = nn.Linear(hidden_size, hidden_size)
self.K = nn.Linear(hidden_size, hidden_size)
self.V = nn.Linear(hidden_size, hidden_size)
self.softmax = nn.Softmax(dim=1)
self.scaling_factor = torch.rsqrt(
    torch.tensor(self.hidden_size, dtype=torch.float)
)

def forward(self, queries, keys, values):
    """The forward pass of the scaled dot attention mechanism.

    Arguments:
        queries: The current decoder hidden state, 2D or 3D
        tensor. (batch_size x (k) x hidden_size)
        keys: The encoder hidden states for each step of the input
        sequence. (batch_size x seq_len x hidden_size)
        values: The encoder hidden states for each step of the
        input sequence. (batch_size x seq_len x hidden_size)

    Returns:
        context: weighted average of the values (batch_size x k x
        hidden_size)
        attention_weights: Normalized attention weights for each
        encoder hidden state. (batch_size x seq_len x k)

    The output must be a softmax weighting over the seq_len
    annotations.
    """

    # -----
    # FILL THIS IN
    # -----

    if len(queries.size()) == 2:
        queries = queries.unsqueeze(1)
        queries = queries.permute(0, 2, 1) # (batch_size, k,
hidden_size)

        q = self.Q(queries).permute(0, 2, 1) # (batch_size,
hidden_size, k)
        k = self.K(keys) # (batch_size, seq_len, hidden_size)
        v = self.V(values)

        # (batch_size x seq_len x k) = (batch_size x seq_len x
hidden_size) @ (batch_size x hidden_size x k)

```

```

        unnormalized_attention = torch.bmm(k * self.scaling_factor, q)
        attention_weights = self.softmax(unnormalized_attention)

        # (batch_size x k x hidden_size) = (batch_size x k x seq_len)
        @ (batch_size x seq_len x hidden_size)
        context = torch.bmm(attention_weights.permute(0, 2, 1), v)
        return context, attention_weights

```

Step 5: Implement causal dot-product Attention

Now, implement the casual scaled dot-product attention mechanism. It will be very similar to your implementation for ScaledDotAttention. The additional step is to mask out the attention to future timesteps so this attention mechanism can be used in a decoder. See the assignment handouts for details.

```

class CausalScaledDotAttention(nn.Module):
    def __init__(self, hidden_size):
        super(CausalScaledDotAttention, self).__init__()

        self.hidden_size = hidden_size
        self.neg_inf = torch.tensor(-1e7)

        self.Q = nn.Linear(hidden_size, hidden_size)
        self.K = nn.Linear(hidden_size, hidden_size)
        self.V = nn.Linear(hidden_size, hidden_size)
        self.softmax = nn.Softmax(dim=1)
        self.scaling_factor = torch.rsqrt(
            torch.tensor(self.hidden_size, dtype=torch.float)
        )

    def forward(self, queries, keys, values):
        """The forward pass of the scaled dot attention mechanism.

        Arguments:
            queries: The current decoder hidden state, 2D or 3D
            tensor. (batch_size x (k) x hidden_size)
            keys: The encoder hidden states for each step of the input
            sequence. (batch_size x seq_len x hidden_size)
            values: The encoder hidden states for each step of the
            input sequence. (batch_size x seq_len x hidden_size)

        Returns:
            context: weighted average of the values (batch_size x k x
            hidden_size)
            attention_weights: Normalized attention weights for each
            encoder hidden state. (batch_size x seq_len x k)

            The output must be a softmax weighting over the seq_len
            annotations.
        """

```

```

# -----
# FILL THIS IN
# -----

if len(queries.size()) == 2:
    queries = queries.unsqueeze(1)
    queries = queries.permute(0, 2, 1) # (batch_size, k,
hidden_size)

    q = self.Q(queries).permute(0, 2, 1) # (batch_size,
hidden_size, k)
    k = self.K(keys) # (batch_size, seq_len, hidden_size)
    v = self.V(values)

    # (batch_size x seq_len x k) = (batch_size x seq_len x
hidden_size) @ (batch_size x hidden_size x k)
    unnormalized_attention = torch.bmm(k * self.scaling_factor, q)

    mask = torch.ones((unnormalized_attention.size()[0],
unnormalized_attention.size()[-1], unnormalized_attention.size()[-2]),
device=torch.device('cuda:0'))
    mask = torch.tril(mask * self.neg_inf)

    attention_weights =
self.softmax(torch.bmm(unnormalized_attention, mask))
    context = torch.bmm(attention_weights.permute(0, 2, 1), v)
    return context, attention_weights

```

Step 6: Attention encoder and decoder

The following cells provide an implementation of an encoder and decoder that use a single ScaledDotAttention block. Please read through them to understand what they are doing.

```

class AttentionEncoder(nn.Module):
    def __init__(self, vocab_size, hidden_size, opts):
        super(AttentionEncoder, self).__init__()

        self.vocab_size = vocab_size
        self.hidden_size = hidden_size
        self.opts = opts

        self.embedding = nn.Embedding(vocab_size, hidden_size)

        self.self_attention = ScaledDotAttention(
            hidden_size=hidden_size,
        )

        self.attention_mlp = nn.Sequential(

```

```

        nn.Linear(hidden_size, hidden_size),
        nn.ReLU(),
    )

    def forward(self, inputs):
        """Forward pass of the encoder scaled dot attention.

        Arguments:
            inputs: Input token indexes across a batch for all time
            steps in the sequence. (batch_size x seq_len)

        Returns:
            annotations: The hidden states computed at each step of
            the input sequence. (batch_size x seq_len x hidden_size)
            None: Used to conform to standard encoder return
            signature.
        """
        batch_size, seq_len = inputs.size()

        encoded = self.embedding(inputs) # batch_size x seq_len x
        hidden_size

        annotations = encoded
        new_annotations, self_attention_weights = self.self_attention(
            annotations, annotations, annotations
        ) # batch_size x seq_len x hidden_size
        residual_annotations = annotations + new_annotations
        new_annotations = self.attention_mlp(residual_annotations)
        annotations = residual_annotations + new_annotations

        return annotations, None

class AttentionDecoder(nn.Module):
    def __init__(self, vocab_size, hidden_size):
        super(AttentionDecoder, self).__init__()
        self.vocab_size = vocab_size
        self.hidden_size = hidden_size

        self.embedding = nn.Embedding(vocab_size, hidden_size)

        self.self_attention = CausalScaledDotAttention(
            hidden_size=hidden_size,
        )

        self.decoder_attention = ScaledDotAttention(
            hidden_size=hidden_size,
        )

        self.attention_mlp = nn.Sequential(
            nn.Linear(hidden_size, hidden_size),

```

```

        nn.ReLU(),
    )

    self.out = nn.Linear(hidden_size, vocab_size)

def forward(self, inputs, annotations, hidden_init):
    """Forward pass of the attention-based decoder RNN.

    Arguments:
        inputs: Input token indexes across a batch for all the
time step. (batch_size x decoder_seq_len)
        annotations: The encoder hidden states for each step of
the input.
                    sequence. (batch_size x seq_len x
hidden_size)
        hidden_init: Not used in the transformer decoder
    Returns:
        output: Un-normalized scores for each token in the
vocabulary, across a batch for all the decoding time steps.
(batch_size x decoder_seq_len x vocab_size)
        attentions: The stacked attention weights applied to the
encoder annotations (batch_size x encoder_seq_len x decoder_seq_len)
    """

    batch_size, seq_len = inputs.size()
    embed = self.embedding(inputs) # batch_size x seq_len x
hidden_size

    encoder_attention_weights_list = []
    self_attention_weights_list = []
    contexts = embed
    new_contexts, self_attention_weights = self.self_attention(
        contexts, contexts, contexts
    ) # batch_size x seq_len x hidden_size
    residual_contexts = contexts + new_contexts
    new_contexts, encoder_attention_weights =
self.decoder_attention(
        residual_contexts, annotations, annotations
    ) # batch_size x seq_len x hidden_size
    residual_contexts = residual_contexts + new_contexts
    new_contexts = self.attention_mlp(residual_contexts)
    contexts = residual_contexts + new_contexts

    encoder_attention_weights_list.append(encoder_attention_weights)
    self_attention_weights_list.append(self_attention_weights)

    output = self.out(contexts)

```

```

        encoder_attention_weights =
torch.stack(encoder_attention_weights_list)
        self_attention_weights =
torch.stack(self_attention_weights_list)

    return output, (encoder_attention_weights,
self_attention_weights)

```

Step 7: Training and analysis (single scaled dot-product attention block)

Now, train the following model, with an encoder and decoder each composed a single ScaledDotAttention block.

```
TEST_SENTENCE = "the air conditioning is working"
```

```

attention_args_s = AttrDict()
args_dict = {
    "data_file_name": "pig_latin_small",
    "cuda": True,
    "nepochs": 100,
    "checkpoint_dir": "checkpoints",
    "learning_rate": 5e-4,
    "early_stopping_patience": 100,
    "lr_decay": 0.99,
    "batch_size": 64,
    "hidden_size": 32,
    "encoder_type": "attention",
    "decoder_type": "attention", # options: rnn / rnn_attention /
attention / transformer
}
attention_args_s.update(args_dict)
print_opts(attention_args_s)

attention_encoder_s, attention_decoder_s, attention_losses_s =
train(attention_args_s)

translated = translate_sentence(
    TEST_SENTENCE, attention_encoder_s, attention_decoder_s, None,
attention_args_s
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```
=====
=====
```

Opts

```
-----
-----
```

data_file_name: pig_latin_small

cuda: 1
nepochs: 100
checkpoint_dir: checkpoints
learning_rate: 0.0005
early_stopping_patience: 100
lr_decay: 0.99
batch_size: 64
hidden_size: 32
encoder_type: attention
decoder_type: attention

```
=====
=====
=====
=====
```

Data Stats

```
-----
-----
('fully', 'ullyfay')
('determining', 'eterminingday')
('confounded', 'onfoundedcay')
('attempting', 'attemptingway')
('darling', 'arlingday')
Num unique word pairs: 3198
Vocabulary: dict_keys(['-', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h',
'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v',
'w', 'x', 'y', 'z', 'SOS', 'EOS'])
Vocab size: 29
```

```
=====
=====
```

Moved models to GPU!

Epoch: 0 | Train loss: 3.047 | Val loss: 2.526 | Gen: ay ay ay ay ay

Epoch: 1 | Train loss: 2.329 | Val loss: 2.239 | Gen:
ayayayayayayayayayay ay ay ay ay

Epoch: 2 | Train loss: 2.123 | Val loss: 2.099 | Gen:
ayayayayayayayayayay ay iny ay ongayy

Epoch: 3 | Train loss: 1.975 | Val loss: 2.007 | Gen: eaay ay

onsiniiniiniiniiniin ay ongayy
Epoch: 4 | Train loss: 1.877 | Val loss: 1.924 | Gen: ay ay
onsnnnnnnnnnnnnnnnnnnnn ay ongmoE0Sooooononooiyoo
Epoch: 5 | Train loss: 1.800 | Val loss: 1.871 | Gen: ay ay
onsnnnnnnnnnnnnnnnnnnnn ay ongrwayoogoE0SE0Synnmoo
Epoch: 6 | Train loss: 1.751 | Val loss: 1.855 | Gen: ay ay
onsnnnnnnnnnnnnnnnnnnnn ay ongongongongongongon
Epoch: 7 | Train loss: 1.719 | Val loss: 1.811 | Gen: ay ay
onsnnnnnnnnnnnnnnnnnnnn ay oorray
Epoch: 8 | Train loss: 1.653 | Val loss: 1.762 | Gen: ay ay
ongnnnnnnnnnnnnnnnnnnnn issssssssssssssssssss oorray
Epoch: 9 | Train loss: 1.615 | Val loss: 1.712 | Gen: ay ay
ongingingingingingon issssssssssssssssssss oorray
Epoch: 10 | Train loss: 1.565 | Val loss: 1.670 | Gen: ay ay
ongioggioggioggiogogo issssssssssssssssssss oorray
Epoch: 11 | Train loss: 1.526 | Val loss: 1.663 | Gen: ay ay
ongingongongongongon issssssssssssssssssss ooorwE0Sy
Epoch: 12 | Train loss: 1.504 | Val loss: 1.642 | Gen: ay
arararararararararar ongingongongongongon issssssssssssssssssss
oornE0Sy
Epoch: 13 | Train loss: 1.475 | Val loss: 1.634 | Gen: ay
arararararararararar ongingongongongongon issssssssssssssssssss oorray
Epoch: 14 | Train loss: 1.453 | Val loss: 1.616 | Gen: ay
arararararararararar onginE0Sinnnnnnnnnnnnnn issssssssssssssssssss
iraryy
Epoch: 15 | Train loss: 1.455 | Val loss: 1.604 | Gen: ay
arararararararararar ongingingingingingon isisissississississ inayyy
Epoch: 16 | Train loss: 1.419 | Val loss: 1.576 | Gen: ay
arararararararararar ongingingingingingon isisissississississ inayyy
Epoch: 17 | Train loss: 1.416 | Val loss: 1.567 | Gen: ay
arararararararararar ongingingingingingon iswsy inayyy
Epoch: 18 | Train loss: 1.397 | Val loss: 1.554 | Gen: ay
arararararararararar ongingingingingingon iswsy inayyy
Epoch: 19 | Train loss: 1.406 | Val loss: 1.582 | Gen: ay
arararararararararar ongingingingingingon issssssssssssssssssss irayay
Epoch: 20 | Train loss: 1.387 | Val loss: 1.554 | Gen: ay
arwarwarwarwarwar ongwngiay isisissississississ irgnyyy
Epoch: 21 | Train loss: 1.373 | Val loss: 1.542 | Gen: ay
arararararararararar ingwnyiny isisissississississ inayyy
Epoch: 22 | Train loss: 1.358 | Val loss: 1.525 | Gen: ay
arararararararararar ongingiay isisissississississ inayyy
Epoch: 23 | Train loss: 1.339 | Val loss: 1.504 | Gen: ay
arararararararararar ongwngnyngwnywny isisissississississ inayyy
Epoch: 24 | Train loss: 1.333 | Val loss: 1.531 | Gen: ay
arwarwarwarwarwar ingwnyngwnywny isisissississississ inayyy
Epoch: 25 | Train loss: 1.340 | Val loss: 1.510 | Gen: ay
arwarwarwarwarwar ingiay isisissississississ irayay
Epoch: 26 | Train loss: 1.314 | Val loss: 1.484 | Gen: ay
arwarwarwarwarwar ingiay isisissississississ irrr
Epoch: 27 | Train loss: 1.285 | Val loss: 1.470 | Gen: ay

arararararararararar ingingingingingingin isisisisisisisisis irayay
Epoch: 28 | Train loss: 1.279 | Val loss: 1.476 | Gen: ay
arwarwarwarwarwar ingiay isisisisisisisisis irayay
Epoch: 29 | Train loss: 1.262 | Val loss: 1.472 | Gen: ay
arwarwarwarwarwar indiay isisisisisisisisis irayay
Epoch: 30 | Train loss: 1.266 | Val loss: 1.485 | Gen: ay ay
inginginginyinginyww iswsy irayay
Epoch: 31 | Train loss: 1.265 | Val loss: 1.465 | Gen: ay
arwarwarwarwarwar ingwwywywywny iswsy irayay
Epoch: 32 | Train loss: 1.250 | Val loss: 1.472 | Gen: ay
arwarwarwarwarwar ingingingingingingww isisisisisisisisis irayay
Epoch: 33 | Train loss: 1.248 | Val loss: 1.450 | Gen: ay
arwarwarwarwarwar ingyagE0SngwngE0Sngy isisisisisisisisis irayay
Epoch: 34 | Train loss: 1.243 | Val loss: 1.461 | Gen: ay ay ingway
isisisisisisisisis irayay
Epoch: 35 | Train loss: 1.245 | Val loss: 1.492 | Gen: ay ay ingway
isisisisisisisisis orayay
Epoch: 36 | Train loss: 1.268 | Val loss: 1.482 | Gen: ay ay ingway
iswsy irayay
Epoch: 37 | Train loss: 1.231 | Val loss: 1.472 | Gen: ay
arwarwarwarwarwar ingway isisisisisisisisis irayay
Epoch: 38 | Train loss: 1.237 | Val loss: 1.482 | Gen: ay ay ingway
isisisisisisisisis irayay
Epoch: 39 | Train loss: 1.229 | Val loss: 1.469 | Gen: ay
arwarwarwarwarwar ingway isssssssssssssssssss irayay
Epoch: 40 | Train loss: 1.231 | Val loss: 1.448 | Gen: ay
arwarwarwarwarwar ingway isisisisisisisisis irayay
Epoch: 41 | Train loss: 1.238 | Val loss: 1.460 | Gen: ay ay ingway
iswsy irayay
Epoch: 42 | Train loss: 1.238 | Val loss: 1.472 | Gen: ay ay ongway
iswsy iraray
Epoch: 43 | Train loss: 1.214 | Val loss: 1.454 | Gen: ay
arwarwarwarwarwar ongway iswsy ingr
Epoch: 44 | Train loss: 1.206 | Val loss: 1.449 | Gen: ay
arwarwarwarwarwar ongway iswsy orgway
Epoch: 45 | Train loss: 1.246 | Val loss: 1.447 | Gen: ethayE0Sety ay
ongway isway owgway
Epoch: 46 | Train loss: 1.197 | Val loss: 1.438 | Gen: etayyE0Sey
arwarwarwarwarwar ongway isway owgway
Epoch: 47 | Train loss: 1.177 | Val loss: 1.434 | Gen: ethy
aywaywayE0SarE0SayE0SarE0Sar ongway isisisisisisisisis owgway
Epoch: 48 | Train loss: 1.187 | Val loss: 1.432 | Gen: etayyE0Sey
arwarwarwarwarwar ongway isisisisisisisisis owgway
Epoch: 49 | Train loss: 1.187 | Val loss: 1.436 | Gen: ethay
arwarwarwarwarwar ongway isisisisisisisisis orarE0Sy
Epoch: 50 | Train loss: 1.242 | Val loss: 1.476 | Gen: ethayy arE0Say
ongway isisisisisisisisis orarE0Sy
Epoch: 51 | Train loss: 1.273 | Val loss: 1.483 | Gen: ethayy
arwarwarwarwarwar ongway isssssssssssssssssss orayE0Sy
Epoch: 52 | Train loss: 1.269 | Val loss: 1.470 | Gen: ethayy aaE0Say

ingwayy isisisisisisisisisis iwaaynyy
Epoch: 53 | Train loss: 1.254 | Val loss: 1.464 | Gen:
ethwhwhwhwhwhwhwhwhh aaE0Say ingwayy isssssssssssssssssss iwaaynyy
Epoch: 54 | Train loss: 1.252 | Val loss: 1.466 | Gen:
ethwhwhwhwhwhwhwhwhh aaE0Say ongwaayE0Sayyy isssssssssssssssssss
inayayy
Epoch: 55 | Train loss: 1.246 | Val loss: 1.461 | Gen:
ethwhwhwhwhwhwhwhwhh iwarwarwwawaawarwar inggway isssssssssssssssssss
owgway
Epoch: 56 | Train loss: 1.223 | Val loss: 1.458 | Gen:
ethwhwhwhwhwhwhwhwhh arwarwarwarwarwar inggway isssssssssssssssssss
owgway
Epoch: 57 | Train loss: 1.220 | Val loss: 1.463 | Gen:
ethwhwhwhwhwhwhwhwhh arwarwarwarwarwar inggway isssssssssssssssssss
owgway
Epoch: 58 | Train loss: 1.221 | Val loss: 1.461 | Gen:
ethwhwhwhwhwhwhwhwhh arwarwarwarwarwar inggway isssssssssssssssssss
owgway
Epoch: 59 | Train loss: 1.214 | Val loss: 1.451 | Gen:
ethwhwhwhwhwhwhwhwhh arwarwarwarwarwar inggway isssssssssssssssssss
owgway
Epoch: 60 | Train loss: 1.203 | Val loss: 1.457 | Gen:
ethwhwhwhwhwhwhwhwhh arwarwarwarwarwar inggway isssssssssssssssssss
iwaray
Epoch: 61 | Train loss: 1.202 | Val loss: 1.448 | Gen:
ethwhwhwhwhwhwhwhwhh arwarwarwarwarwar inggway isssssssssssssssssss
owgway
Epoch: 62 | Train loss: 1.192 | Val loss: 1.444 | Gen:
ethwhwhwhwhwhwhwhwhh ayE0Say inggway isssssssssssssssssss owgway
Epoch: 63 | Train loss: 1.187 | Val loss: 1.434 | Gen:
ethwhwhwhwhwhwhwhwhh arwarwarwarwarwar inggway isssssssssssssssssss
owgway
Epoch: 64 | Train loss: 1.166 | Val loss: 1.410 | Gen: ethwy
arwirrrrrirrrrrrrrrr inggway isisisisisisisisisis owgway
Epoch: 65 | Train loss: 1.146 | Val loss: 1.388 | Gen: ethy
arE0SarE0SarE0SarE0SarE0SarE0Sar inggway isway owgway
Epoch: 66 | Train loss: 1.134 | Val loss: 1.385 | Gen: ethy
arrrrrrrrrrrrrrrrrrrr inggway isway owgway
Epoch: 67 | Train loss: 1.135 | Val loss: 1.378 | Gen: ethy ay
inggway isway owgway
Epoch: 68 | Train loss: 1.140 | Val loss: 1.390 | Gen: ethy ay
inggway isway owgway
Epoch: 69 | Train loss: 1.128 | Val loss: 1.391 | Gen: ethy ay
inggway isway owgway
Epoch: 70 | Train loss: 1.166 | Val loss: 1.417 | Gen: ethy ay
inggway isway owgway
Epoch: 71 | Train loss: 1.163 | Val loss: 1.413 | Gen: ethy
arrrrrrrrrrrrrrrrrrrr inggway isway owgway
Epoch: 72 | Train loss: 1.162 | Val loss: 1.422 | Gen: ethy ay
inggway isway owgry

Epoch: 73 | Train loss: 1.155 | Val loss: 1.408 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingway isway owgry
Epoch: 74 | Train loss: 1.134 | Val loss: 1.411 | Gen: ethy ay ingayy
isway ingry
Epoch: 75 | Train loss: 1.145 | Val loss: 1.418 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingayy isway ingry
Epoch: 76 | Train loss: 1.144 | Val loss: 1.407 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingayy isway ingry
Epoch: 77 | Train loss: 1.123 | Val loss: 1.384 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingayy isway ingry
Epoch: 78 | Train loss: 1.103 | Val loss: 1.373 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingwayE0Sngyy isssssssssssssssssss ingry
Epoch: 79 | Train loss: 1.110 | Val loss: 1.405 | Gen: ethy ay ingayy
isssssssssssssssssss ingry
Epoch: 80 | Train loss: 1.133 | Val loss: 1.372 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingwayE0Sngyy isssssssssssssssssss iogwwy
Epoch: 81 | Train loss: 1.101 | Val loss: 1.374 | Gen: ethy
arE0Saraaraaraarywyw ingwayE0Sngyy isssssssssssssssssss ingwagy
Epoch: 82 | Train loss: 1.092 | Val loss: 1.361 | Gen: ethy
arE0Saraaraaraarywyw ingwayE0Sngyy isway ingwagy
Epoch: 83 | Train loss: 1.083 | Val loss: 1.363 | Gen: ethy
arE0Saraaraaraarywyw ingwayE0Sngyyyy isway ingwaay
Epoch: 84 | Train loss: 1.077 | Val loss: 1.361 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingwayE0Sngyyyy isway ingwaay
Epoch: 85 | Train loss: 1.078 | Val loss: 1.360 | Gen: ethy
arE0Saraaraaraarywyw ingwayE0Sngyyyy isway ingwagy
Epoch: 86 | Train loss: 1.071 | Val loss: 1.350 | Gen: ethy
arE0Saraaraaraarywyw ingwayE0Sngyyyy isway ingwagy
Epoch: 87 | Train loss: 1.067 | Val loss: 1.347 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingwayE0Sngyy isway ingrwgy
Epoch: 88 | Train loss: 1.069 | Val loss: 1.355 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr ingwayE0Sngyyyy isway ingrnay
Epoch: 89 | Train loss: 1.066 | Val loss: 1.351 | Gen: ethy
arrrrrrrrrrrrrrrrrrrrrr indddy isway ingrnay
Epoch: 90 | Train loss: 1.066 | Val loss: 1.349 | Gen: ay
aywaywayiyywywywww ingwayE0Sngy isway ingray
Epoch: 91 | Train loss: 1.061 | Val loss: 1.342 | Gen: ay
ayE0Sayaiywywywywww ingwayE0Sngy isway ingray
Epoch: 92 | Train loss: 1.057 | Val loss: 1.338 | Gen: ay
ayE0Sayaayaayaayyywyw ingwayE0Sngy isway ingray
Epoch: 93 | Train loss: 1.057 | Val loss: 1.331 | Gen: ay
arrrrrrrrrrrrrrrrrrrrrr ingwayE0Sngy isway ingray
Epoch: 94 | Train loss: 1.066 | Val loss: 1.337 | Gen: ethy
arE0Saraaraayaayyywyw ongwayE0Sngy isway ingray
Epoch: 95 | Train loss: 1.051 | Val loss: 1.324 | Gen: ethy
arE0Saraayaayaayyywyw ingwayE0Sngy isway ingray
Epoch: 96 | Train loss: 1.038 | Val loss: 1.323 | Gen: ethy
arE0Saraaraaraarywyw ingway isway ingray
Epoch: 97 | Train loss: 1.037 | Val loss: 1.329 | Gen: ethy
arE0Saraaraaraarywyw ingway isway ingray

```
Epoch: 98 | Train loss: 1.056 | Val loss: 1.326 | Gen: ethy  
arrrrrrrrrrrrrrrrrrrrr ingway isway ingrayer  
Epoch: 99 | Train loss: 1.054 | Val loss: 1.358 | Gen: ethy  
arE0Saraaraaraarywyw ingway isway ingrayer  
Obtained lowest validation loss of: 1.3230563144426088  
source:          the air conditioning is working  
translated:      ethy arE0Saraaraaraarywyw ingway isway ingrayer
```

Step 8: Transformer encoder and decoder

The following cells provide an implementation of the transformer encoder and decoder that use your `ScaledDotAttention` and `CausalScaledDotAttention`. Please read through them to understand what they are doing.

```

class TransformerEncoder(nn.Module):
    def __init__(self, vocab_size, hidden_size, num_layers, opts):
        super(TransformerEncoder, self).__init__()

        self.vocab_size = vocab_size
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        self.opts = opts

        self.embedding = nn.Embedding(vocab_size, hidden_size)

        self.self_attentions = nn.ModuleList(
            [
                ScaledDotAttention(
                    hidden_size=hidden_size,
                )
                for i in range(self.num_layers)
            ]
        )
        self.attention_mlp = nn.ModuleList(
            [
                nn.Sequential(
                    nn.Linear(hidden_size, hidden_size),
                    nn.ReLU(),
                )
                for i in range(self.num_layers)
            ]
        )

        self.positional_encodings = self.create_positional_encodings()

    def forward(self, inputs):
        """Forward pass of the encoder RNN.

        Arguments:
            inputs: Input token indexes across a batch for all time

```

steps in the sequence. (batch_size x seq_len)

Returns:

annotations: The hidden states computed at each step of the input sequence. (batch_size x seq_len x hidden_size)

signature.
None: Used to conform to standard encoder return signature.

signature.
None: Used to conform to standard encoder return signature.

```
"""
    batch_size, seq_len = inputs.size()

    encoded = self.embedding(inputs) # batch_size x seq_len x
hidden_size
```

```
    # Add positinal embeddings from
self.create_positional_encodings. (a'la
https://arxiv.org/pdf/1706.03762.pdf, section 3.5)
    encoded = encoded + self.positional_encodings[:seq_len]
```

```
    annotations = encoded
    for i in range(self.num_layers):
        new_annotations, self_attention_weights =
self.self_attentions[i](
            annotations, annotations, annotations
        ) # batch_size x seq_len x hidden_size
        residual_annotations = annotations + new_annotations
        new_annotations = self.attention_mlps[i]
(residual_annotations)
        annotations = residual_annotations + new_annotations
```

Transformer encoder does not have a last hidden or cell layer.

```
    return annotations, None
    # return annotations, None, None
def create_positional_encodings(self, max_seq_len=1000):
    """Creates positional encodings for the inputs.
```

Arguments:

max_seq_len: a number larger than the maximum string length we expect to encounter during training

Returns:

pos_encodings: (max_seq_len, hidden_dim) Positional encodings for a sequence with length max_seq_len.

"""

```
pos_indices = torch.arange(max_seq_len)[..., None]
dim_indices = torch.arange(self.hidden_size // 2)[None, ...]
exponents = (2 * dim_indices).float() / (self.hidden_size)
```

```

    trig_args = pos_indices / (10000**exponents)
    sin_terms = torch.sin(trig_args)
    cos_terms = torch.cos(trig_args)

    pos_encodings = torch.zeros((max_seq_len, self.hidden_size))
    pos_encodings[:, 0::2] = sin_terms
    pos_encodings[:, 1::2] = cos_terms

    if self.opts.cuda:
        pos_encodings = pos_encodings.cuda()

    return pos_encodings

class TransformerDecoder(nn.Module):
    def __init__(self, vocab_size, hidden_size, num_layers):
        super(TransformerDecoder, self).__init__()
        self.vocab_size = vocab_size
        self.hidden_size = hidden_size

        self.embedding = nn.Embedding(vocab_size, hidden_size)
        self.num_layers = num_layers

        self.self_attentions = nn.ModuleList(
            [
                CausalScaledDotAttention(
                    hidden_size=hidden_size,
                )
                for i in range(self.num_layers)
            ]
        )
        self.encoder_attentions = nn.ModuleList(
            [
                ScaledDotAttention(
                    hidden_size=hidden_size,
                )
                for i in range(self.num_layers)
            ]
        )
        self.attention_mlps = nn.ModuleList(
            [
                nn.Sequential(
                    nn.Linear(hidden_size, hidden_size),
                    nn.ReLU(),
                )
                for i in range(self.num_layers)
            ]
        )
        self.out = nn.Linear(hidden_size, vocab_size)

        self.positional_encodings = self.create_positional_encodings()

```

```

def forward(self, inputs, annotations, hidden_init):
    """Forward pass of the attention-based decoder RNN.

    Arguments:
        inputs: Input token indexes across a batch for all the
time step. (batch_size x decoder_seq_len)
        annotations: The encoder hidden states for each step of
the input.
sequence. (batch_size x seq_len x
hidden_size)
        hidden_init: Not used in the transformer decoder
    Returns:
        output: Un-normalized scores for each token in the
vocabulary, across a batch for all the decoding time steps.
(batch_size x decoder_seq_len x vocab_size)
        attentions: The stacked attention weights applied to the
encoder annotations (batch_size x encoder_seq_len x decoder_seq_len)
    """

    batch_size, seq_len = inputs.size()
    embed = self.embedding(inputs) # batch_size x seq_len x
hidden_size

    embed = embed + self.positional_encodings[:seq_len]

    encoder_attention_weights_list = []
    self_attention_weights_list = []
    contexts = embed
    for i in range(self.num_layers):
        new_contexts, self_attention_weights =
self.self_attentions[i](
            contexts, contexts, contexts
        ) # batch_size x seq_len x hidden_size
        residual_contexts = contexts + new_contexts
        new_contexts, encoder_attention_weights =
self.encoder_attentions[i](
            residual_contexts, annotations, annotations
        ) # batch_size x seq_len x hidden_size
        residual_contexts = residual_contexts + new_contexts
        new_contexts = self.attention_mlps[i](residual_contexts)
        contexts = residual_contexts + new_contexts

    encoder_attention_weights_list.append(encoder_attention_weights)
    self_attention_weights_list.append(self_attention_weights)

    output = self.out(contexts)
    encoder_attention_weights =

```

```

torch.stack(encoder_attention_weights_list)
    self_attention_weights =
torch.stack(self_attention_weights_list)

    return output, (encoder_attention_weights,
self_attention_weights)

def create_positional_encodings(self, max_seq_len=1000):
    """Creates positional encodings for the inputs.

    Arguments:
        max_seq_len: a number larger than the maximum string
length we expect to encounter during training

    Returns:
        pos_encodings: (max_seq_len, hidden_dim) Positional
encodings for a sequence with length max_seq_len.
    """
    pos_indices = torch.arange(max_seq_len)[..., None]
    dim_indices = torch.arange(self.hidden_size // 2)[None, ...]
    exponents = (2 * dim_indices).float() / (self.hidden_size)
    trig_args = pos_indices / (10000**exponents)
    sin_terms = torch.sin(trig_args)
    cos_terms = torch.cos(trig_args)

    pos_encodings = torch.zeros((max_seq_len, self.hidden_size))
    pos_encodings[:, 0::2] = sin_terms
    pos_encodings[:, 1::2] = cos_terms

    pos_encodings = pos_encodings.cuda()

    return pos_encodings

```

Step 9: Training and analysis (with scaled dot-product attention)

Now we will train a (simplified) transformer encoder-decoder model.

First, we train our smaller model on the small dataset. Use this model to answer Question 4 in the handout.

```
TEST_SENTENCE = "the air conditioning is working"
```

```

trans32_args_s = AttrDict()
args_dict = {
    "data_file_name": "pig_latin_small",
    "cuda": True,
    "nepochs": 100,
    "checkpoint_dir": "checkpoints",
    "learning_rate": 5e-4,
    "early_stopping_patience": 100,

```



```

        "lr_decay": 0.99,
        "batch_size": 64,
        "hidden_size": 32,
        "encoder_type": "transformer",
        "decoder_type": "transformer", # options: rnn / rnn_attention /
transformer
        "num_transformer_layers": 4,
    }
    trans32_args_s.update(args_dict)
    print_opts(trans32_args_s)

```

```

trans32_encoder_s, trans32_decoder_s, trans32_losses_s =
train(trans32_args_s)

```

```

translated = translate_sentence(
    TEST_SENTENCE, trans32_encoder_s, trans32_decoder_s, None,
    trans32_args_s
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```

=====
=====

```

Opts

```

-----
-----

```

```

data_file_name: pig_latin_small

cuda: 1

nepochs: 100

checkpoint_dir: checkpoints

learning_rate: 0.0005

early_stopping_patience: 100

lr_decay: 0.99

batch_size: 64

hidden_size: 32

encoder_type: transformer

decoder_type: transformer

```

num_transformer_layers: 4

```
=====
=====
=====
=====
```

Data Stats

```
-----
-----
```

```
('whisper', 'isperwhay')
('prayers', 'ayerspray')
('sketch', 'etchskay')
('disinclination', 'isinclinationday')
('tranquilize', 'anquilizetray')
```

Num unique word pairs: 3198

Vocabulary: dict_keys(['-', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z', 'SOS', 'EOS'])

Vocab size: 29

```
=====
=====
```

Moved models to GPU!

Epoch: 0 | Train loss: 3.680 | Val loss: 2.709 | Gen: ayayay ay

ayayay ay ay

Epoch: 1 | Train loss: 2.436 | Val loss: 2.318 | Gen: iyy iwy iy

itititoty if

Epoch: 2 | Train loss: 2.102 | Val loss: 2.074 | Gen: ay iay

itcE0SE0Scy irrrrrrrrrrrrrrrrrrwy iowywnyyyy

Epoch: 3 | Train loss: 1.907 | Val loss: 1.954 | Gen: ay iyy

iaannayay iwiyyyE0Syay ofofgwwwwy

Epoch: 4 | Train loss: 1.781 | Val loss: 1.868 | Gen: ay iay aanyyyy

iway onangwy

Epoch: 5 | Train loss: 1.691 | Val loss: 1.812 | Gen: ay iwy intai-

y isisisisisisisisE0Sy inoiyy

Epoch: 6 | Train loss: 1.636 | Val loss: 1.744 | Gen: ay iway

ongaaningiggyyy iway ongogoahy

Epoch: 7 | Train loss: 1.571 | Val loss: 1.761 | Gen: ay iway

insncaayyiy-y-ydiyy iswy onhohhygk

Epoch: 8 | Train loss: 1.494 | Val loss: 1.737 | Gen: eway ayay

onggggagiy iswwy onkkkkkkkkkk

Epoch: 9 | Train loss: 1.449 | Val loss: 1.630 | Gen: eyE0Sy iwwy

onggggnniiga-n--yaoyi isway onhkkhkkk

Epoch: 10 | Train loss: 1.396 | Val loss: 1.593 | Gen: ewE0Sy iwwy

ongna-aiy iswsy ongngngniE0Sgknhnn

Epoch: 11 | Train loss: 1.351 | Val loss: 1.689 | Gen: ebay iway

ongay--ynaayE0SE0SyyE0Siai iwwwy onhhohhhhhhhy

Epoch: 12 | Train loss: 1.340 | Val loss: 1.579 | Gen: essay irwy

oogyyayyyyyyy iswwwy onhnonnnrrrooonohrr

Epoch: 13 | Train loss: 1.305 | Val loss: 1.580 | Gen: ey iririrwry

ingay-iaay is onkgngigkE0Say
 Epoch: 14 | Train loss: 1.304 | Val loss: 1.489 | Gen: eay aywy oogay
 iwwy onokonokikwwy
 Epoch: 15 | Train loss: 1.240 | Val loss: 1.605 | Gen: eyy ay ayayy
 iwwy orahE0Saaay
 Epoch: 16 | Train loss: 1.207 | Val loss: 1.482 | Gen: eay ay ongny-
 iiy isway okayykakeE0SkE0Siiiayy
 Epoch: 17 | Train loss: 1.157 | Val loss: 1.464 | Gen: eay ayE0Siy
 ondwyE0SyE0Sy isaay irayahayaaaE0Syiiyy
 Epoch: 18 | Train loss: 1.140 | Val loss: 1.475 | Gen: eay ayE0Siy
 oogay iwwy onioiiyyhnwyyhiyygh
 Epoch: 19 | Train loss: 1.112 | Val loss: 1.461 | Gen: eay ayE0Siriwy
 ongiyyiy isaay ohgghahagyyyyhy
 Epoch: 20 | Train loss: 1.096 | Val loss: 1.448 | Gen: eay ayE0Siy
 ongny-ony iwwy onggaagagaggggE0Swwy
 Epoch: 21 | Train loss: 1.076 | Val loss: 1.451 | Gen: ehy aywywywy
 ongay iswwy ikkkkkkkkkyhhE0SE0Shaaaa
 Epoch: 22 | Train loss: 1.054 | Val loss: 1.414 | Gen: ettty ayE0Siry
 ongigggayaagay isay okiiiiignrnwhhkk
 Epoch: 23 | Train loss: 1.049 | Val loss: 1.402 | Gen: ety aywayE0Sy
 inaanniy isway okkkkiairrk
 Epoch: 24 | Train loss: 1.020 | Val loss: 1.443 | Gen: eththay
 ayiirirry ingny-i-gy iway ohghhhghhhiyhy
 Epoch: 25 | Train loss: 1.006 | Val loss: 1.418 | Gen: etty
 aywayE0Siy ongny isway okkgikyyygyiiykiyygy
 Epoch: 26 | Train loss: 0.991 | Val loss: 1.399 | Gen: etwy
 iriririrry ingngngngiy iwi ivivagE0SE0SE0Sgyy
 Epoch: 27 | Train loss: 0.971 | Val loss: 1.416 | Gen: ththtay ayway
 ongngngngnggynny iwwy okvgigaE0SkE0SE0SaE0SyaE0Skyy
 Epoch: 28 | Train loss: 0.960 | Val loss: 1.402 | Gen: etwy arryyriyy
 ongannyy iswwy okgghgwyaawy
 Epoch: 29 | Train loss: 0.930 | Val loss: 1.352 | Gen: tthahE0Sy
 airiy ingninnniy iwwy oghnhhrrgggy
 Epoch: 30 | Train loss: 1.349 | Val loss: 1.549 | Gen: ehay ayay
 iinihhhay iway okgkggggyy
 Epoch: 31 | Train loss: 1.169 | Val loss: 1.512 | Gen: ehathy arway
 ongngdy iway ogggaaggy
 Epoch: 32 | Train loss: 1.120 | Val loss: 1.478 | Gen: thhhhay
 ayaawwy ongniniiiy isayE0SE0Ssy ovonwwwywhhohry
 Epoch: 33 | Train loss: 1.137 | Val loss: 1.401 | Gen: ehaay ayE0Sy
 ongaayaiy isway ovingggy
 Epoch: 34 | Train loss: 1.044 | Val loss: 1.379 | Gen: ehathay ay
 ondniwgndy isway okikgwkkE0Syy
 Epoch: 35 | Train loss: 0.985 | Val loss: 1.371 | Gen: ehathay ayE0Sy
 ongnnynny isway ovingggy
 Epoch: 36 | Train loss: 0.960 | Val loss: 1.369 | Gen: ehathay ay
 ondwnyaiyy isway okingwgy
 Epoch: 37 | Train loss: 0.931 | Val loss: 1.355 | Gen: ehathay ay
 oninnyoiyny isway okingwgy
 Epoch: 38 | Train loss: 0.919 | Val loss: 1.352 | Gen: ehthay ay

ingnnydny isway okiinny
 Epoch: 39 | Train loss: 0.904 | Val loss: 1.354 | Gen: ehathay ay
 ondiyy isway okikgwwyE0Syy
 Epoch: 40 | Train loss: 0.941 | Val loss: 1.357 | Gen: ehathay arE0Sy
 ondindnny isway oviiway
 Epoch: 41 | Train loss: 1.302 | Val loss: 1.574 | Gen: ewwwwy iway
 ondondiiy iway okgaaaaE0Swy
 Epoch: 42 | Train loss: 1.175 | Val loss: 1.430 | Gen: ettwtwE0Sy
 iray ondnnddddgdgdy isway owwnaE0Syy
 Epoch: 43 | Train loss: 1.101 | Val loss: 1.419 | Gen: ettttty irayay
 ontingindiigggy isway ogwaagggy
 Epoch: 44 | Train loss: 1.059 | Val loss: 1.388 | Gen: etttway iray
 ondindindindyy isway ogiaagggy
 Epoch: 45 | Train loss: 1.021 | Val loss: 1.376 | Gen: etttway iray
 ondindindindndnnny isway okwwnny
 Epoch: 46 | Train loss: 0.997 | Val loss: 1.366 | Gen: etttway iray
 ondindindindanE0Snnny isway okwngny
 Epoch: 47 | Train loss: 0.978 | Val loss: 1.370 | Gen: etttway iray
 ondindindindndnnny isway ogwngny
 Epoch: 48 | Train loss: 0.964 | Val loss: 1.370 | Gen: etttway aray
 ondindindindiyE0SE0Sy isway ogwngwyy
 Epoch: 49 | Train loss: 0.949 | Val loss: 1.362 | Gen: etttay iray
 ondindindidy isway ogiwnny
 Epoch: 50 | Train loss: 0.934 | Val loss: 1.361 | Gen: etttway irwy
 ondindindinay iway ogwaaay
 Epoch: 51 | Train loss: 0.920 | Val loss: 1.352 | Gen: ettttay irwyy
 ondindindinay isway ogwgggy
 Epoch: 52 | Train loss: 0.912 | Val loss: 1.359 | Gen: ettttay irwyy
 ondindindinay iway oggkway
 Epoch: 53 | Train loss: 0.900 | Val loss: 1.357 | Gen: ettttay iriyy
 ondindindinwy isway oggkway
 Epoch: 54 | Train loss: 0.895 | Val loss: 1.349 | Gen: ettttay irayy
 ondindindinyy isway oggkway
 Epoch: 55 | Train loss: 0.884 | Val loss: 1.343 | Gen: ettthey irayy
 ondindindinay isway oggkwhay
 Epoch: 56 | Train loss: 0.872 | Val loss: 1.332 | Gen: ettthey iriy
 ondindindindidy isway ogggay
 Epoch: 57 | Train loss: 0.866 | Val loss: 1.330 | Gen: ettthey arwy
 ondindindinay isway okwngwyy
 Epoch: 58 | Train loss: 0.858 | Val loss: 1.329 | Gen: ettttay iriy
 ondindindinay isway oggawwaygy
 Epoch: 59 | Train loss: 0.850 | Val loss: 1.339 | Gen: ettthey iriry
 ondindindnnay isway okwngwyy
 Epoch: 60 | Train loss: 0.840 | Val loss: 1.332 | Gen: ettthey irway
 ondindindinay isway owwwwrwywy
 Epoch: 61 | Train loss: 0.831 | Val loss: 1.334 | Gen: ettthey arwry
 ondindindinay isway okwngwyy
 Epoch: 62 | Train loss: 0.823 | Val loss: 1.335 | Gen: ettthey
 irwayE0Sy ondindindnnay isway ogkkwawE0Say
 Epoch: 63 | Train loss: 0.819 | Val loss: 1.339 | Gen: etthay arway

ondindindiaay isway okgiawyy
Epoch: 64 | Train loss: 1.251 | Val loss: 1.672 | Gen: ewwy iraywwy
onininingny iwyyyy ogginy
Epoch: 65 | Train loss: 1.122 | Val loss: 1.443 | Gen: etway ieay
ongngny iway ogkwy
Epoch: 66 | Train loss: 1.017 | Val loss: 1.429 | Gen: eway awayay
ondniny iway ogiy
Epoch: 67 | Train loss: 0.970 | Val loss: 1.385 | Gen: eway iywy
ongagayE0Sy isway ogiy
Epoch: 68 | Train loss: 0.940 | Val loss: 1.368 | Gen: etwy aewy
ongnggday isway okkkgy
Epoch: 69 | Train loss: 0.920 | Val loss: 1.353 | Gen: ethay
aeE0SaE0SE0Sy onddiidy isway ogkwngy
Epoch: 70 | Train loss: 0.901 | Val loss: 1.357 | Gen: ethey irwy
ondididinE0Sdy isway ogknnny
Epoch: 71 | Train loss: 0.890 | Val loss: 1.366 | Gen: ethey
aewaE0Say oniniddndE0Sy isway ogkigggE0Sy
Epoch: 72 | Train loss: 0.880 | Val loss: 1.353 | Gen: ethey arwy
onininingE0Sy isway ogkigggE0Sy
Epoch: 73 | Train loss: 0.865 | Val loss: 1.350 | Gen: ethey iywy
ongngggay isway ogignwyy
Epoch: 74 | Train loss: 0.852 | Val loss: 1.345 | Gen: ethey iywy
ondndndndady isway ogkigggE0Sy
Epoch: 75 | Train loss: 0.844 | Val loss: 1.353 | Gen: ethey irwy
ondniidndwy isway ogknnnyaiy
Epoch: 76 | Train loss: 0.836 | Val loss: 1.347 | Gen: ethey iywy
ongngdgy isway ogkigggE0Sy
Epoch: 77 | Train loss: 0.827 | Val loss: 1.342 | Gen: ethey iywy
onindnnndway isway okkkay
Epoch: 78 | Train loss: 0.819 | Val loss: 1.333 | Gen: ethey iywy
ondniddy isway ogkigggE0Sy
Epoch: 79 | Train loss: 0.813 | Val loss: 1.339 | Gen: ethey iywy
oniningnday isway ogknnnyaky
Epoch: 80 | Train loss: 0.807 | Val loss: 1.334 | Gen: ethty aywy
onddiddy isway oginnwyy
Epoch: 81 | Train loss: 0.797 | Val loss: 1.337 | Gen: ethey aywy
oniningddwyy isway oginnwgy
Epoch: 82 | Train loss: 0.791 | Val loss: 1.332 | Gen: ethey aywy
oniningddwyy isway ogiigwyy
Epoch: 83 | Train loss: 0.784 | Val loss: 1.334 | Gen: ethey aywy
ongngdindwy isway ogggagay
Epoch: 84 | Train loss: 0.805 | Val loss: 1.367 | Gen: etwy aywy
ongdy isway ogwaaway
Epoch: 85 | Train loss: 0.832 | Val loss: 1.340 | Gen: ethty aywy
onddin-idyy isway ogkkkaE0Swy
Epoch: 86 | Train loss: 0.810 | Val loss: 1.333 | Gen: etway aywy
ongiiignnyny isway ogkwaagay
Epoch: 87 | Train loss: 0.799 | Val loss: 1.336 | Gen: etway aywy
ongiinginyny isway ogkwaagay
Epoch: 88 | Train loss: 0.788 | Val loss: 1.320 | Gen: ethty aywy

```

onddiinnyny isway ogkwaagay
Epoch: 89 | Train loss: 0.778 | Val loss: 1.325 | Gen: etway aywy
ongiingnnyny isway ogkwaagay
Epoch: 90 | Train loss: 0.771 | Val loss: 1.317 | Gen: etway aywy
ongny isway owwwagay
Epoch: 91 | Train loss: 0.766 | Val loss: 1.319 | Gen: etway aywy
ongdiinyyny isway ogkingway
Epoch: 92 | Train loss: 0.758 | Val loss: 1.320 | Gen: ethety aywy
ongny isway ogkingway
Epoch: 93 | Train loss: 0.761 | Val loss: 1.303 | Gen: etway aywy
ongniggoyyny isway oggiggway
Epoch: 94 | Train loss: 0.748 | Val loss: 1.315 | Gen: ethey aiaay
ongngdy isway oggiggway
Epoch: 95 | Train loss: 0.743 | Val loss: 1.311 | Gen: etway aywy
oniiinindaEOSy isway oggigway
Epoch: 96 | Train loss: 0.732 | Val loss: 1.305 | Gen: ethty aiaay
ongnigddndEOSy isway oggiggway
Epoch: 97 | Train loss: 0.727 | Val loss: 1.306 | Gen: ethty aywy
ondiininyyny isway oggigway
Epoch: 98 | Train loss: 0.719 | Val loss: 1.294 | Gen: ethty iywy
ongny isway ogkiggway
Epoch: 99 | Train loss: 0.716 | Val loss: 1.308 | Gen: etway aywy
ondiinanyyny isway ogkinway
Obtained lowest validation loss of: 1.2940169721841812
source:          the air conditioning is working
translated:      etway aywy ondiinanyyny isway ogkinway

```

```

TEST_SENTENCE = "the air conditioning is working"
translated = translate_sentence(
    TEST_SENTENCE, trans32_encoder_s, trans32_decoder_s, None,
    trans32_args_s
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```

source:          the air conditioning is working
translated:      etway aywy ondiinanyyny isway ogkinway

```

In the following cells, we investigate the effects of increasing model size and dataset size on the training / validation curves and generalization of the Transformer. We will increase hidden size to 64, and also increase dataset size. Include the best achieved validation loss in your report.

```
TEST_SENTENCE = "the air conditioning is working"
```

```

trans32_args_l = AttrDict()
args_dict = {
    "data_file_name": "pig_latin_large", # Increased data set size
    "cuda": True,
    "nepochs": 100,
    "checkpoint_dir": "checkpoints",

```

```

    "learning_rate": 5e-4,
    "early_stopping_patience": 10,
    "lr_decay": 0.99,
    "batch_size": 512,
    "hidden_size": 32,
    "encoder_type": "transformer",
    "decoder_type": "transformer", # options: rnn / rnn_attention /
transformer
    "num_transformer_layers": 3,
}
trans32_args_l.update(args_dict)
print_opts(trans32_args_l)

```

```

trans32_encoder_l, trans32_decoder_l, trans32_losses_l =
train(trans32_args_l)

```

```

translated = translate_sentence(
    TEST_SENTENCE, trans32_encoder_l, trans32_decoder_l, None,
trans32_args_l
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```

=====
=====

```

Opts

```

-----
-----

```

data_file_name: pig_latin_large

cuda: 1

nepochs: 100

checkpoint_dir: checkpoints

learning_rate: 0.0005

early_stopping_patience: 10

lr_decay: 0.99

batch_size: 512

hidden_size: 32

encoder_type: transformer

decoder_type: transformer

num_transformer_layers: 3

```
=====
=====
=====
=====
```

Data Stats

```
-----
-----
('brass', 'assbray')
('limo', 'imolay')
('nb', 'nbay')
('race', 'aceray')
('clarify', 'arifyclay')
Num unique word pairs: 22402
Vocabulary: dict_keys(['-', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h',
'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v',
'w', 'x', 'y', 'z', 'S0S', 'E0S'])
Vocab size: 29
=====
=====
```

Moved models to GPU!

```
Epoch:  0 | Train loss: 3.143 | Val loss: 2.665 | Gen: ayE0Sy
ayayayayayayayayayay ayay ayay
Epoch:  1 | Train loss: 2.474 | Val loss: 2.392 | Gen: ete-t-
ayayayay oooooooooooooooooooooo ay oooooooooooooooooooooo
Epoch:  2 | Train loss: 2.255 | Val loss: 2.250 | Gen: ayay aay
oooooooooooooooooooooooooooo ay oooooooooooooooooooooooooo
Epoch:  3 | Train loss: 2.117 | Val loss: 2.146 | Gen: - aaay inty ay
ooooooooooooooooooooonooooon
Epoch:  4 | Train loss: 2.016 | Val loss: 2.112 | Gen: ay ay indy ay
aaE0SaE0SaE0SE0Sy
Epoch:  5 | Train loss: 1.939 | Val loss: 2.024 | Gen: o- ay innny
ay onnnnnnnnnnnnnnnnnnnnn
Epoch:  6 | Train loss: 1.856 | Val loss: 1.959 | Gen: adhtyyE0Sy ay
innntay ay onnnnnnnnnnnnnnnnnnnnn
Epoch:  7 | Train loss: 1.791 | Val loss: 1.912 | Gen:
ayhtE0SttyyyE0Syy ay innndtay ay onnnnnnnnnnnnnnnnnnnnn
Epoch:  8 | Train loss: 1.737 | Val loss: 1.873 | Gen: ay ay
innntaay ay onnnnnnnnnnnnyE0Syaay
Epoch:  9 | Train loss: 1.690 | Val loss: 1.854 | Gen: ay ay indddy
ay onnoony
Epoch: 10 | Train loss: 1.685 | Val loss: 1.914 | Gen: ay
ayayayayayayayayay indy-y isayay onnnnnnnnnnnnnnnny
Epoch: 11 | Train loss: 1.674 | Val loss: 1.823 | Gen: ay aray inty
iy onrgggay
Epoch: 12 | Train loss: 1.620 | Val loss: 1.789 | Gen: ay ay indcy iy
```


ongrgggry
Epoch: 13 | Train loss: 1.578 | Val loss: 1.744 | Gen: ay aray
intintntntnnay iy ongrgggay
Epoch: 14 | Train loss: 2.129 | Val loss: 2.063 | Gen: aay ararsy
ondaE0Sy ayayayE0SE0Sy inyay
Epoch: 15 | Train loss: 1.833 | Val loss: 1.937 | Gen: ----- ayarh-
aaE0Syy onddyaay ayyE0Sy usyay
Epoch: 16 | Train loss: 1.755 | Val loss: 1.883 | Gen: --- arsytsy
onc sssssssssssssssssssssss onc
Epoch: 17 | Train loss: 1.707 | Val loss: 1.850 | Gen: ----- arstssy
onddyy asyE0Syy onc
Epoch: 18 | Train loss: 1.669 | Val loss: 1.826 | Gen: ----- arsyssy
onddyy ayyE0Sy ongpay
Epoch: 19 | Train loss: 1.635 | Val loss: 1.794 | Gen: ehhtttt-
arayssy ongccy ayyE0Sy ongy
Epoch: 20 | Train loss: 1.606 | Val loss: 1.772 | Gen: ----- arayssy
intiy ayyE0Sy ongy
Epoch: 21 | Train loss: 1.626 | Val loss: 1.781 | Gen:
athhhhhththhhhhhhhty arsssssssssshshssss ingcy ay onay
Epoch: 22 | Train loss: 1.607 | Val loss: 1.753 | Gen: ay
arsssssssssshssss ingcy aay onay
Epoch: 23 | Train loss: 1.578 | Val loss: 1.750 | Gen: ay ararE0Sy
inaay ay ongyE0Syy
Epoch: 24 | Train loss: 1.554 | Val loss: 1.711 | Gen: ay
arE0Sriiasryyyy inaay ay ogE0SaE0SE0SE0SE0SE0Sy
Epoch: 25 | Train loss: 1.533 | Val loss: 1.714 | Gen: ay ayE0Sraaayy
inaay ay ongray
Epoch: 26 | Train loss: 1.514 | Val loss: 1.677 | Gen: ay aysyyyy
inaay isay ongayE0SyE0SE0Sy
Epoch: 27 | Train loss: 1.496 | Val loss: 1.682 | Gen: ay arayE0Sy
inaay isay ongaE0Syy
Epoch: 28 | Train loss: 1.489 | Val loss: 1.632 | Gen: ay arayE0Sy
inaay aaay ongayE0SE0Sy
Epoch: 29 | Train loss: 1.469 | Val loss: 1.651 | Gen: ayE0StE0SyE0Sy
araray ingay asay ongiay
Epoch: 30 | Train loss: 1.452 | Val loss: 1.612 | Gen: ayE0StE0SyE0Sy
ararE0Sy ingay asay odraay
Epoch: 31 | Train loss: 1.437 | Val loss: 1.634 | Gen:
ehhhhhhhhhhhhhhhhhhh arararE0Srayayy ingay asay ongyy
Epoch: 32 | Train loss: 1.421 | Val loss: 1.594 | Gen: ayE0StE0SyE0Sy
ararE0Sy ingay asay odraay
Epoch: 33 | Train loss: 1.407 | Val loss: 1.619 | Gen:
ehhhhhhhhhhhhhhhhhhh arayE0Syayy ingay ay ongyy
Epoch: 34 | Train loss: 1.427 | Val loss: 1.592 | Gen:
ehhhhhhhhhhhhhhhhyE0Syh aray ongcayyE0Sy sy oggrgay
Epoch: 35 | Train loss: 1.401 | Val loss: 1.611 | Gen: - aray
ongE0Say ay ongy
Epoch: 36 | Train loss: 1.382 | Val loss: 1.583 | Gen: ehhhhayE0Sy
aray ondingy ay oorgay
Epoch: 37 | Train loss: 1.367 | Val loss: 1.580 | Gen:

ehhhhhhhhhhhhhhhhhhhhh array ondingy ay ooorgay
Epoch: 38 | Train loss: 1.356 | Val loss: 1.567 | Gen: ehhhhyyE0Sy
array ondingy ay oogry
Epoch: 39 | Train loss: 1.345 | Val loss: 1.567 | Gen:
eththhhE0SyyyE0SE0SE0Shhhayy array onciy ay oogry
Epoch: 40 | Train loss: 1.335 | Val loss: 1.538 | Gen: adhE0Syy array
ondingcy ay ongrgy
Epoch: 41 | Train loss: 1.329 | Val loss: 1.565 | Gen:
ehhhhhhhhhhhhhhhhhhhhh array onciy ay ogggrrgy
Epoch: 42 | Train loss: 1.320 | Val loss: 1.518 | Gen: ethhtE0Say
array ondingcy ay ongrgy
Epoch: 43 | Train loss: 1.309 | Val loss: 1.537 | Gen:
ehhhhhhhhhhhhhhhhhhhha araray onciy ay ogggrrgy
Epoch: 44 | Train loss: 1.296 | Val loss: 1.508 | Gen: ethhyay array
oncingcy ay oonrgry
Epoch: 45 | Train loss: 1.285 | Val loss: 1.520 | Gen:
ehhhhhhhhhhhhhhhhhhhhy array onciy ay ogggrrgy
Epoch: 46 | Train loss: 1.276 | Val loss: 1.494 | Gen: ahhhhaay array
ongcngciy ay ogggrrgy
Epoch: 47 | Train loss: 1.266 | Val loss: 1.504 | Gen:
ehhhhhhhhhhhhhhhhhhhhy array onciy ay oggggggy
Epoch: 48 | Train loss: 1.257 | Val loss: 1.483 | Gen: ahhhhaay array
ongcngciy ay ogggrrgy
Epoch: 49 | Train loss: 1.250 | Val loss: 1.494 | Gen:
ettttayE0SyyyE0SyhE0Shhaay araray onciy ay oggggggy
Epoch: 50 | Train loss: 1.369 | Val loss: 1.535 | Gen: ewhhawy away
oncingcy isy oggggggyyy
Epoch: 51 | Train loss: 1.319 | Val loss: 1.499 | Gen: ewhhhaay array
oncongcy isy oggggggaE0Sy
Epoch: 52 | Train loss: 1.296 | Val loss: 1.477 | Gen: ewhthaay array
onccndcky isy oggggggaE0Sy
Epoch: 53 | Train loss: 1.279 | Val loss: 1.465 | Gen:
ewhyhaytyhhyhhy array ongcngggy isy opgrry
Epoch: 54 | Train loss: 1.266 | Val loss: 1.464 | Gen: ewhtawy array
ongcngggy isy oggggggaE0Sy
Epoch: 55 | Train loss: 1.256 | Val loss: 1.448 | Gen: ethhahy array
ongcngggy isy opgrry
Epoch: 56 | Train loss: 1.243 | Val loss: 1.439 | Gen: ethtawy array
ongcngggy iay oggggggyyy
Epoch: 57 | Train loss: 1.233 | Val loss: 1.432 | Gen: ethtahay
araray ongcngggy iay oggggggyay
Epoch: 58 | Train loss: 1.223 | Val loss: 1.428 | Gen: eththaay
araray ongcngggy isy oggggggyy
Epoch: 59 | Train loss: 1.214 | Val loss: 1.418 | Gen: ethtahay
araray ongcngggy isy odshrrss
Epoch: 60 | Train loss: 1.204 | Val loss: 1.417 | Gen: eththaay
araway ondindcdy iay ogggiiggyy
Epoch: 61 | Train loss: 1.197 | Val loss: 1.409 | Gen: ethtaway
araray ongcngggy isy odgrrrrss
Epoch: 62 | Train loss: 1.189 | Val loss: 1.406 | Gen: eththayaay

```

araway ondinidnccy iay oggiiggyy
Epoch: 63 | Train loss: 1.181 | Val loss: 1.400 | Gen: eththaayE0Sy
araray ongingtty iay odgrry
Epoch: 64 | Train loss: 1.174 | Val loss: 1.401 | Gen: eththayaay
araray ontinidnndy iay ogrigggyy
Epoch: 65 | Train loss: 1.168 | Val loss: 1.400 | Gen: eththaayE0Sy
araray ongingtty iay opgrggy
Epoch: 66 | Train loss: 1.199 | Val loss: 1.428 | Gen: eththayaay
aray ontcittiigE0Scy iay opsry
Epoch: 67 | Train loss: 1.668 | Val loss: 1.810 | Gen: thththy away
oy ieysstay o
Epoch: 68 | Train loss: 1.564 | Val loss: 1.650 | Gen:
ththttththththththhaa awaywy onE0SnE0Syyy iay o
Epoch: 69 | Train loss: 1.464 | Val loss: 1.597 | Gen:
eththththththththhhty aray ociiyy iay onny
Epoch: 70 | Train loss: 1.414 | Val loss: 1.562 | Gen:
eththththththththhhhhhy aray ontioyiy iay onE0Sy
Epoch: 71 | Train loss: 1.378 | Val loss: 1.542 | Gen:
ethththththhaytthhhay aray ontioy iay oniy
Epoch: 72 | Train loss: 1.354 | Val loss: 1.528 | Gen:
tththththththththhhhhhy aray ontioy iay oniy
Epoch: 73 | Train loss: 1.333 | Val loss: 1.513 | Gen:
eththththththththhhhhhy aray ontiny isy oniy
Validation loss has not improved in 10 epochs, stopping early
Obtained lowest validation loss of: 1.3996964582690485
source:         the air conditioning is working
translated:     ehthththththhaE0Shhhhy aray ontiny isy onay

```

```
TEST_SENTENCE = "the air conditioning is working"
```

```

trans64_args_s = AttrDict()
args_dict = {
    "data_file_name": "pig_latin_small",
    "cuda": True,
    "nepochs": 50,
    "checkpoint_dir": "checkpoints",
    "learning_rate": 5e-4,
    "early_stopping_patience": 20,
    "lr_decay": 0.99,
    "batch_size": 64,
    "hidden_size": 64, # Increased model size
    "encoder_type": "transformer",
    "decoder_type": "transformer", # options: rnn / rnn_attention /
    transformer
    "num_transformer_layers": 3,
}
trans64_args_s.update(args_dict)
print_opts(trans64_args_s)

```

```
trans64_encoder_s, trans64_decoder_s, trans64_losses_s =
```

```

train(trans64_args_s)

translated = translate_sentence(
    TEST_SENTENCE, trans64_encoder_s, trans64_decoder_s, None,
    trans64_args_s
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```

=====
=====

```

Opts

```

-----
-----

```

```

                data_file_name: pig_latin_small
                    cuda: 1
                    nepochs: 50
                checkpoint_dir: checkpoints
                    learning_rate: 0.0005
        early_stopping_patience: 20
                    lr_decay: 0.99
                    batch_size: 64
                    hidden_size: 64
                    encoder_type: transformer
                    decoder_type: transformer
        num_transformer_layers: 3

```

```

=====
=====
=====
=====

```

Data Stats

```

-----
-----

```

```

('whisper', 'isperwhay')
('prayers', 'ayerspray')

```


Epoch: 19 | Train loss: 0.572 | Val loss: 1.061 | Gen: eteteay iraway
ondtitiaanay isway orikgiay
Epoch: 20 | Train loss: 0.540 | Val loss: 1.002 | Gen: ehay
awayE0Swayyyy ondctitgnawwy isway okiaiggggy
Epoch: 21 | Train loss: 0.489 | Val loss: 0.965 | Gen: ethay iwayay
ondtiinainy isway okkigwnay
Epoch: 22 | Train loss: 0.449 | Val loss: 0.944 | Gen: ethay iwayay
ondtittngny isway okkiggwgy
Epoch: 23 | Train loss: 0.424 | Val loss: 0.958 | Gen: ethay
awaaayE0Syy onatiitnany isway okkingway
Epoch: 24 | Train loss: 0.525 | Val loss: 1.089 | Gen: etetay iriyay
oniitininiiggiyy isway okkiiggkky
Epoch: 25 | Train loss: 0.576 | Val loss: 0.967 | Gen: etethay irwy
indtittdnwcnccaiy isway okrknnkjkjy
Epoch: 26 | Train loss: 0.489 | Val loss: 0.870 | Gen: ethay irayy
ondititintotoay isway okiningway
Epoch: 27 | Train loss: 0.425 | Val loss: 0.909 | Gen: ethay irway
onditinininnE0Sy isway orkingwayE0Sy
Epoch: 28 | Train loss: 0.392 | Val loss: 0.824 | Gen: ethay iray
ondtitiniicgy isway okingwgyay
Epoch: 29 | Train loss: 0.339 | Val loss: 0.805 | Gen: ethay iriaE0Sy
onditiningnny isway okingwgyay
Epoch: 30 | Train loss: 0.311 | Val loss: 0.777 | Gen: ethay iray
onditionoooy isway okingway
Epoch: 31 | Train loss: 0.293 | Val loss: 0.762 | Gen: ethay irayy
ondititiningcayy isway okingwgyay
Epoch: 32 | Train loss: 0.286 | Val loss: 0.934 | Gen: ethay iray
onditiningay isway okwakway
Epoch: 33 | Train loss: 0.329 | Val loss: 0.835 | Gen: ethay irayy
onditiionngoy isway okikgkngway
Epoch: 34 | Train loss: 0.315 | Val loss: 0.855 | Gen: ehay irayy
onditinininnccy isway okkingway
Epoch: 35 | Train loss: 0.352 | Val loss: 0.804 | Gen: ethay arawy
onditionioiy isway okikiwgyy
Epoch: 36 | Train loss: 0.265 | Val loss: 0.726 | Gen: ethay irway
onditioniongaay isway okikiayE0SE0Sy
Epoch: 37 | Train loss: 0.219 | Val loss: 0.754 | Gen: ethay irway
onditioningcay isway okingwgyay
Epoch: 38 | Train loss: 0.236 | Val loss: 0.723 | Gen: ethyy irray
onditionioncay isway okingwgyay
Epoch: 39 | Train loss: 0.223 | Val loss: 0.816 | Gen: ethyy arway
onditiononoocyaE0Sy isway okingway
Epoch: 40 | Train loss: 0.239 | Val loss: 0.727 | Gen: ethyy ariray
ondititingincy isway orkingway
Epoch: 41 | Train loss: 0.194 | Val loss: 0.686 | Gen: ethyy arwayy
ondititionggay isway orkingway
Epoch: 42 | Train loss: 0.158 | Val loss: 0.731 | Gen: ethyy ariaay
onditioioiiaay isway orkingway
Epoch: 43 | Train loss: 0.154 | Val loss: 0.685 | Gen: ethyy irrwy
onditioionoicayE0Sy isway orkingway

```

Epoch: 44 | Train loss: 0.148 | Val loss: 0.779 | Gen: ethhy ariryay
onnditiningcaay isway orkingway
Epoch: 45 | Train loss: 0.154 | Val loss: 0.723 | Gen: ethyy iiiay
onnditioioncny isway orkingway
Epoch: 46 | Train loss: 0.127 | Val loss: 0.749 | Gen: ethyy ariwyay
onddiioonoonay isway orkingway
Epoch: 47 | Train loss: 0.133 | Val loss: 0.671 | Gen: ethyy arrwaay
onniitooooonocyy isway orkingway
Epoch: 48 | Train loss: 0.126 | Val loss: 0.724 | Gen: eehyy arrwyay
onnditionioncy isway orkingway
Epoch: 49 | Train loss: 0.134 | Val loss: 0.790 | Gen: ethyy airway
ondditiiooay isway orkingway
Obtained lowest validation loss of: 0.6713013056861726
source:          the air conditioning is working
translated:      ethyy airway ondditiiooay isway orkingway

```

```
TEST_SENTENCE = "the air conditioning is working"
```

```

trans64_args_l = AttrDict()
args_dict = {
    "data_file_name": "pig_latin_large", # Increased data set size
    "cuda": True,
    "nepochs": 50,
    "checkpoint_dir": "checkpoints",
    "learning_rate": 5e-4,
    "early_stopping_patience": 20,
    "lr_decay": 0.99,
    "batch_size": 512,
    "hidden_size": 64, # Increased model size
    "encoder_type": "transformer",
    "decoder_type": "transformer", # options: rnn / rnn_attention /
transformer
    "num_transformer_layers": 3,
}
trans64_args_l.update(args_dict)
print_opts(trans64_args_l)

```

```

trans64_encoder_l, trans64_decoder_l, trans64_losses_l =
train(trans64_args_l)

```

```

translated = translate_sentence(
    TEST_SENTENCE, trans64_encoder_l, trans64_decoder_l, None,
    trans64_args_l
)
print("source:\t\t{} \ntranslated:\t{}".format(TEST_SENTENCE,
translated))

```

```

=====
=====

```

Opts

```

-----
-----
data_file_name: pig_latin_large
    cuda: 1
    nepochs: 50
checkpoint_dir: checkpoints
    learning_rate: 0.0005
early_stopping_patience: 20
    lr_decay: 0.99
    batch_size: 512
    hidden_size: 64
    encoder_type: transformer
    decoder_type: transformer
num_transformer_layers: 3

=====
=====
=====
=====

Data Stats

-----
-----
('brass', 'assbray')
('limo', 'imolay')
('nb', 'nbay')
('race', 'aceray')
('clarify', 'arifyclay')
Num unique word pairs: 22402
Vocabulary: dict_keys(['-', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h',
'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v',
'w', 'x', 'y', 'z', 'SOS', 'EOS'])
Vocab size: 29
=====
=====
Moved models to GPU!
Epoch:  0 | Train loss: 2.732 | Val loss: 2.294 | Gen: ey ay

```


iaaaaaaaaaadaaaayy inginiin nyy
Epoch: 1 | Train loss: 2.024 | Val loss: 2.081 | Gen: t ay intntay
int igy
Epoch: 2 | Train loss: 1.806 | Val loss: 1.884 | Gen: e ay
ontnnoyooooaEOS-oyEOSSE0Soon isisaaaassssssyy inay
Epoch: 3 | Train loss: 1.660 | Val loss: 1.756 | Gen: etey ay
ontintintintintiy isisisisisisiway ingy
Epoch: 4 | Train loss: 1.571 | Val loss: 1.673 | Gen: etay ay
ondadaaddnnnnnddy isssssssssssssssssss ingggy
Epoch: 5 | Train loss: 1.499 | Val loss: 1.612 | Gen: etayhy ay
ontiiiiiyy isississsssiiiiissi ingaaay
Epoch: 6 | Train loss: 1.434 | Val loss: 1.550 | Gen: etay
ararrrrrrrrrrrrrrarra indayEOSayEOSSE0Sywyayy isssssssssssssssssss ongiy
Epoch: 7 | Train loss: 1.315 | Val loss: 1.473 | Gen: eththt ay
ontintintintintintin isayEOSSE0SyEOSaaay ingiyy
Epoch: 8 | Train loss: 1.331 | Val loss: 1.484 | Gen: eay ay
ongoogoogonnniaanyyn isayyyyy y
Epoch: 9 | Train loss: 1.852 | Val loss: 2.158 | Gen: ttway ay
ommmmmmmmmmmommammmm iay owaywaywaywaywaywayw
Epoch: 10 | Train loss: 1.629 | Val loss: 1.872 | Gen: ey ay
ondiiddyddydydddiyn iay onaaygy-y
Epoch: 11 | Train loss: 1.518 | Val loss: 1.690 | Gen: ay ay inniniy
iy y
Epoch: 12 | Train loss: 1.425 | Val loss: 1.562 | Gen: etway ay
ondin-inEOSinEOSinEOSnnind iy ay
Epoch: 13 | Train loss: 1.358 | Val loss: 1.510 | Gen: ey ay
ondindindindinaindin iy ay
Epoch: 14 | Train loss: 1.320 | Val loss: 1.503 | Gen: twaay ay
ondindgnEOSinEOSinninind iy yy
Epoch: 15 | Train loss: 1.282 | Val loss: 1.468 | Gen: eway ay
ondindindindindindin isaay ay
Epoch: 16 | Train loss: 1.258 | Val loss: 1.466 | Gen: eway ay
ondindindindindindin isway ay
Epoch: 17 | Train loss: 1.233 | Val loss: 1.428 | Gen: eway ay
ondindindindindindin isway ay
Epoch: 18 | Train loss: 1.220 | Val loss: 1.420 | Gen: eway ay
ondindiddincincincin isaay ay
Epoch: 19 | Train loss: 1.199 | Val loss: 1.387 | Gen: eway aray
ondindindindway iway ay
Epoch: 20 | Train loss: 1.188 | Val loss: 1.407 | Gen: eway ay
ondindincincincincin isway oy
Epoch: 21 | Train loss: 1.164 | Val loss: 1.370 | Gen: eway irrrray
ondindindindway isayy oy
Epoch: 22 | Train loss: 1.154 | Val loss: 1.364 | Gen: eway ay
ondindindincincincin isway oy
Epoch: 23 | Train loss: 1.127 | Val loss: 1.325 | Gen: eway irrrray
ondindindindwny isayy oway
Epoch: 24 | Train loss: 1.108 | Val loss: 1.315 | Gen: eway ay
ondindindindwnnEOSnndn isway oway
Epoch: 25 | Train loss: 1.092 | Val loss: 1.334 | Gen: eway irryy

ondiddiddindidyE0Snycc isayy owiy
Epoch: 26 | Train loss: 1.081 | Val loss: 1.319 | Gen: eway array
oncincindincwacwincn isway owiy
Epoch: 27 | Train loss: 1.068 | Val loss: 1.269 | Gen: eway irrroy
ondindiddidday isway owiy
Epoch: 28 | Train loss: 1.069 | Val loss: 1.325 | Gen: ewty array
oncincindincwncinnin isway orgiggay
Epoch: 29 | Train loss: 1.061 | Val loss: 1.261 | Gen: eway array
ondindandiniindnE0Scni isayy owiy
Epoch: 30 | Train loss: 1.055 | Val loss: 1.314 | Gen: eththtty array
ondindindindincindnc isayy oniy
Epoch: 31 | Train loss: 1.028 | Val loss: 1.247 | Gen: eththaay
array ondicid-y isayy oniy
Epoch: 32 | Train loss: 1.009 | Val loss: 1.272 | Gen: eththtwy ay
ondindindinciyacay isayy okiy
Epoch: 33 | Train loss: 0.998 | Val loss: 1.247 | Gen: eththtay
iraraE0SE0Siirriirry oncincicccccccciiii isayE0Siy oniy
Epoch: 34 | Train loss: 0.989 | Val loss: 1.246 | Gen: ethththay
array oncccciccinwncinnin isayy okiy
Epoch: 35 | Train loss: 0.966 | Val loss: 1.185 | Gen: eththaay
array onginginainny isayy oknnkky
Epoch: 36 | Train loss: 0.974 | Val loss: 1.379 | Gen: eththtty array
ongingiggingngggnegg isway okngaay
Epoch: 37 | Train loss: 0.985 | Val loss: 1.334 | Gen: eththaay array
oncicczcaaE0Scnycy isisE0Siy owiy
Epoch: 38 | Train loss: 0.975 | Val loss: 1.296 | Gen: ethtay array
ongingdggiinoyyiygid isway okiy
Epoch: 39 | Train loss: 0.974 | Val loss: 1.277 | Gen: eththtay
array oncincindiicwdy issyyiiiiiissyy owiy
Epoch: 40 | Train loss: 0.943 | Val loss: 1.157 | Gen: eththaay array
ongingnndndwnnoingen isway orinni
Epoch: 41 | Train loss: 0.917 | Val loss: 1.173 | Gen: ethththtty
arrrrarraaararararaar ondincindwE0Sy isayyy oriingay
Epoch: 42 | Train loss: 0.913 | Val loss: 1.232 | Gen: eththaay
array ongingingingingingnc isway orkingway
Epoch: 43 | Train loss: 0.894 | Val loss: 1.190 | Gen: eththaay
irrayE0Siay oncincinE0SiaE0Sineinccc isway orkingy
Epoch: 44 | Train loss: 0.882 | Val loss: 1.179 | Gen: eththaay array
ongnngonnnnnnicyiwn isway orkngwyE0Sy
Epoch: 45 | Train loss: 0.864 | Val loss: 1.133 | Gen: eththaay
array ondiniiiiwywywy isayyy orkingway
Epoch: 46 | Train loss: 0.851 | Val loss: 1.179 | Gen: eththaay array
ondindiinwicacny isway orkrgwyE0Sy
Epoch: 47 | Train loss: 0.842 | Val loss: 1.129 | Gen: eththaay
array onddndngw-cconeccei issway orkingway
Epoch: 48 | Train loss: 0.833 | Val loss: 1.145 | Gen: ethtay array
ondingooaaaE0Swnniinii isway oriingway
Epoch: 49 | Train loss: 0.823 | Val loss: 1.115 | Gen: eththaay
array ondinddny isswy oriingway
Obtained lowest validation loss of: 1.115473040552051

source: the air conditioning is working
translated: eththaay array ondinddny isswy oriingway

The following cell generates two loss plots. In the first plot, we compare the effects of increasing dataset size. In the second plot, we compare the effects of increasing model size. Include both plots in your report, and include your analysis of the results.

```
save_loss_comparison_by_dataset(  
    trans32_losses_s,  
    trans32_losses_l,  
    trans64_losses_s,  
    trans64_losses_l,  
    trans32_args_s,  
    trans32_args_l,  
    trans64_args_s,  
    trans64_args_l,  
    "trans_by_dataset",  
)  
save_loss_comparison_by_hidden(  
    trans32_losses_s,  
    trans32_losses_l,  
    trans64_losses_s,  
    trans64_losses_l,  
    trans32_args_s,  
    trans32_args_l,  
    trans64_args_s,  
    trans64_args_l,  
    "trans_by_hidden",  
)
```

<Figure size 432x288 with 0 Axes>

<Figure size 432x288 with 0 Axes>

Scaled Dot Product Attention

Part 2.1: Additive Attention

It takes around four and a half minutes to run the training loops in Part 1 but it takes around three and a half minutes to run the training loop in Part 2. Additive attention trains faster as it converges faster, which triggers early stopping. Therefore, we would not need to go through all 50 epochs like the GRUs.

(More comments regarding training time between RNNAttention and ScaledDotAttention can be found in Part 2.2 Question 3)

Part 2.2: Scaled Dot Product Attention

Question 1 & 2

Code for `ScaledDotProduct` and `CausalScaledDotProduct` can be found in Step 4 and Step 5 respectively.

Question 3

The model trained using `ScaledDotAttention` performs worse than `RNNAttention`. However, it took much less time to run 100 epoches of `ScaledDotAttention` than 35 epoches of `RNNAttention`. This reveals that `ScaledDotAttention` is more computationally efficient and space-efficient. `ScaledDotAttention` is also easier to implement as it uses highly optimized matrix multiplication (whereas additive attention uses a two layer fully-connected network). This could be the reason that `RNNAttention` outperforms `ScaledDotAttention`.

Question 4

We chose to represent word positions in this manner because we want the positions to be bounded and unique for each time step. Another advantage of using sine and cosine is that we could retrieve information regarding the relative positions between 2 words through linear transformation (matrix).

Although one-hot encoding can be used for positional encoding, one hot encoding is a long list/array where every bit is 1 or 0. Flipping bits for a long array becomes a waste of space when you could utilize floats to represent positions.

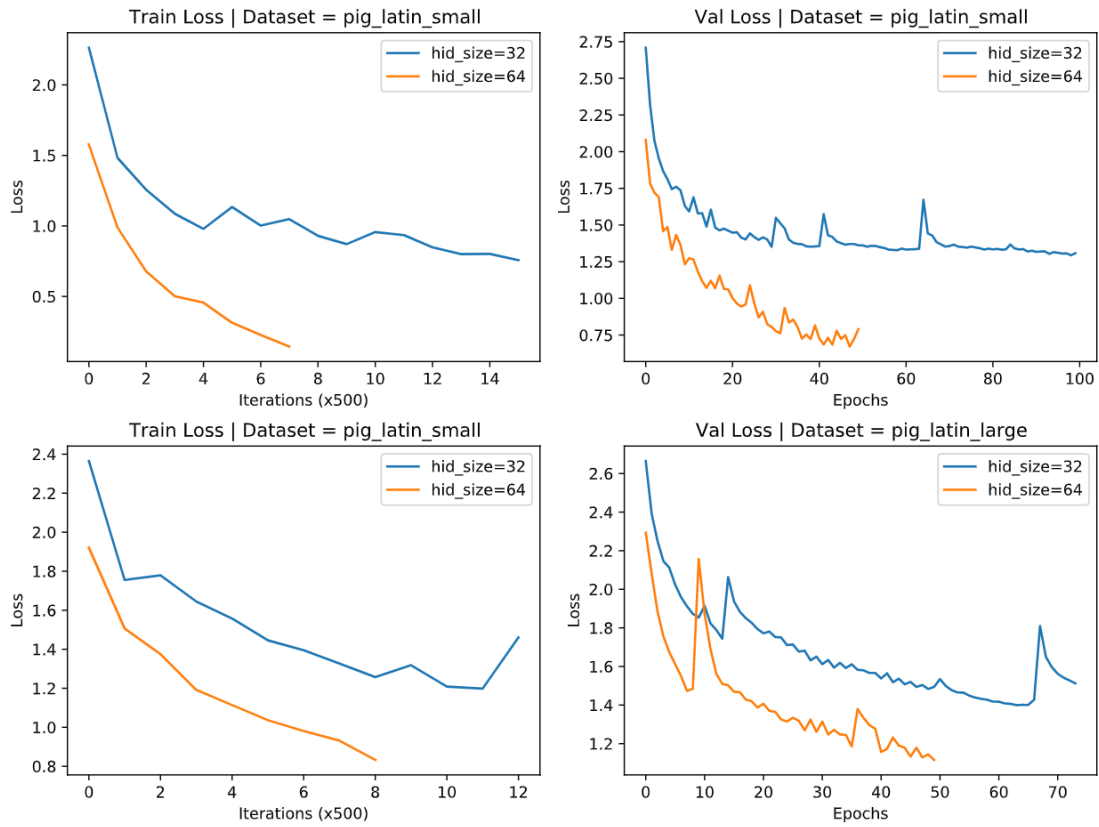
Question 5

The `TransformerEncoder` and `TransformerDecoder` obtains the lowest validation accuracy of 1.02. It performs better than single block `Attention` that obtains the lowest validation accuracy of 1.32. However, it is worse than `RNNAttention` model, which obtains the lowest validation accuracy of about 0.33. Additionally, both `Transformer` and single-layer attention are trained using all 100 epoches but the `RNN Attention` model reached early stopping condition. Therefore, `RNN Attention` is the best model among all these 3 models.

Question 6

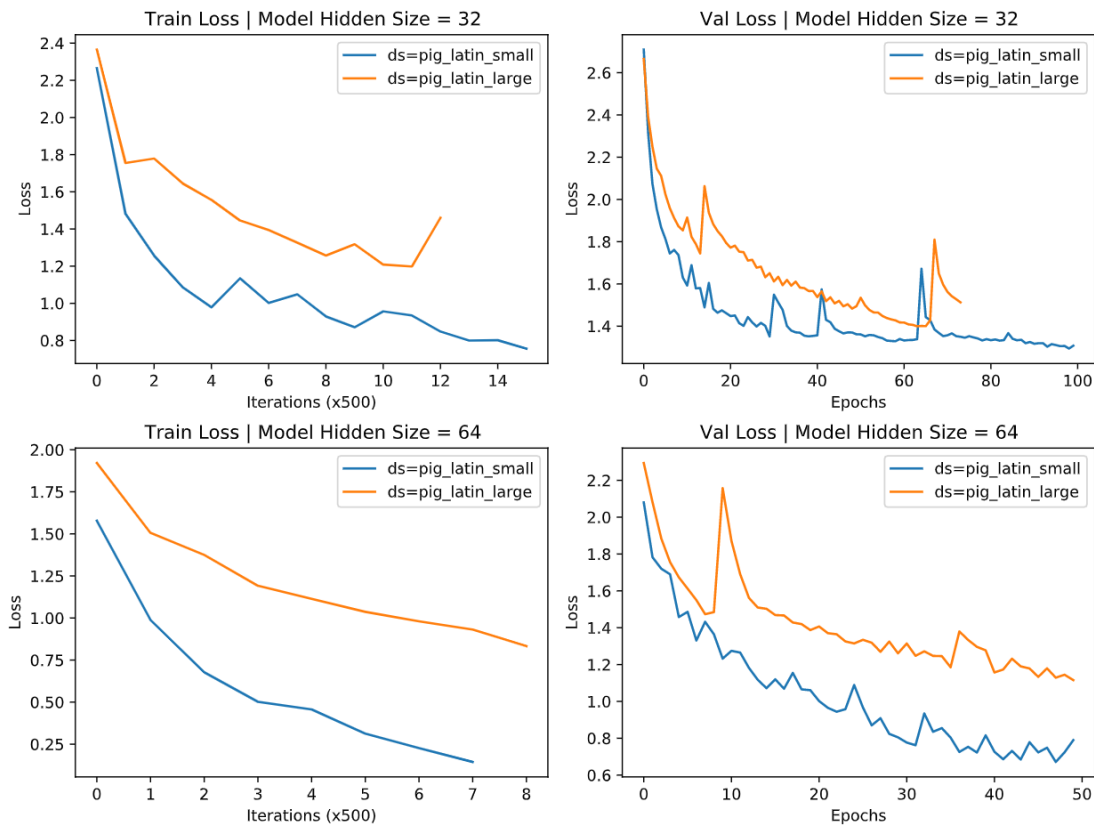
Loss Curve for save_loss_comparison_by_hidden

Performance by Hidden State Size



Loss Curve for save_loss_comparison_by_dataset

Performance by Dataset Size



Lowest Validation Loss

- small datasize, small hidden/model size: 1.29
- small datasize, large hidden/model size: 0.67
- large datasize, small hidden/model size: 1.39
- large datasize, large hidden/model size: 1.12

Analysis

When the datasize remains, large model size lowers validation loss. When the model size is constant, smaller datasize leads to lower validation loss. This behavior adheres to the intuition that large hidden layers in Transformers have the ability to extract and detect patterns, which is why larger model size decreases validation loss given the same datasize. If the size of the hidden layer is fixed, the model can achieve better result with small dataset as the Transformer is not overwhelmed by the number of training examples.

Colab FAQ

For some basic overview and features offered in Colab notebooks, check out: [Overview of Colaboratory Features](#)

You need to use the colab GPU for this assignment by selecting:

Runtime → Change runtime type → Hardware Accelerator: GPU

Part 4: Fine-tuning pretrained language models

Acknowledgement: This notebook is based on the code from <https://mccormickml.com/2019/07/22/BERT-fine-tuning/>. Credit to Chris McCormick and Nick Ryan.

Background

Fine-tuning BERT on our task of interest takes some setup. Although these steps are done for you, please take a moment to look through them and make sure you understand their purpose.

Install the [HuggingFace Transformers](#) package that contains the pretrained BERT models.

```
!pip install --upgrade transformers
```

```
Collecting transformers
```

```
  Downloading transformers-4.17.0-py3-none-any.whl (3.8 MB)
```

```
Requirement already satisfied: packaging>=20.0 in
```

```
/usr/local/lib/python3.7/dist-packages (from transformers) (21.3)
```

```
Requirement already satisfied: filelock in
```

```
/usr/local/lib/python3.7/dist-packages (from transformers) (3.6.0)
```

```
Requirement already satisfied: tqdm>=4.27 in
```

```
/usr/local/lib/python3.7/dist-packages (from transformers) (4.63.0)
```

```
Requirement already satisfied: numpy>=1.17 in
```

```
/usr/local/lib/python3.7/dist-packages (from transformers) (1.21.5)
```

```
Collecting sacremoses
```

```
  Downloading sacremoses-0.0.49-py3-none-any.whl (895 kB)
```

```
Requirement already satisfied: importlib-metadata in
```

```
/usr/local/lib/python3.7/dist-packages (from transformers) (4.11.2)
```

```
Collecting pyyaml>=5.1
```

```
  Downloading PyYAML-6.0-cp37-cp37m-
```

```
manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_12_x86_64.manylinux_2010_x86_64.whl (596 kB)
```

```
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from transformers) (2.23.0)
```

```
Collecting tokenizers!=0.11.3,>=0.11.1
```

```
  Downloading tokenizers-0.11.6-cp37-cp37m-
```

```

manylinux_2_12_x86_64.manylinux2010_x86_64.whl (6.5 MB)
Requirement already satisfied: regex!=2019.12.17 in
/usr/local/lib/python3.7/dist-packages (from transformers)
(2019.12.20)
Requirement already satisfied: typing-extensions>=3.7.4.3 in
/usr/local/lib/python3.7/dist-packages (from huggingface-
hub<1.0,>=0.1.0->transformers) (3.10.0.2)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in
/usr/local/lib/python3.7/dist-packages (from packaging>=20.0-
>transformers) (3.0.7)
Requirement already satisfied: zipp>=0.5 in
/usr/local/lib/python3.7/dist-packages (from importlib-metadata-
>transformers) (3.7.0)
Requirement already satisfied: idna<3,>=2.5 in
/usr/local/lib/python3.7/dist-packages (from requests->transformers)
(2.10)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1
in /usr/local/lib/python3.7/dist-packages (from requests-
>transformers) (1.24.3)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.7/dist-packages (from requests->transformers)
(2021.10.8)
Requirement already satisfied: chardet<4,>=3.0.2 in
/usr/local/lib/python3.7/dist-packages (from requests->transformers)
(3.0.4)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-
packages (from sacremoses->transformers) (1.15.0)
Requirement already satisfied: click in /usr/local/lib/python3.7/dist-
packages (from sacremoses->transformers) (7.1.2)
Requirement already satisfied: joblib in
/usr/local/lib/python3.7/dist-packages (from sacremoses->transformers)
(1.1.0)
Installing collected packages: pyyaml, tokenizers, sacremoses,
huggingface-hub, transformers
  Attempting uninstall: pyyaml
    Found existing installation: PyYAML 3.13
    Uninstalling PyYAML-3.13:
      Successfully uninstalled PyYAML-3.13
Successfully installed huggingface-hub-0.4.0 pyyaml-6.0 sacremoses-
0.0.49 tokenizers-0.11.6 transformers-4.17.0

```

Set the random seeds for reproducibility.

```

import os
import random

```

```

import numpy as np
import torch

```

SEED = 42


```
torch.manual_seed(SEED)
torch.cuda.manual_seed_all(SEED)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False
np.random.seed(SEED)
random.seed(SEED)
os.environ['PYTHONHASHSEED'] = str(SEED)
```

Run the following cells to download the verbal arithmetic dataset from the CSC413 webpage and load it into a DataFrame

```
!pip install wget
```

```
Collecting wget
```

```
  Downloading wget-3.2.zip (10 kB)
```

```
Building wheels for collected packages: wget
```

```
  Building wheel for wget (setup.py) ... e=wget-3.2-py3-none-any.whl
size=9675
```

```
sha256=680307cee4380a314ad0fff12fe37a5ef705fe6130a3dc9de21fedd745aa816
1
```

```
  Stored in directory:
```

```
/root/.cache/pip/wheels/a1/b6/7c/0e63e34eb06634181c63adacca38b79ff8f35
c37e3c13e3c02
```

```
Successfully built wget
```

```
Installing collected packages: wget
```

```
Successfully installed wget-3.2
```

```
import wget
import os
```

```
print('Downloading verbal arithmetic dataset')
```

```
# The URL for the dataset zip file.
```

```
url = 'https://csc413-uoft.github.io/2021/assets/misc/'
```

```
# Download the file (if we haven't already)
```

```
if not os.path.exists('./PA03_data_20_train.csv'):
    wget.download(url + 'PA03_data_20_train.csv',
                  './PA03_data_20_train.csv')
    print('Done downloading training data')
```

```
else:
    print('Already downloaded training data')
```

```
if not os.path.exists('./PA03_data_20_test.csv'):
    wget.download(url + 'PA03_data_20_test.csv',
                  './PA03_data_20_test.csv')
    print('Done downloading test data')
```

```
else:
    print('Already downloaded test data')
```

Downloading verbal arithmetic dataset
Done downloading training data
Done downloading test data

```
import pandas as pd
```

```
df = pd.read_csv("./PA03_data_20_train.csv", header=0, names=["index",  
"input", "label"])
```

```
print("Number of data points: ", df.shape[0])
```

```
sampled = df.sample(10)
```

```
# Display 10 random rows from the data.
```

```
df.sample(10)
```

Number of data points: 640

	index	input	label
20	264	thirteen plus four	2
180	712	fifteen minus twelve	2
270	327	sixteen plus seven	2
221	58	two plus eighteen	2
542	646	twelve minus six	2
143	626	eleven minus six	2
247	759	seventeen minus nineteen	0
467	691	fourteen minus eleven	2
283	372	eighteen plus twelve	2
259	274	thirteen plus fourteen	2

Tokenizer

To feed our text to BERT, it must be split into tokens, and then these tokens must be mapped to their index in the tokenizer vocabulary. For this we can use the AutoTokenizer from the transformers library.

As mentioned in the assignment handout, we will use [MathBERT](#), which uses the same architecture as BERT, but has been pretrained on text from pre-kindergarten, high-school, and college graduate level mathematical content.

```
from transformers import AutoTokenizer
```

```
bert_tokenizer = AutoTokenizer.from_pretrained('tbs17/MathBERT',  
do_lower_case=True)
```

```
{"version_major":2,"version_minor":0,"model_id":"b7c00bcb2523480cb2681  
c8816b0a647"}
```

```
{"version_major":2,"version_minor":0,"model_id":"d1c5814c3ed34dfffb4596  
0558fd2c71e"}
```

```
{"version_major":2,"version_minor":0,"model_id":"425cc192d58b4e269ad6f  
6c95bdfbdb8"}
```

```
{"version_major":2,"version_minor":0,"model_id":"bb7dd10304d64f6ebe8f8b4ff7621660"}
```

```
inputs = df.input.values
labels = df.label.values
print("Train data size ", len(inputs))
print('* Original: ', inputs[0])
# Print the sentence split into tokens.
print('* Tokenized: ', bert_tokenizer.tokenize(inputs[0]))
# Print the sentence mapped to token ids.
print('* Token IDs: ',
      bert_tokenizer.convert_tokens_to_ids(bert_tokenizer.tokenize(inputs[0]
)))
```

```
Train data size 640
* Original:    five minus twelve
* Tokenized:   ['five', 'minus', 'twelve']
* Token IDs:   [2274, 15718, 4376]
```

Formatting the inputs

In order to use BERT for fine-tuning, we need to format the inputs in a way that matches the inputs of the pretraining step. In short, we need to:

1. Add special tokens to the start and end of each sentence.
2. Pad & truncate all sentences to a single constant length.
3. Explicitly differentiate real tokens from padding tokens with the "attention mask".

Special Tokens

[SEP]

At the end of every sentence, we need to append the special [SEP] token.

This token is an artifact of two-sentence tasks, where BERT is given two separate sentences and asked to determine something (e.g., can the answer to the question in sentence A be found in sentence B?).

[CLS]

For classification tasks, we must prepend the special [CLS] token to the beginning of every sentence.

This token has special significance. BERT consists of 12 Transformer layers. Each transformer takes in a list of token embeddings, and produces the same number of embeddings on the output.

On the output of the final transformer, *only the first embedding (corresponding to the [CLS] token) is used by the classifier.*

"The first token of every sequence is always a special classification token ([CLS]). The final hidden state corresponding to this token is used as the aggregate sequence representation for classification tasks." (from the [BERT paper](#))

Also, because BERT is trained to only use this [CLS] token for classification, we know that the model has been motivated to encode everything it needs for the classification step into that single 768-value embedding vector.

Sentence Length & Attention Mask

The sentences in our dataset obviously have varying lengths, so how does BERT handle this?

BERT has two constraints:

1. All sentences must be padded or truncated to a single, fixed length.
2. The maximum sentence length is 512 tokens.

Padding is done with a special [PAD] token, which is at index 0 in the BERT vocabulary.

The "Attention Mask" is simply an array of 0s and 1s indicating which tokens are padding and which aren't.

In our dataset, all sentences have three word tokens. However, we set the max length of sentence to 7 in this example to show what paddings will be in real world applications.

```
# Set the maximum sequence length.
```

```
MAX_LEN = 7
```

```
# Print BERTs special PAD token and its index in the vocabulary
```

```
print(f'Padding token: "{bert_tokenizer.pad_token}", ID: {bert_tokenizer.pad_token_id}')
```

```
Padding token: "[PAD]", ID: 0
```

Luckily, the BertTokenizer object from the transformers library makes it easy to preprocess our input text correctly

```
tokenized_inputs = bert_tokenizer(  
    inputs.tolist(),           # Input text  
    add_special_tokens=True,   # add '[CLS]' and '[SEP]'  
    padding='max_length',     # pad to a length specified by the  
max_length  
    max_length=MAX_LEN,       # truncate all sentences longer than  
max_length  
    return_tensors='pt',      # return everything we need as PyTorch  
tensors  
)
```

```
input_ids = tokenized_inputs['input_ids']  
attention_masks = tokenized_inputs['attention_mask']
```

```

# Print sentence 0, now as a list of IDs.
print('Original: ', tokenized_inputs['input_ids'][0])
print('* Token IDs:', tokenized_inputs['attention_mask'][0])
print('* Tokenized:',
bert_tokenizer.decode(tokenized_inputs['input_ids'][0]))
print('* Attention_mask', tokenized_inputs['attention_mask'][0])

Original:  tensor([ 101, 2274, 15718, 4376, 102, 0, 0])
* Token IDs: tensor([1, 1, 1, 1, 1, 0, 0])
* Tokenized: [CLS] five minus twelve [SEP] [PAD] [PAD]
* Attention_mask tensor([1, 1, 1, 1, 1, 0, 0])

```

Training & Validation Split

Let's divide up our data into a train set (80%) and a validation set (20%).

We'll also create an iterator for our dataset using the torch DataLoader class. This helps save on memory during training because, unlike a for loop, with an iterator the entire dataset does not need to be loaded into memory.

```

from sklearn.model_selection import train_test_split
import torch
from torch.utils.data import TensorDataset, DataLoader, RandomSampler,
SequentialSampler

def train_valid_split(input_ids, attention_masks, labels,
batch_size=32):
    # Use 80% for training and 20% for validation.
    train_inputs, validation_inputs, train_masks, validation_masks,
train_labels, validation_labels = train_test_split(
        input_ids, attention_masks, labels, random_state=SEED,
test_size=0.2, stratify=labels
    )

    print('example train_input: ', train_inputs[0])
    print('example attention_mask: ', train_masks[0])

    train_labels = torch.tensor(train_labels)
    validation_labels = torch.tensor(validation_labels)

    # Create the DataLoader for our training set.
    train_data = TensorDataset(train_inputs, train_masks,
train_labels)
    train_dataloader = DataLoader(train_data, shuffle=True,
batch_size=batch_size)

    # Create the DataLoader for our validation set.
    validation_data = TensorDataset(validation_inputs,
validation_masks, validation_labels)

```

```
validation_dataloader = DataLoader(validation_data, shuffle=False,
batch_size=batch_size)
```

```
    return train_dataloader, validation_dataloader
```

```
bert_train_dataloader, bert_validation_dataloader = train_valid_split(
    input_ids=input_ids,
    attention_masks=attention_masks,
    labels=labels,
    batch_size=32
)
```

```
example train_input:    tensor([ 101, 11977, 15718, 4376, 102,
0,      0])
```

```
example attention_mask: tensor([1, 1, 1, 1, 1, 0, 0])
```

Questions

Question 1: Add a classifier to BERT [1pts]

Here, we will add a simple classifier to the [BertModel](#) provided by the Transformers library.

Your tasks are:

1. In `__init__`, add a linear classifier that will map BERTs [CLS] token representation to the unnormalized output probabilities for each class (logits).
2. In forward, pass BERTs [CLS] token representation to this new classifier to produce the logits.

In total, you won't have to write more than three new lines of code. See the comments in the code for help!

```
from transformers import BertModel
import torch.nn as nn
```

```
class BertForSentenceClassification(BertModel):
    def __init__(self, config):
        super().__init__(config)

        ##### START YOUR CODE HERE #####
        # Add a linear classifier that map BERTs [CLS] token
representation to the unnormalized
        # output probabilities for each class (logits).
        # Notes:
        # * See the documentation for torch.nn.Linear
        # * You do not need to add a softmax, as this is included in
the loss function
        # * The size of BERTs token representation can be accessed at
config.hidden_size
        # * The number of output classes can be accessed at
```

```

config.num_labels
    self.classifier = torch.nn.Linear(config.hidden_size,
config.num_labels)
    ##### END YOUR CODE HERE #####
    self.loss = torch.nn.CrossEntropyLoss()

    def forward(self, labels=None, **kwargs):
        outputs = super().forward(**kwargs)
        ##### START YOUR CODE HERE #####
        # Pass BERTs [CLS] token representation to this new classifier
        to produce the logits.
        # Notes:
        # * The [CLS] token representation can be accessed at
        outputs.pooler_output
        cls_token_repr = outputs.pooler_output
        logits = self.classifier(cls_token_repr)
        ##### END YOUR CODE HERE #####
        if labels is not None:
            outputs = (logits, self.loss(logits, labels))
        else:
            outputs = (logits,)
        return outputs

```

Question 2: Fine-tune BERT [0pts]

In this section, we will instantiate our pretrained BERT model + the new classifier, and train both on our verbal arithmetic dataset for a few epochs.

As mentioned in the assignment handout, we will use [MathBERT](#), which uses the same architecture as BERT, but has been pretrained on text from pre-kindergarten, high-school, and college graduate level mathematical content.

Although the code is written for you, please read it first to understand what it is doing. Additionally, running this code and making sure the model can be fine-tuned helps you check your implementation from Question 1. **Note:** This may print a warning: "Some weights of the model checkpoint at..." which you can ignore.

```

mathbert = BertForSentenceClassification.from_pretrained(
    "tbs17/MathBERT", # the name of the pretrained model
    num_labels=3,      # the number of classes in our downstream task
)

{"version_major":2,"version_minor":0,"model_id":"9ec178a6d4114b6c9be92f4c90aleb71"}

```

Some weights of the model checkpoint at tbs17/MathBERT were not used when initializing BertForSentenceClassification:

```

['cls.predictions.transform.LayerNorm.bias', 'cls.predictions.bias',
'cls.seq_relationship.weight', 'cls.predictions.decoder.bias',
'cls.predictions.transform.dense.weight',

```

```
'cls.predictions.transform.LayerNorm.weight',
'cls.seq_relationship.bias', 'cls.predictions.transform.dense.bias',
'cls.predictions.decoder.weight']
- This IS expected if you are initializing
BertForSentenceClassification from the checkpoint of a model trained
on another task or with another architecture (e.g. initializing a
BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing
BertForSentenceClassification from the checkpoint of a model that you
expect to be exactly identical (initializing a
BertForSequenceClassification model from a
BertForSequenceClassification model).
Some weights of BertForSentenceClassification were not initialized
from the model checkpoint at tbs17/MathBERT and are newly initialized:
['bert.classifier.bias', 'bert.classifier.weight']
You should probably TRAIN this model on a down-stream task to be able
to use it for predictions and inference.
```

The following cell prints information about the models parameters

```
# Model parameters visualization
params = list(mathbert.named_parameters())

print('The BERT model has {:} different named parameters.\n'
      .format(len(params)))

print('==== Embedding Layer ====')

for p in params[0:5]:
    print("{:<55} {:>12}".format(p[0], str(tuple(p[1].size()))))

print('\n==== First Transformer Layer ====')

for p in params[5:21]:
    print("{:<55} {:>12}".format(p[0], str(tuple(p[1].size()))))

print('\n==== Output Layer ====')

for p in params[-4:]:
    print("{:<55} {:>12}".format(p[0], str(tuple(p[1].size()))))
```

The BERT model has 201 different named parameters.

==== Embedding Layer ====

embeddings.word_embeddings.weight	(30522, 768)
embeddings.position_embeddings.weight	(512, 768)
embeddings.token_type_embeddings.weight	(2, 768)
embeddings.LayerNorm.weight	(768,)
embeddings.LayerNorm.bias	(768,)

==== First Transformer Layer ====

```
encoder.layer.0.attention.self.query.weight      (768, 768)
encoder.layer.0.attention.self.query.bias        (768,)
encoder.layer.0.attention.self.key.weight        (768, 768)
encoder.layer.0.attention.self.key.bias          (768,)
encoder.layer.0.attention.self.value.weight      (768, 768)
encoder.layer.0.attention.self.value.bias        (768,)
encoder.layer.0.attention.output.dense.weight    (768, 768)
encoder.layer.0.attention.output.dense.bias      (768,)
encoder.layer.0.attention.output.LayerNorm.weight (768,)
encoder.layer.0.attention.output.LayerNorm.bias  (768,)
encoder.layer.0.intermediate.dense.weight        (3072, 768)
encoder.layer.0.intermediate.dense.bias          (3072,)
encoder.layer.0.output.dense.weight              (768, 3072)
encoder.layer.0.output.dense.bias                (768,)
encoder.layer.0.output.LayerNorm.weight          (768,)
encoder.layer.0.output.LayerNorm.bias            (768,)
```

==== Output Layer ====

```
pooler.dense.weight      (768, 768)
pooler.dense.bias        (768,)
classifier.weight         (3, 768)
classifier.bias           (3,)
```

The next cell defines fairly standard train and evaluation loops in PyTorch

```
from torch.optim import AdamW
import time
import datetime
from transformers import get_linear_schedule_with_warmup
from tqdm import tqdm

def flat_accuracy(preds, labels):
    pred_flat = np.argmax(preds, axis=1).flatten()
    labels_flat = labels.flatten()
    return np.sum(pred_flat == labels_flat) / len(labels_flat)

def format_time(elapsed):
    elapsed_rounded = int(round((elapsed)))
    return str(datetime.timedelta(seconds=elapsed_rounded))

def get_optimizer_and_scheduler(model, total_steps, lr=2e-5,
weight_decay=0.01):
    # Apply weight decay to all parameters beside the biases or
LayerNorm weights
    no_decay = ['bias', 'LayerNorm.weight']
    optimizer_grouped_parameters = [
```

```

        {
            'params': [p for n, p in model.named_parameters() if not
any(nd in n for nd in no_decay)],
            'weight_decay': weight_decay},
        {
            'params': [p for n, p in model.named_parameters() if
any(nd in n for nd in no_decay)],
            'weight_decay': 0.0
        }
    ]
    optimizer = AdamW(model.parameters(), lr=lr)
    scheduler = get_linear_schedule_with_warmup(
        optimizer,
        # Warmup learning rate for first 10% of training steps
        num_warmup_steps=int(0.10 * total_steps),
        num_training_steps=total_steps,
    )
    return optimizer, scheduler

def train_model(model, epochs, train_dataloader,
validation_dataloader):
    # Use GPU, if available
    device = torch.device("cuda" if torch.cuda.is_available() else
"cpu")
    model = model.to(device)

    # Setup optimizer and LR scheduler
    total_steps = len(train_dataloader) * epochs
    optimizer, scheduler = get_optimizer_and_scheduler(
        model, total_steps, lr=5e-5, weight_decay=0.01
    )

    loss_values = []
    eval_accs = []

    for epoch in range(0, epochs):
        t0 = time.time()

        total_loss = 0
        model.train()

        with tqdm(train_dataloader, unit="batch") as train_pbar:
            for batch in train_pbar:
                train_pbar.set_description(f"Training (epoch {epoch +
1}))")

                b_input_ids = batch[0].to(device)
                b_input_mask = batch[1].to(device)
                b_labels = batch[2].to(device)

```

```

model.zero_grad()

# Perform a forward pass (evaluate the model on this
training batch).
# This will return the loss because we have provided
the `labels`.
outputs = model(
    input_ids=b_input_ids,
    attention_mask=b_input_mask,
    labels=b_labels
)

# The call to `model` always returns a tuple, so we
need to pull the
# loss value out of the tuple.
_, loss = outputs

# Accumulate the training loss over all of the batches
so that we can
# calculate the average loss at the end. `loss` is a
Tensor containing a
# single value; the `.item()` function just returns
the Python value
# from the tensor.
total_loss += loss.item()

# Perform a backward pass to calculate the gradients.
loss.backward()

# Clip the norm of the gradients to 1.0.
# This is to help prevent the "exploding gradients"
problem.
torch.nn.utils.clip_grad_norm_(model.parameters(),
1.0)

# Update parameters and take a step using the computed
gradient.
# The optimizer dictates the "update rule"--how the
parameters are
# modified based on their gradients, the learning
rate, etc.
optimizer.step()

# Update the learning rate.
scheduler.step()

# Calculate the average loss over the training data.
avg_train_loss = total_loss / len(train_dataloader)

```

```

    # Store the loss value for plotting the learning curve.
    loss_values.append(avg_train_loss)

    print(" * Average training loss:
{0:.2f}".format(avg_train_loss))
    print(" * Training epoch took:
{:}".format(format_time(time.time() - t0)))

    print("Running Validation...")

    t0 = time.time()
    model.eval()

    eval_loss, eval_accuracy = 0, 0
    nb_eval_steps, nb_eval_examples = 0, 0

# Evaluate data for one epoch
    for batch in validation_dataloader:
        batch = tuple(t.to(device) for t in batch)
        b_input_ids, b_input_mask, b_labels = batch

        with torch.no_grad():
            # Forward pass, calculate logit predictions.
            # This will return the logits rather than the loss
            because we have
            # not provided labels.
            # token_type_ids is the same as the "segment ids",
            which
            # differentiates sentence 1 and 2 in 2-sentence tasks.
            outputs = model(
                input_ids=b_input_ids,
                attention_mask=b_input_mask
            )

            # Get the "logits" output by the model. The "logits" are
            the output
            # values prior to applying an activation function like the
            softmax.

            logits = outputs[0]
            # Move logits and labels to CPU
            logits = logits.detach().cpu().numpy()
            label_ids = b_labels.to('cpu').numpy()
            # Calculate the accuracy for this batch of test sentences.
            tmp_eval_accuracy = flat_accuracy(logits, label_ids)
            # Accumulate the total accuracy.
            eval_accuracy += tmp_eval_accuracy
            # Track the number of batches
            nb_eval_steps += 1

```

```

        avg_eval_acc = eval_accuracy/nb_eval_steps
        print(" * Accuracy: {0:.2f}".format(avg_eval_acc))
        print(" * Validation took:
{:} ".format(format_time(time.time() - t0)))
        eval_accs.append(avg_eval_acc)
        print("Training complete!")
    return loss_values, eval_accs

```

Finally, run the following cell to fine-tune the model

```

# About 2-3 seconds per epoch using GPU
mathbert_loss_vals, mathbert_eval_accs = train_model(
    model=mathbert,
    epochs=3,
    train_dataloader=bert_train_dataloader,
    validation_dataloader=bert_validation_dataloader
)

```

```

Training (epoch 1): 100%|██████████| 16/16 [00:03<00:00, 4.78batch/s]

```

```

    * Average training loss: 0.68
    * Training epoch took: 0:00:03
Running Validation...
    * Accuracy: 0.85
    * Validation took: 0:00:00

```

```

Training (epoch 2): 100%|██████████| 16/16 [00:03<00:00, 5.02batch/s]

```

```

    * Average training loss: 0.37
    * Training epoch took: 0:00:03
Running Validation...
    * Accuracy: 0.91
    * Validation took: 0:00:00

```

```

Training (epoch 3): 100%|██████████| 16/16 [00:03<00:00, 4.99batch/s]

```

```

    * Average training loss: 0.22
    * Training epoch took: 0:00:03
Running Validation...
    * Accuracy: 0.90
    * Validation took: 0:00:00
Training complete!

```

Once the model is trained, we can plot some performance metrics

```

%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns

def plot_loss_and_acc(loss_vals, eval_accs):
    sns.set(style='darkgrid')
    sns.set(font_scale=1.5)

```

```

plt.rcParams["figure.figsize"] = (12,6)
fig, ax1 = plt.subplots(1,1)
ax1.plot(loss_vals, 'b-o', label = 'training loss')
ax2 = ax1.twinx()
ax2.plot(eval_accs, 'y-o', label = 'validation accuracy')
ax2.set_title("Training loss and validation accuracy")
ax2.set_xlabel("Epoch")
ax1.set_ylabel("Loss", color='b')
ax2.set_ylabel("Accuracy", color='y')
ax1.tick_params(axis='y', rotation=0, labelcolor='b' )
ax2.tick_params(axis='y', rotation=0, labelcolor='y' )
plt.show()

```

plot_loss_and_acc(mathbert_loss_vals, mathbert_eval_accs)



Question 3: Freezing the pretrained weights [0.5pts]

Now, let's try training the model again, except this time we will *not* fine-tune BERT's weights (we sometimes say these weights are "frozen"). To do this, we will only compute gradients for the classifiers parameters.

We can do this in pytorch by setting the `requires_grad` attribute to `False` for all parameters beside the classifiers.

Run the following cells to instantiate the model and train only the classifier. Then answer the follow-up questions in the assignment handout.

Note: This may print a warning: "Some weights of the model checkpoint at..." which you can ignore.

```

mathbert_frozen = BertForSentenceClassification.from_pretrained(
    "tbs17/MathBERT", # the name of the pretrained model
    num_labels=3,      # the number of classes in our downstream task
)

```

Some weights of the model checkpoint at tbs17/MathBERT were not used when initializing BertForSentenceClassification:

```
['cls.predictions.transform.LayerNorm.bias', 'cls.predictions.bias',  
'cls.seq_relationship.weight', 'cls.predictions.decoder.bias',  
'cls.predictions.transform.dense.weight',  
'cls.predictions.transform.LayerNorm.weight',  
'cls.seq_relationship.bias', 'cls.predictions.transform.dense.bias',  
'cls.predictions.decoder.weight']
```

- This IS expected if you are initializing BertForSentenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertForSentenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSentenceClassification were not initialized from the model checkpoint at tbs17/MathBERT and are newly initialized: ['bert.classifier.bias', 'bert.classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
for name, param in mathbert_frozen.named_parameters():  
    # Only compute gradients for parameters of our  
    # newly added classifier. BERT will not be trained.  
    if 'classifier' not in name:  
        param.requires_grad = False
```

About 1 second per epoch on GPU

```
mathbert_frozen_loss_vals, mathbert_frozen_eval_accs = train_model(  
    model=mathbert_frozen,  
    epochs=3,  
    train_dataloader=bert_train_dataloader,  
    validation_dataloader=bert_validation_dataloader  
)
```

Training (epoch 1): 100%|██████████| 16/16 [00:00<00:00, 16.33batch/s]

```
* Average training loss: 1.36  
* Training epoch took: 0:00:01
```

Running Validation...

```
* Accuracy: 0.05  
* Validation took: 0:00:00
```

Training (epoch 2): 100%|██████████| 16/16 [00:00<00:00, 17.70batch/s]

```
* Average training loss: 1.19  
* Training epoch took: 0:00:01
```

Running Validation...

```
* Accuracy: 0.12
* Validation took: 0:00:00
```

```
Training (epoch 3): 100%|██████████| 16/16 [00:00<00:00, 17.73batch/s]
```

```
* Average training loss: 1.11
* Training epoch took: 0:00:01
Running Validation...
```

```
* Accuracy: 0.30
* Validation took: 0:00:00
```

```
Training complete!
```

```
plot_loss_and_acc(mathbert_frozen_loss_vals,
mathbert_frozen_eval_accs)
```



Response to Question 3

It takes about 3 seconds to train when using fine-tuning and 1 second to train when BERT's weights are frozen. As fine-tuning without freezing the pre-trained weights has more weights, it would take more time to compute and updates all the weight parameters.

However, as no layers are frozen in fine-tuning, all the parameters are adjusted for this single sentence classification task. Therefore, the validation accuracy with fine-tuning is much higher than training with pretrained weights frozen.

Question 4: Effect of pretraining data [0.5pts]

Now, let's try fine-tuning the model again, except this time we will use [BERTtweets](#) pretrained weights. BERTtweets uses the same architecture as BERT (and MathBERT), but has been pretrained on 100s of millions of *tweets*.

Run the following cells to instantiate our model with BERTtweets pretrained weights and fine-tune it. Then answer the follow-up questions in the assignment handout.

Note: This may print a warning: "You are using a model of type..." which you can ignore.

```
bertweet = BertForSentenceClassification.from_pretrained(  
    "vinai/bertweet-base", # the name of the pretrained model  
    num_labels=3,         # the number of classes in our downstream  
    task  
)  
  
{  
    "version_major":2,"version_minor":0,"model_id":"38f5446897ef411bb1553398468ba266"}  
}
```

You are using a model of type roberta to instantiate a model of type bert. This is not supported for all configurations of models and can yield errors.

```
{  
    "version_major":2,"version_minor":0,"model_id":"d502f19f09e5417fadf95b5d58687fcc"}  
}
```

Some weights of the model checkpoint at vinai/bertweet-base were not used when initializing BertForSentenceClassification:

```
['roberta.encoder.layer.10.intermediate.dense.weight',  
'roberta.encoder.layer.9.intermediate.dense.weight',  
'roberta.encoder.layer.10.attention.output.LayerNorm.weight',  
'roberta.encoder.layer.2.attention.output.dense.weight',  
'roberta.encoder.layer.2.output.dense.bias',  
'roberta.encoder.layer.9.attention.self.value.weight',  
'roberta.encoder.layer.11.attention.output.LayerNorm.bias',  
'roberta.encoder.layer.6.output.LayerNorm.bias',  
'roberta.encoder.layer.3.attention.output.LayerNorm.bias',  
'lm_head.decoder.weight',  
'roberta.encoder.layer.3.output.LayerNorm.weight',  
'roberta.encoder.layer.3.output.dense.weight',  
'roberta.encoder.layer.1.output.LayerNorm.bias',  
'roberta.encoder.layer.9.output.dense.bias',  
'roberta.encoder.layer.11.attention.self.key.bias',  
'roberta.encoder.layer.3.output.LayerNorm.bias',  
'roberta.encoder.layer.9.intermediate.dense.bias',  
'roberta.encoder.layer.5.intermediate.dense.bias',  
'roberta.encoder.layer.2.attention.output.dense.bias',  
'roberta.encoder.layer.0.attention.self.value.weight',  
'roberta.encoder.layer.2.output.LayerNorm.weight',  
'roberta.encoder.layer.5.output.dense.bias',  
'roberta.encoder.layer.5.output.LayerNorm.bias',  
'roberta.encoder.layer.7.attention.self.key.weight',  
'roberta.encoder.layer.7.attention.self.value.weight',  
'roberta.encoder.layer.11.attention.self.query.weight',  
'roberta.encoder.layer.9.attention.output.dense.weight',  
'roberta.encoder.layer.2.output.dense.weight',  
'roberta.encoder.layer.2.attention.self.value.bias',  
'roberta.encoder.layer.3.attention.output.LayerNorm.weight',
```

'roberta.encoder.layer.9.attention.output.LayerNorm.bias',
'roberta.encoder.layer.1.attention.self.value.weight',
'roberta.encoder.layer.9.attention.self.key.bias',
'roberta.encoder.layer.3.intermediate.dense.weight',
'roberta.encoder.layer.5.output.LayerNorm.weight',
'roberta.encoder.layer.3.output.dense.bias',
'roberta.embeddings.token_type_embeddings.weight',
'roberta.encoder.layer.1.attention.output.dense.weight',
'roberta.encoder.layer.3.attention.output.dense.bias',
'roberta.encoder.layer.6.attention.self.value.weight',
'roberta.encoder.layer.4.attention.self.query.weight',
'roberta.encoder.layer.9.attention.output.dense.bias',
'roberta.encoder.layer.10.attention.self.query.weight',
'roberta.encoder.layer.1.attention.self.value.bias',
'roberta.encoder.layer.10.attention.output.dense.bias',
'roberta.pooler.dense.weight',
'roberta.encoder.layer.7.output.LayerNorm.bias',
'roberta.encoder.layer.1.intermediate.dense.weight',
'roberta.encoder.layer.0.output.dense.bias',
'roberta.encoder.layer.1.output.dense.weight',
'roberta.encoder.layer.5.attention.self.key.bias',
'roberta.encoder.layer.6.output.LayerNorm.weight',
'roberta.encoder.layer.6.attention.output.dense.bias',
'roberta.encoder.layer.0.attention.self.key.weight',
'roberta.encoder.layer.9.attention.self.key.weight',
'roberta.encoder.layer.11.output.dense.bias',
'roberta.encoder.layer.7.attention.output.LayerNorm.bias',
'roberta.encoder.layer.10.output.dense.weight',
'roberta.encoder.layer.2.attention.self.value.weight',
'roberta.encoder.layer.1.attention.output.dense.bias',
'roberta.encoder.layer.11.output.LayerNorm.weight',
'roberta.encoder.layer.4.attention.self.value.bias',
'roberta.encoder.layer.7.intermediate.dense.weight',
'roberta.encoder.layer.4.attention.self.query.bias',
'roberta.encoder.layer.11.attention.self.value.weight',
'roberta.encoder.layer.2.attention.output.LayerNorm.bias',
'roberta.encoder.layer.0.attention.output.dense.bias',
'roberta.encoder.layer.7.attention.self.value.bias',
'roberta.encoder.layer.3.attention.self.key.bias',
'roberta.encoder.layer.0.output.dense.weight',
'roberta.encoder.layer.4.intermediate.dense.weight',
'roberta.encoder.layer.6.output.dense.bias',
'roberta.encoder.layer.2.intermediate.dense.weight',
'roberta.encoder.layer.4.attention.output.dense.weight',
'roberta.encoder.layer.8.attention.output.dense.weight',
'roberta.encoder.layer.4.attention.self.value.weight',
'roberta.encoder.layer.0.attention.self.key.bias',
'roberta.encoder.layer.8.attention.self.value.bias',
'roberta.encoder.layer.8.output.LayerNorm.weight',
'roberta.encoder.layer.10.attention.self.value.weight',

'roberta.encoder.layer.8.attention.self.key.bias',
'roberta.encoder.layer.8.attention.output.LayerNorm.bias',
'roberta.encoder.layer.8.intermediate.dense.weight',
'lm_head.layer_norm.weight',
'roberta.encoder.layer.6.attention.output.LayerNorm.weight',
'roberta.encoder.layer.8.output.LayerNorm.bias',
'roberta.encoder.layer.6.attention.self.query.bias',
'roberta.encoder.layer.6.attention.output.dense.weight',
'roberta.encoder.layer.10.attention.self.key.weight',
'roberta.encoder.layer.5.attention.output.LayerNorm.bias',
'roberta.encoder.layer.6.intermediate.dense.bias',
'roberta.encoder.layer.5.attention.self.value.bias',
'roberta.encoder.layer.10.output.LayerNorm.weight',
'roberta.encoder.layer.11.output.dense.weight',
'roberta.encoder.layer.3.attention.self.value.weight',
'roberta.encoder.layer.1.attention.output.LayerNorm.weight',
'roberta.encoder.layer.10.output.dense.bias',
'roberta.encoder.layer.2.intermediate.dense.bias',
'roberta.encoder.layer.4.intermediate.dense.bias',
'roberta.encoder.layer.6.intermediate.dense.weight',
'roberta.encoder.layer.4.attention.self.key.bias',
'roberta.encoder.layer.5.intermediate.dense.weight',
'roberta.encoder.layer.7.output.dense.weight',
'roberta.encoder.layer.10.intermediate.dense.bias',
'roberta.encoder.layer.0.output.LayerNorm.weight',
'roberta.encoder.layer.4.output.LayerNorm.bias',
'roberta.encoder.layer.7.attention.output.dense.bias',
'roberta.encoder.layer.3.attention.self.query.bias',
'roberta.encoder.layer.4.attention.output.LayerNorm.weight',
'roberta.encoder.layer.1.attention.output.LayerNorm.bias',
'roberta.encoder.layer.8.attention.output.dense.bias',
'roberta.encoder.layer.9.attention.self.value.bias',
'roberta.encoder.layer.7.attention.self.query.bias',
'roberta.pooler.dense.bias', 'roberta.embeddings.LayerNorm.weight',
'roberta.encoder.layer.5.attention.self.value.weight',
'roberta.encoder.layer.0.output.LayerNorm.bias',
'roberta.encoder.layer.2.attention.self.query.bias',
'roberta.embeddings.position_ids',
'roberta.encoder.layer.1.intermediate.dense.bias',
'roberta.encoder.layer.11.attention.output.dense.bias',
'roberta.encoder.layer.8.output.dense.bias',
'roberta.encoder.layer.0.attention.output.LayerNorm.bias',
'roberta.encoder.layer.5.attention.output.LayerNorm.weight',
'roberta.encoder.layer.11.intermediate.dense.bias',
'roberta.encoder.layer.1.attention.self.key.weight',
'roberta.encoder.layer.4.attention.output.LayerNorm.bias',
'roberta.encoder.layer.7.output.LayerNorm.weight',
'roberta.encoder.layer.6.attention.self.key.weight',
'roberta.encoder.layer.8.attention.self.query.weight',
'lm_head.dense.bias',

'roberta.encoder.layer.6.attention.output.LayerNorm.bias',
'roberta.encoder.layer.3.attention.self.query.weight',
'roberta.encoder.layer.11.attention.output.LayerNorm.weight',
'roberta.encoder.layer.2.output.LayerNorm.bias',
'roberta.encoder.layer.6.attention.self.query.weight',
'roberta.embeddings.LayerNorm.bias',
'roberta.encoder.layer.11.intermediate.dense.weight',
'roberta.encoder.layer.7.attention.output.dense.weight',
'roberta.encoder.layer.9.attention.output.LayerNorm.weight',
'roberta.encoder.layer.10.attention.output.dense.weight',
'lm_head.layer_norm.bias',
'roberta.encoder.layer.4.attention.self.key.weight',
'roberta.encoder.layer.4.attention.output.dense.bias',
'roberta.encoder.layer.4.output.dense.weight',
'roberta.encoder.layer.2.attention.output.LayerNorm.weight',
'roberta.encoder.layer.1.attention.self.query.weight',
'roberta.encoder.layer.5.attention.self.key.weight',
'roberta.encoder.layer.0.attention.output.LayerNorm.weight',
'roberta.encoder.layer.1.output.LayerNorm.weight',
'roberta.encoder.layer.6.attention.self.value.bias',
'roberta.encoder.layer.0.attention.self.value.bias',
'roberta.encoder.layer.7.intermediate.dense.bias',
'roberta.encoder.layer.9.output.LayerNorm.bias',
'roberta.encoder.layer.10.attention.output.LayerNorm.bias',
'roberta.encoder.layer.6.output.dense.weight',
'roberta.encoder.layer.4.output.LayerNorm.weight',
'roberta.encoder.layer.11.attention.self.value.bias',
'roberta.encoder.layer.6.attention.self.key.bias',
'roberta.encoder.layer.5.attention.output.dense.bias',
'roberta.encoder.layer.9.attention.self.query.bias',
'roberta.encoder.layer.10.output.LayerNorm.bias',
'roberta.encoder.layer.3.attention.output.dense.weight',
'roberta.encoder.layer.1.attention.self.query.bias',
'roberta.encoder.layer.0.attention.self.query.bias',
'roberta.encoder.layer.8.attention.self.query.bias',
'roberta.encoder.layer.5.attention.output.dense.weight',
'roberta.encoder.layer.11.attention.self.query.bias',
'roberta.encoder.layer.7.attention.output.LayerNorm.weight',
'roberta.encoder.layer.8.intermediate.dense.bias',
'roberta.encoder.layer.2.attention.self.key.weight',
'roberta.encoder.layer.5.attention.self.query.weight',
'roberta.encoder.layer.7.output.dense.bias',
'roberta.encoder.layer.8.attention.self.key.weight',
'roberta.encoder.layer.11.output.LayerNorm.bias',
'roberta.encoder.layer.0.attention.output.dense.weight',
'roberta.encoder.layer.9.output.LayerNorm.weight',
'roberta.encoder.layer.1.output.dense.bias',
'roberta.encoder.layer.4.output.dense.bias',
'roberta.encoder.layer.5.output.dense.weight',
'roberta.encoder.layer.10.attention.self.key.bias',

```

'roberta.encoder.layer.8.output.dense.weight',
'roberta.encoder.layer.5.attention.self.query.bias',
'roberta.encoder.layer.0.attention.self.query.weight',
'roberta.encoder.layer.10.attention.self.value.bias',
'roberta.encoder.layer.9.attention.self.query.weight',
'roberta.encoder.layer.11.attention.self.key.weight',
'roberta.encoder.layer.0.intermediate.dense.weight',
'roberta.embeddings.position_embeddings.weight',
'roberta.encoder.layer.7.attention.self.key.bias',
'roberta.encoder.layer.7.attention.self.query.weight',
'lm_head.dense.weight', 'roberta.embeddings.word_embeddings.weight',
'roberta.encoder.layer.3.intermediate.dense.bias',
'roberta.encoder.layer.2.attention.self.key.bias',
'roberta.encoder.layer.3.attention.self.key.weight',
'roberta.encoder.layer.8.attention.output.LayerNorm.weight',
'roberta.encoder.layer.2.attention.self.query.weight',
'roberta.encoder.layer.9.output.dense.weight',
'roberta.encoder.layer.8.attention.self.value.weight',
'roberta.encoder.layer.3.attention.self.value.bias',
'roberta.encoder.layer.11.attention.output.dense.weight',
'roberta.encoder.layer.1.attention.self.key.bias', 'lm_head.bias',
'roberta.encoder.layer.0.intermediate.dense.bias',
'roberta.encoder.layer.10.attention.self.query.bias',
'lm_head.decoder.bias']

```

- This IS expected if you are initializing

BertForSentenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing

BertForSentenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSentenceClassification were not initialized from the model checkpoint at vinai/bertweet-base and are newly initialized: ['encoder.layer.11.output.dense.weight',

```

'encoder.layer.0.output.LayerNorm.bias',
'encoder.layer.3.output.LayerNorm.bias',
'encoder.layer.11.output.LayerNorm.bias',
'encoder.layer.1.attention.self.key.bias',
'encoder.layer.6.attention.output.dense.weight',
'encoder.layer.3.attention.output.LayerNorm.weight',
'embeddings.word_embeddings.weight',
'encoder.layer.5.attention.output.LayerNorm.weight',
'encoder.layer.11.attention.output.dense.weight',
'encoder.layer.10.attention.self.key.weight',
'encoder.layer.4.attention.self.key.weight',
'embeddings.LayerNorm.bias',
'embeddings.token_type_embeddings.weight',
'encoder.layer.0.attention.self.key.weight',

```

'encoder.layer.1.attention.output.LayerNorm.weight',
'encoder.layer.11.attention.self.value.weight',
'encoder.layer.9.output.dense.weight',
'encoder.layer.8.intermediate.dense.weight',
'encoder.layer.0.attention.self.key.bias',
'encoder.layer.4.output.LayerNorm.weight',
'encoder.layer.2.attention.self.value.bias',
'encoder.layer.3.attention.self.query.weight',
'encoder.layer.7.attention.self.key.weight',
'encoder.layer.8.attention.output.LayerNorm.bias',
'encoder.layer.8.attention.self.key.bias',
'encoder.layer.0.attention.output.LayerNorm.bias',
'encoder.layer.0.attention.self.query.weight',
'encoder.layer.2.intermediate.dense.bias',
'encoder.layer.3.output.LayerNorm.weight',
'encoder.layer.8.attention.output.dense.bias',
'encoder.layer.7.attention.self.value.weight',
'encoder.layer.1.output.LayerNorm.bias',
'encoder.layer.7.output.dense.weight',
'encoder.layer.11.intermediate.dense.weight',
'encoder.layer.11.output.LayerNorm.weight', 'pooler.dense.weight',
'encoder.layer.1.intermediate.dense.weight',
'encoder.layer.4.attention.output.LayerNorm.bias',
'encoder.layer.6.attention.self.value.bias',
'encoder.layer.6.output.dense.bias',
'encoder.layer.7.attention.self.query.bias',
'encoder.layer.2.attention.self.query.bias',
'encoder.layer.0.attention.output.dense.weight',
'encoder.layer.6.attention.output.LayerNorm.bias',
'encoder.layer.10.output.LayerNorm.bias',
'encoder.layer.8.attention.output.dense.weight',
'encoder.layer.2.attention.self.key.weight',
'encoder.layer.11.attention.self.query.weight',
'encoder.layer.2.attention.output.LayerNorm.bias',
'encoder.layer.1.attention.self.query.weight',
'encoder.layer.10.intermediate.dense.bias',
'encoder.layer.11.attention.self.key.bias',
'encoder.layer.10.intermediate.dense.weight',
'encoder.layer.11.attention.self.key.weight',
'encoder.layer.0.attention.self.value.bias',
'encoder.layer.5.output.LayerNorm.bias',
'encoder.layer.3.output.dense.weight',
'encoder.layer.4.attention.self.value.weight',
'encoder.layer.11.attention.output.dense.bias',
'encoder.layer.8.intermediate.dense.bias',
'encoder.layer.0.attention.output.LayerNorm.weight',
'encoder.layer.1.output.dense.weight',
'encoder.layer.8.attention.self.key.weight',
'encoder.layer.4.output.dense.weight',
'encoder.layer.1.attention.self.query.bias',

'encoder.layer.10.attention.output.LayerNorm.bias',
'encoder.layer.5.attention.self.key.bias',
'encoder.layer.3.attention.output.dense.weight',
'encoder.layer.1.attention.self.key.weight',
'encoder.layer.3.attention.output.LayerNorm.bias',
'encoder.layer.3.attention.self.value.weight',
'encoder.layer.6.output.LayerNorm.bias',
'encoder.layer.8.attention.output.LayerNorm.weight',
'encoder.layer.4.attention.self.query.weight',
'encoder.layer.0.intermediate.dense.weight',
'encoder.layer.4.output.dense.bias',
'encoder.layer.8.attention.self.query.bias',
'encoder.layer.2.output.dense.bias',
'encoder.layer.7.attention.output.LayerNorm.weight',
'encoder.layer.6.intermediate.dense.bias',
'encoder.layer.2.attention.self.key.bias',
'encoder.layer.10.output.dense.weight',
'encoder.layer.2.attention.output.LayerNorm.weight',
'encoder.layer.7.attention.output.dense.bias',
'encoder.layer.9.output.LayerNorm.weight',
'encoder.layer.2.attention.self.value.weight',
'encoder.layer.8.attention.self.value.bias',
'encoder.layer.4.attention.output.dense.bias',
'encoder.layer.3.intermediate.dense.weight',
'encoder.layer.0.attention.self.value.weight',
'encoder.layer.3.attention.self.query.bias',
'encoder.layer.5.intermediate.dense.bias',
'encoder.layer.11.attention.output.LayerNorm.bias',
'encoder.layer.8.output.LayerNorm.bias',
'encoder.layer.1.attention.self.value.bias',
'encoder.layer.4.intermediate.dense.weight',
'encoder.layer.0.attention.output.dense.bias',
'encoder.layer.5.attention.self.value.bias',
'encoder.layer.6.attention.self.query.bias',
'encoder.layer.5.attention.self.value.weight',
'encoder.layer.6.attention.self.value.weight',
'encoder.layer.7.attention.output.LayerNorm.bias',
'encoder.layer.9.attention.output.dense.bias',
'encoder.layer.10.output.LayerNorm.weight',
'encoder.layer.9.attention.self.query.bias',
'encoder.layer.0.intermediate.dense.bias',
'encoder.layer.1.output.dense.bias',
'encoder.layer.5.output.dense.weight',
'encoder.layer.7.attention.self.key.bias',
'encoder.layer.9.output.LayerNorm.bias',
'encoder.layer.2.attention.self.query.weight',
'encoder.layer.7.attention.output.dense.weight',
'encoder.layer.5.attention.self.query.weight',
'encoder.layer.5.output.LayerNorm.weight',
'encoder.layer.5.attention.output.LayerNorm.bias',

'encoder.layer.11.attention.self.value.bias',
'encoder.layer.10.attention.self.value.weight',
'encoder.layer.9.attention.output.LayerNorm.bias',
'encoder.layer.2.output.LayerNorm.weight',
'encoder.layer.3.attention.self.key.weight',
'encoder.layer.5.attention.output.dense.bias',
'encoder.layer.5.output.dense.bias',
'encoder.layer.8.output.dense.weight', 'embeddings.LayerNorm.weight',
'encoder.layer.6.attention.self.key.weight',
'encoder.layer.7.output.LayerNorm.weight',
'encoder.layer.4.attention.self.key.bias',
'embeddings.position_embeddings.weight',
'encoder.layer.6.attention.self.key.bias',
'encoder.layer.9.output.dense.bias',
'encoder.layer.10.attention.self.query.weight',
'encoder.layer.4.attention.self.query.bias',
'encoder.layer.4.attention.output.dense.weight',
'encoder.layer.8.output.LayerNorm.weight', 'classifier.weight',
'encoder.layer.5.attention.output.dense.weight',
'encoder.layer.0.output.dense.weight',
'encoder.layer.3.attention.output.dense.bias',
'encoder.layer.2.attention.output.dense.weight',
'encoder.layer.7.attention.self.value.bias',
'encoder.layer.11.attention.output.LayerNorm.weight',
'encoder.layer.6.attention.output.dense.bias',
'encoder.layer.0.output.LayerNorm.weight',
'encoder.layer.10.attention.output.dense.weight',
'encoder.layer.0.attention.self.query.bias',
'encoder.layer.0.output.dense.bias',
'encoder.layer.6.attention.self.query.weight',
'encoder.layer.10.attention.self.value.bias',
'encoder.layer.10.attention.output.dense.bias',
'encoder.layer.9.attention.self.value.bias',
'encoder.layer.11.intermediate.dense.bias',
'encoder.layer.1.attention.self.value.weight',
'encoder.layer.7.output.dense.bias',
'encoder.layer.7.output.LayerNorm.bias',
'encoder.layer.9.attention.output.LayerNorm.weight',
'encoder.layer.4.output.LayerNorm.bias',
'encoder.layer.8.output.dense.bias',
'encoder.layer.1.output.LayerNorm.weight',
'encoder.layer.6.attention.output.LayerNorm.weight',
'encoder.layer.2.output.dense.weight',
'encoder.layer.5.intermediate.dense.weight',
'encoder.layer.6.output.LayerNorm.weight',
'encoder.layer.9.attention.self.key.weight',
'encoder.layer.1.attention.output.dense.bias',
'encoder.layer.9.attention.self.key.bias',
'encoder.layer.9.attention.self.query.weight',
'encoder.layer.10.attention.self.query.bias',


```

'encoder.layer.8.attention.self.query.weight',
'encoder.layer.3.intermediate.dense.bias',
'encoder.layer.9.intermediate.dense.weight',
'encoder.layer.11.output.dense.bias',
'encoder.layer.1.attention.output.LayerNorm.bias',
'encoder.layer.9.intermediate.dense.bias',
'encoder.layer.10.attention.output.LayerNorm.weight',
'encoder.layer.11.attention.self.query.bias',
'encoder.layer.8.attention.self.value.weight',
'encoder.layer.7.intermediate.dense.weight',
'encoder.layer.6.output.dense.weight',
'encoder.layer.2.attention.output.dense.bias',
'encoder.layer.9.attention.output.dense.weight',
'encoder.layer.6.intermediate.dense.weight',
'encoder.layer.10.output.dense.bias', 'pooler.dense.bias',
'encoder.layer.4.attention.self.value.bias',
'encoder.layer.4.attention.output.LayerNorm.weight',
'encoder.layer.5.attention.self.query.bias',
'encoder.layer.4.intermediate.dense.bias',
'encoder.layer.3.attention.self.key.bias',
'encoder.layer.3.output.dense.bias', 'classifier.bias',
'encoder.layer.5.attention.self.key.weight',
'encoder.layer.1.attention.output.dense.weight',
'encoder.layer.2.output.LayerNorm.bias',
'encoder.layer.7.intermediate.dense.bias',
'encoder.layer.10.attention.self.key.bias',
'encoder.layer.7.attention.self.query.weight',
'encoder.layer.3.attention.self.value.bias',
'encoder.layer.1.intermediate.dense.bias',
'encoder.layer.2.intermediate.dense.weight',
'encoder.layer.9.attention.self.value.weight']

```

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

BERTweets has its own tokenizer, so we have to repeat the data loading process

```
from transformers import AutoTokenizer
```

```
bertweet_tokenizer = AutoTokenizer.from_pretrained('vinai/bertweet-
base', do_lower_case=True)
```

```
tokenized_inputs = bertweet_tokenizer(
    inputs.tolist(),
    add_special_tokens=True,
    padding='max_length',
    max_length=MAX_LEN,
    return_tensors='pt',
)
```

```
bert_train_dataloader, bert_validation_dataloader = train_valid_split(
```

```

        input_ids=tokenized_inputs['input_ids'],
        attention_masks=tokenized_inputs['attention_mask'],
        labels=labels,
        batch_size=32
    )

{"version_major":2,"version_minor":0,"model_id":"5923335ee9754c179af07b12d5224fc1"}

{"version_major":2,"version_minor":0,"model_id":"7ec021c2f46842f4a0d8976367b4f35b"}

```

emoji is not installed, thus not converting emoticons or emojis into text. Please install emoji: `pip3 install emoji`
 Special tokens have been added in the vocabulary, make sure the associated word embeddings are fine-tuned or trained.

```

example train_input:    tensor([    0, 57641, 12309, 15103,     2,
1,     1])
example attention_mask: tensor([1, 1, 1, 1, 1, 0, 0])

```

```

# About 2-3 seconds per epoch on GPU
bertweet_loss_vals, bertweet__eval_accs = train_model(
    model=bertweet,
    epochs=3,
    train_dataloader=bert_train_dataloader,
    validation_dataloader=bert_validation_dataloader
)

```

```

Training (epoch 1): 100%|██████████| 16/16 [00:03<00:00,  4.63batch/s]

```

```

    * Average training loss: 0.70
    * Training epoch took: 0:00:03
Running Validation...
    * Accuracy: 0.73
    * Validation took: 0:00:00

```

```

Training (epoch 2): 100%|██████████| 16/16 [00:03<00:00,  4.30batch/s]

```

```

    * Average training loss: 0.47
    * Training epoch took: 0:00:04
Running Validation...
    * Accuracy: 0.75
    * Validation took: 0:00:00

```

```

Training (epoch 3): 100%|██████████| 16/16 [00:04<00:00,  3.96batch/s]

```

```

    * Average training loss: 0.46
    * Training epoch took: 0:00:04
Running Validation...
    * Accuracy: 0.76
    * Validation took: 0:00:00

```

Training complete!

```
plot_loss_and_acc(bertweet_loss_vals, bertweet_eval_accs)
```



Response to Question 4

When fine-tuning BERT with BERTweet, its validation accuracy is lower than fine-tuning than MathBERT. Fine-tuning with MathBERT performs better because it contains words and formulas, which flows better logically. Therefore, BERT can extract more relevant information from MathBERT than English Tweets (BERTweet).

Question 5: Inspect models predictions [0pts]

In the following cell, we have provided a function that allows you to inspect the models predictions. Given an input, e.g. "three minus two minus two", it will return a trained models prediction i.e. "negative", "zero", or "positive".

Compare the performance of mathbert, mathbert_frozen and bertweet. Try a few unseen examples of arithmetic questions using all models. Can you find examples where one model clearly outperforms the others? Can you find examples where all models perform poorly?

```
def what_is(input, model, tokenizer):
    # Use GPU, if available
    device = torch.device("cuda" if torch.cuda.is_available() else
"cpu")
    model = model.to(device)

    # Get map of human readable outputs
    index_to_sentiment_map = {0: "negative", 1: "zero", 2: "positive"}

    tokenized_inputs = tokenizer(
        input,                # Input text
        add_special_tokens=True, # add '[CLS]' and '[SEP]'
        padding='max_length',   # pad to a length specified by the
max_length
```

```

        max_length=MAX_LEN,          # truncate all sentences longer than
max_length
        return_tensors='pt',         # return everything we need as
PyTorch tensors
    )

    input_ids = tokenized_inputs['input_ids'].to(device)
    attention_masks = tokenized_inputs['attention_mask'].to(device)

    with torch.no_grad():
        outputs = model(input_ids=input_ids,
attention_mask=attention_masks)
        logits = outputs[0]
        logits = logits.detach().cpu().numpy()
        print(index_to_sentiment_map[np.argmax(logits, axis=1)[0]])

what_is("three minus five", model=mathbert, tokenizer=bert_tokenizer)

negative

what_is("three minus five", model=mathbert_frozen,
tokenizer=bert_tokenizer)

zero

what_is("three minus five", model=bertweet, tokenizer=bert_tokenizer)

positive

```

Colab FAQ

For some basic overview and features offered in Colab notebooks, check out: [Overview of Colaboratory Features](#)

You need to use the colab GPU for this assignment by selecting:

Runtime → Change runtime type → Hardware Accelerator: GPU

Part 4: Connecting Text and Images with CLIP

Acknowledgement: This notebook is based on the code from https://colab.research.google.com/github/openai/clip/blob/master/notebooks/Interacting_with_CLIP.ipynb. Credit to OpenAI.

Section I: Interacting with CLIP

This is a self-contained notebook that shows how to download and run CLIP models, calculate the similarity between arbitrary image and text inputs, and perform zero-shot image classifications. The next cells will install the clip package and its dependencies, and check if PyTorch 1.7.1 or later is installed.

```
! pip install ftfy regex tqdm
! pip install git+https://github.com/openai/CLIP.git
```

Collecting ftfy

Downloading ftfy-6.1.1-py3-none-any.whl (53 kB)
Requirement already satisfied: regex in /usr/local/lib/python3.7/dist-packages (2019.12.20)

Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (4.63.0)

Requirement already satisfied: wcwidth>=0.2.5 in /usr/local/lib/python3.7/dist-packages (from ftfy) (0.2.5)

Installing collected packages: ftfy

Successfully installed ftfy-6.1.1

Collecting git+https://github.com/openai/CLIP.git

Cloning https://github.com/openai/CLIP.git to /tmp/pip-req-build-ic_9stf7

Running command git clone -q https://github.com/openai/CLIP.git /tmp/pip-req-build-ic_9stf7

Requirement already satisfied: ftfy in /usr/local/lib/python3.7/dist-packages (from clip==1.0) (6.1.1)

Requirement already satisfied: regex in /usr/local/lib/python3.7/dist-packages (from clip==1.0) (2019.12.20)

Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from clip==1.0) (4.63.0)

```

Requirement already satisfied: torch in /usr/local/lib/python3.7/dist-packages (from clip==1.0) (1.10.0+cu111)
Requirement already satisfied: torchvision in /usr/local/lib/python3.7/dist-packages (from clip==1.0) (0.11.1+cu111)
Requirement already satisfied: wcwidth>=0.2.5 in /usr/local/lib/python3.7/dist-packages (from ftfy->clip==1.0) (0.2.5)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch->clip==1.0) (3.10.0.2)
Requirement already satisfied: pillow!=8.3.0,>=5.3.0 in /usr/local/lib/python3.7/dist-packages (from torchvision->clip==1.0) (7.1.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torchvision->clip==1.0) (1.21.5)
Building wheels for collected packages: clip
  Building wheel for clip (setup.py) ... e=clip-1.0-py3-none-any.whl
size=1369221
sha256=2a4f078a68b65e2692c4f9487733386c8364e755604554c16e5ef0b0cd509fb
c
  Stored in directory:
/tmp/pip-ephem-wheel-cache-5_13w4i1/wheels/fd/b9/c3/5b4470e35ed76e174b
ff77c92f91da82098d5e35fd5bc8cdac
Successfully built clip
Installing collected packages: clip
Successfully installed clip-1.0

```

```

import numpy as np
import torch

```

```

print("Torch version:", torch.__version__)
torch_version = torch.__version__.split(".")
assert (int(torch_version[0]) == 1 and int(torch_version[1]) >= 7) or
int(torch_version[0]) > 1, "PyTorch 1.7.1 or later is required"

```

```
Torch version: 1.10.0+cu111
```

Loading the model

clip.available_models() will list the names of available CLIP models.

```
import clip
```

```
clip.available_models()
```

```

['RN50',
 'RN101',
 'RN50x4',
 'RN50x16',
 'RN50x64',
 'ViT-B/32',

```

```

'ViT-B/16',
'ViT-L/14']

model, preprocess = clip.load("ViT-B/32")
model.cuda().eval()
input_resolution = model.visual.input_resolution
context_length = model.context_length
vocab_size = model.vocab_size

print("Model parameters:", f"{np.sum([int(np.prod(p.shape)) for p in
model.parameters()]):,}")
print("Input resolution:", input_resolution)
print("Context length:", context_length)
print("Vocab size:", vocab_size)

100%|████████████████████████████████████████| 338M/338M [00:04<00:00,
77.8MiB/s]

Model parameters: 151,277,313
Input resolution: 224
Context length: 77
Vocab size: 49408

```

Image Preprocessing

We resize the input images and center-crop them to conform with the image resolution that the model expects. Before doing so, we will normalize the pixel intensity using the dataset mean and standard deviation.

The second return value from `clip.load()` contains a torchvision Transform that performs this preprocessing.

```

preprocess

Compose(
  Resize(size=224, interpolation=bicubic, max_size=None,
  antialias=None)
  CenterCrop(size=(224, 224))
  <function _convert_image_to_rgb at 0x7f8774a089e0>
  ToTensor()
  Normalize(mean=(0.48145466, 0.4578275, 0.40821073),
  std=(0.26862954, 0.26130258, 0.27577711))
)

```

Text Preprocessing

We use a case-insensitive tokenizer, which can be invoked using `clip.tokenize()`. By default, the outputs are padded to become 77 tokens long, which is what the CLIP models expects.

```
clip.tokenize("Hello World!")
```

```

tensor([[49406, 3306, 1002, 256, 49407, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0]])

```

Setting up input images and texts

We are going to feed 8 example images and their textual descriptions to the model, and compare the similarity between the corresponding features.

The tokenizer is case-insensitive, and we can freely give any suitable textual descriptions.

```

import os
import skimage
import IPython.display
import matplotlib.pyplot as plt
from PIL import Image
import numpy as np

from collections import OrderedDict
import torch

%matplotlib inline
%config InlineBackend.figure_format = 'retina'

# images in skimage to use and their textual descriptions
descriptions = {
    "page": "a page of text about segmentation",
    "chelsea": "a facial photo of a tabby cat",
    "astronaut": "a portrait of an astronaut with the American flag",
    "rocket": "a rocket standing on a launchpad",
    "motorcycle_right": "a red motorcycle standing in a garage",
    "camera": "a person looking at a camera on a tripod",
    "horse": "a black-and-white silhouette of a horse",
    "coffee": "a cup of coffee on a saucer"
}

original_images = []
images = []
texts = []

```



```
plt.figure(figsize=(16, 5))


for filename in [filename for filename in os.listdir(skimage.data_dir)
if filename.endswith(".png") or filename.endswith(".jpg")]:
    name = os.path.splitext(filename)[0]
    if name not in descriptions:
        continue

    image = Image.open(os.path.join(skimage.data_dir,
filename)).convert("RGB")
    #print(image.__dict__.keys())
    #print(image._size)
    #image_sequence = image.getdata()
    #image_array = np.array(image_sequence)
    #print(image_array.shape)

    plt.subplot(2, 4, len(images) + 1)
    plt.imshow(image)
    plt.title(f"{filename}\n{descriptions[name]}")
    plt.xticks([])
    plt.yticks([])

    original_images.append(image)
    images.append(preprocess(image))
    texts.append(descriptions[name])

plt.tight_layout()
```



camera.png
a person looking at a camera on a tripod



astronaut.png
a portrait of an astronaut with the American flag



coffee.png
a cup of coffee on a saucer




chelsea.png
a facial photo of a tabby cat



motorcycle_right.png
a red motorcycle standing in a garage



horse.png
a black-and-white silhouette of a horse



rocket.jpg
a rocket standing on a launchpad



Building features

We normalize the images, tokenize each text input, and run the forward pass of the model to get the image and text features.

```
image_input = torch.tensor(np.stack(images)).cuda()
text_tokens = clip.tokenize(["This is " + desc for desc in
texts]).cuda()
```

```

with torch.no_grad():
    image_features = model.encode_image(image_input).float()
    text_features = model.encode_text(text_tokens).float()

```

Calculating cosine similarity

We normalize the features and calculate the dot product of each pair.

```

image_features /= image_features.norm(dim=-1, keepdim=True)
text_features /= text_features.norm(dim=-1, keepdim=True)
similarity = text_features.cpu().numpy() @
image_features.cpu().numpy().T

count = len(descriptions)

plt.figure(figsize=(20, 14))
plt.imshow(similarity, vmin=0.1, vmax=0.3)
# plt.colorbar()
plt.yticks(range(count), texts, fontsize=18)
plt.xticks([])
for i, image in enumerate(original_images):
    plt.imshow(image, extent=(i - 0.5, i + 0.5, -1.6, -0.6),
origin="lower")
for x in range(similarity.shape[1]):
    for y in range(similarity.shape[0]):
        plt.text(x, y, f"{similarity[y, x]:.2f}", ha="center",
va="center", size=12)

for side in ["left", "top", "right", "bottom"]:
    plt.gca().spines[side].set_visible(False)

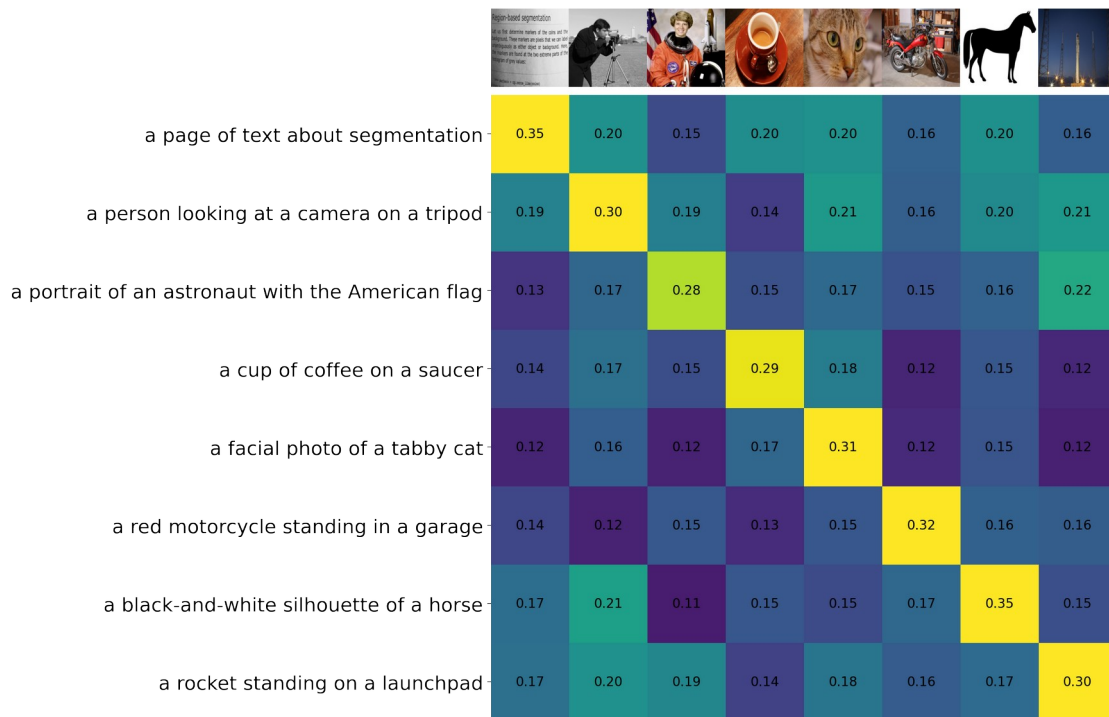
plt.xlim([-0.5, count - 0.5])
plt.ylim([count + 0.5, -2])

plt.title("Cosine similarity between text and image features",
size=20)

Text(0.5, 1.0, 'Cosine similarity between text and image features')

```

Cosine similarity between text and image features



Zero-Shot Image Classification

You can classify images using the cosine similarity (times 100) as the logits to the softmax operation.

```
from torchvision.datasets import CIFAR100
```

```
cifar100 = CIFAR100(os.path.expanduser("~/cache"),
transform=preprocess, download=True)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-100-python.tar.gz>
to /root/.cache/cifar-100-python.tar.gz

```
{"version_major":2,"version_minor":0,"model_id":"dea5e0c9f5014e4a9c1015cab538d1e9"}
```

Extracting /root/.cache/cifar-100-python.tar.gz to /root/.cache

```
text_descriptions = [f"This is a photo of a {label}" for label in
cifar100.classes]
text_tokens = clip.tokenize(text_descriptions).cuda()
```

```
with torch.no_grad():
    text_features = model.encode_text(text_tokens).float()
    text_features /= text_features.norm(dim=-1, keepdim=True)
```

```

text_probs = (100.0 * image_features @ text_features.T).softmax(dim=-1)
top_probs, top_labels = text_probs.cpu().topk(5, dim=-1)

plt.figure(figsize=(16, 16))











for i, image in enumerate(original_images):
    plt.subplot(4, 4, 2 * i + 1)
    plt.imshow(image)
    plt.axis("off")

    plt.subplot(4, 4, 2 * i + 2)
    y = np.arange(top_probs.shape[-1])
    plt.grid()
    plt.barh(y, top_probs[i])
    plt.gca().invert_yaxis()
    plt.gca().set_axisbelow(True)
    plt.yticks(y, [cifar100.classes[index] for index in
top_labels[i].numpy()])
    plt.xlabel("probability")

plt.subplots_adjust(wspace=0.5)
plt.show()

```

Region-based segmentation
Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of gray values.



Section II: Now let's do a Scavenger Hunt!

We want you to figure out what caption best describes the image below. We will run your caption against images in ImageNet and display the image with the highest network probability. The goal is that your caption paired with the image below will give the highest network output.



We will download a subset of ImageNet called Tiny ImageNet. Tiny ImageNet has only 200 classes, with each class having 500 training images, 50 validation images and 50 test images.

```
! git clone https://github.com/seshuad/IMagenet
```

```
Cloning into 'IMagenet'...
```

```
remote: Enumerating objects: 120594, done.ote: Total 120594 (delta 0),  
reused 0 (delta 0), pack-reused 120594
```

In order to reduce time and memory consumption, we will only consider the first 1000 images in the test set as the possible search space.

```
import os  
img_paths = []  
for rootdir, subdir, filenames in os.walk("IMagenet/tiny-imagenet-  
200/test/images"):  
    for file_ in sorted(filenames)[:1000]:  
        img_paths.append(os.path.join(rootdir, file_))
```

```
'''
```

```
TO DO: change caption below to produce target image
```

```
'''
```

```
# caption = "butterfly on a/one flower"  
caption = "butterfly on purple flower"
```

```
# the search process can be found at the bottom of this file
```

Now, we will run the model for the first 1000 images in the Tiny ImageNet test set. We will display the image that produces the highest network probability with your written caption

```
original_images = []  
images = []
```

```

for img_path in img_paths:

    image = Image.open(img_path).convert("RGB")
    original_images.append(image)
    images.append(preprocess(image))

image_input = torch.tensor(np.stack(images)).cuda()
with torch.no_grad():
    image_features = model.encode_image(image_input).float()
    image_features /= image_features.norm(dim=-1, keepdim=True)

text_tokens = clip.tokenize(caption).cuda()

with torch.no_grad():
    text_features = model.encode_text(text_tokens).float()
    text_features /= text_features.norm(dim=-1, keepdim=True)

text_probs = (100.0 * image_features @
text_features.T).softmax(dim=0).cpu().detach().numpy()
highest_prob = np.argmax(text_probs)
plt.axis('off')
plt.imshow(original_images[highest_prob])

<matplotlib.image.AxesImage at 0x7f875b003250>

```



Search Process

To generate the image, I used caption = "butterfly on purple flower".

After trying butterfly on a flower and butterfly on one flower (failed attempts), I realized I need to include more details of the image (which is try I tried specifying the color of the flower and it worked).